

DEVIEW  
2018

# Druid로 쉽고 빠르게 데이터 분석하기

송은혜

NAVER

**NAVER**

# CONTENS

1. 데이터 분석의 필요성

2. Druid? 그게 뭔데?

- Apache Druid vs. Elasticsearch vs. Apache Kudu

3. Druid로 분석하기 – 기본편

4. Druid로 성능 향상시키기 – 고급편

- 질의성능 향상시키기

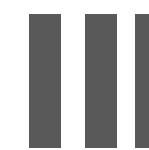
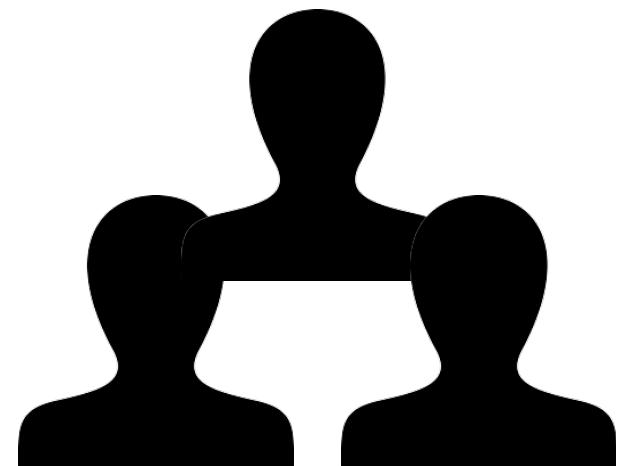
- 저장성능 향상시키기

- 편한 운영을 위한 노력

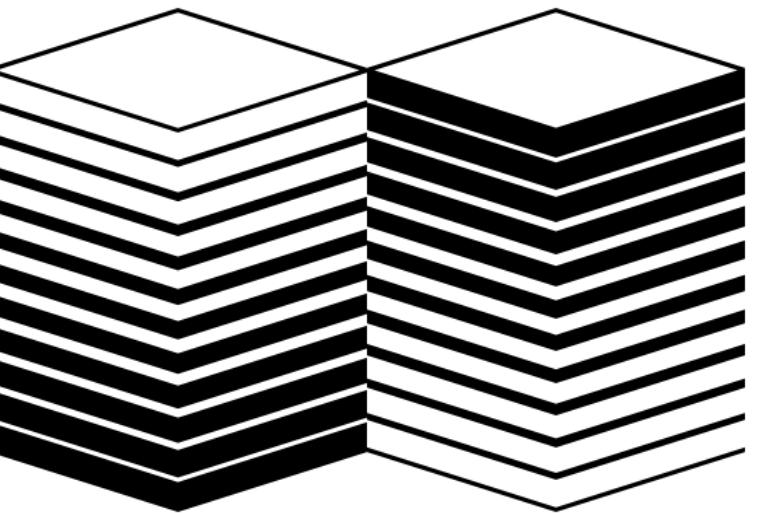
1.

# 데이터 분석의 필요성

# 1. 데이터 분석의 필요성



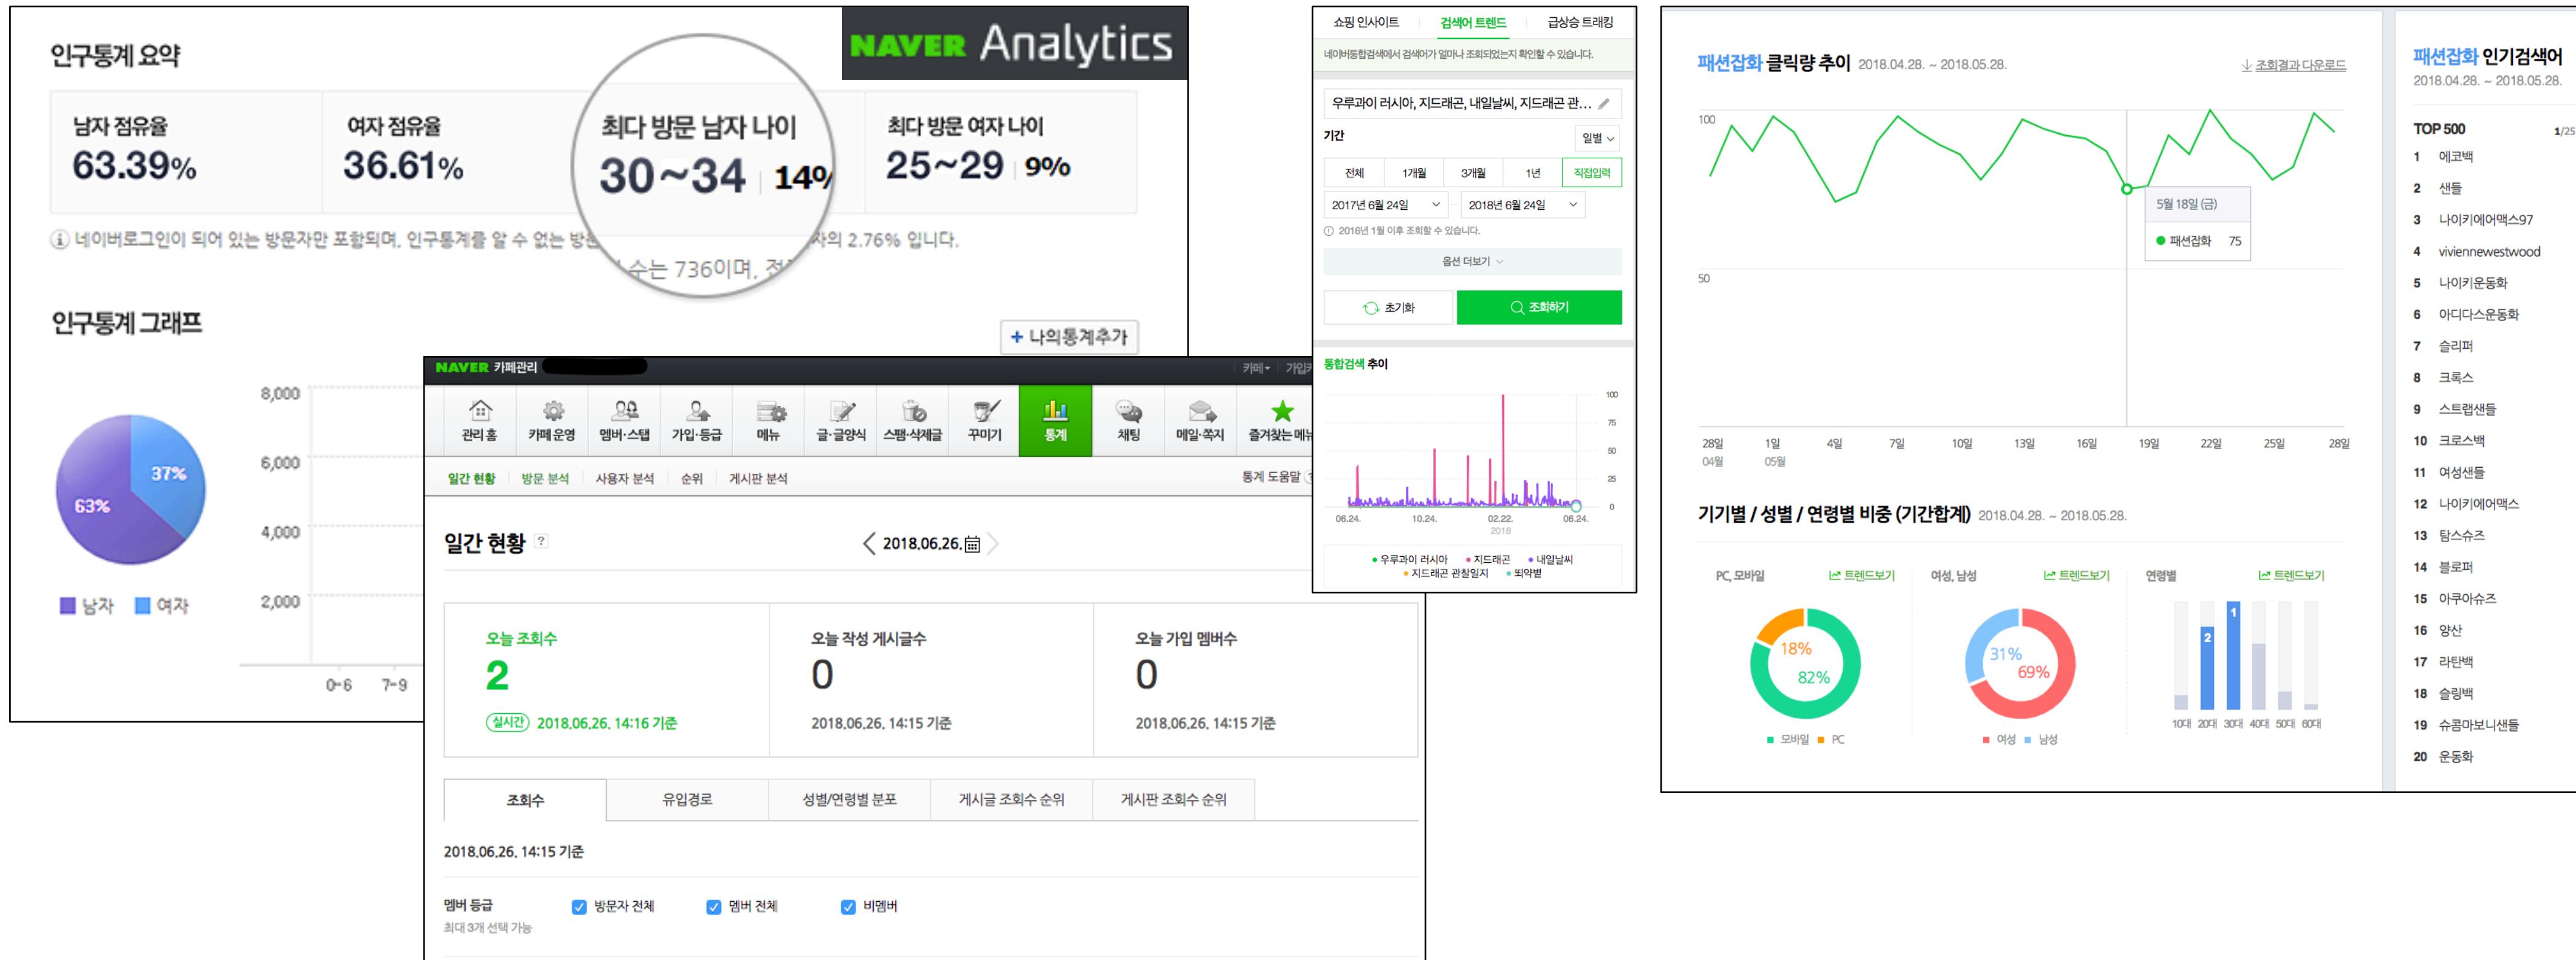
로그



Created by Magicon  
from Noun Project

현재에 대한 평가  
서비스 품질향상  
Data-Driven ...

# 1. 데이터 분석의 필요성



2.

# Druid? 그게 뭔데?



## 2.1 druid , 그게 뭔데?

- 컬럼기반 분산저장소
- Time 기반 partitioning
- 실시간/배치 데이터 저장기능 제공
- 데이터 질의기능 제공

Next Session.  
"스트림 저장소 최적화 이야기"

## 2.2 Druid vs. Elasticsearch vs. Kudu

### 검색어트렌드/쇼핑인사이트

The screenshot shows a search interface with the following sections:

- 기간**: Date range selection from 2017-06-25 to 2018-06-25. A note below says "· 2016년 1월 이후 조회할 수 있습니다." (Can be viewed from January 2016).
- 범위**: Category selection with checkboxes for 전체 (All), 모바일 (Mobile), and PC.
- 성별**: Gender selection with checkboxes for 전체 (All), 여성 (Female), and 남성 (Male).
- 연령선택**: Age group selection with checkboxes for various ranges: ~12, 13~18, 19~24, 25~29, 30~34, 35~39, 40~44, 45~49, 50~54, 55~60, and 60~.
- 검색**: A large green button labeled "네이버 검색 데이터 조회" (Search data query).

### 요구사항

- 임의의 조건 조합 분석 (원본 < 조합)
- 데이터 유효기간기간 없음  
(데이터 감소X)
- 내부 분석환경
- 개발/운영비용 최소화 (관리인원 5명)

## 2.2 Druid vs. Elasticsearch vs. Kudu

후보 선정이유 – 데이터 Size, 운영비용

Apache Druid  
(v0.10.0)

- 저장/질의기능 제공
- 빠른 필터처리

Elasticsearch  
(v2.3)

- 기존 사용 플랫폼
- 저장/질의기능 제공

Apache Kudu  
(v1.4)

- 가장 적은 저장공간
- 빠른 조회(scan)기능
- 질의기능 미제공  
(질의엔진 필요)

# 2.2 Druid vs. Elasticsearch vs. Kudu

응답시간 비교실험 – 질의 패턴

(2.1억건, 동일 병렬환경)

Random Access

```
SELECT count(*)  
      FROM logs  
     WHERE dimA = ?
```

Ranking

```
SELECT url,  
       count(*)  
     FROM logs  
    GROUP BY url  
   ORDER BY  
         count(*) desc  
        LIMIT 10
```

Group By

```
SELECT user,  
       count(*)  
     FROM logs  
    GROUP BY user
```

Join

```
SELECT ...  
      FROM logs  
INNER JOIN users  
        ON ...
```

# 2.2 Druid vs. Elasticsearch vs. Kudu

## 응답시간 비교실험

(2.1억건, 동일 병렬환경)

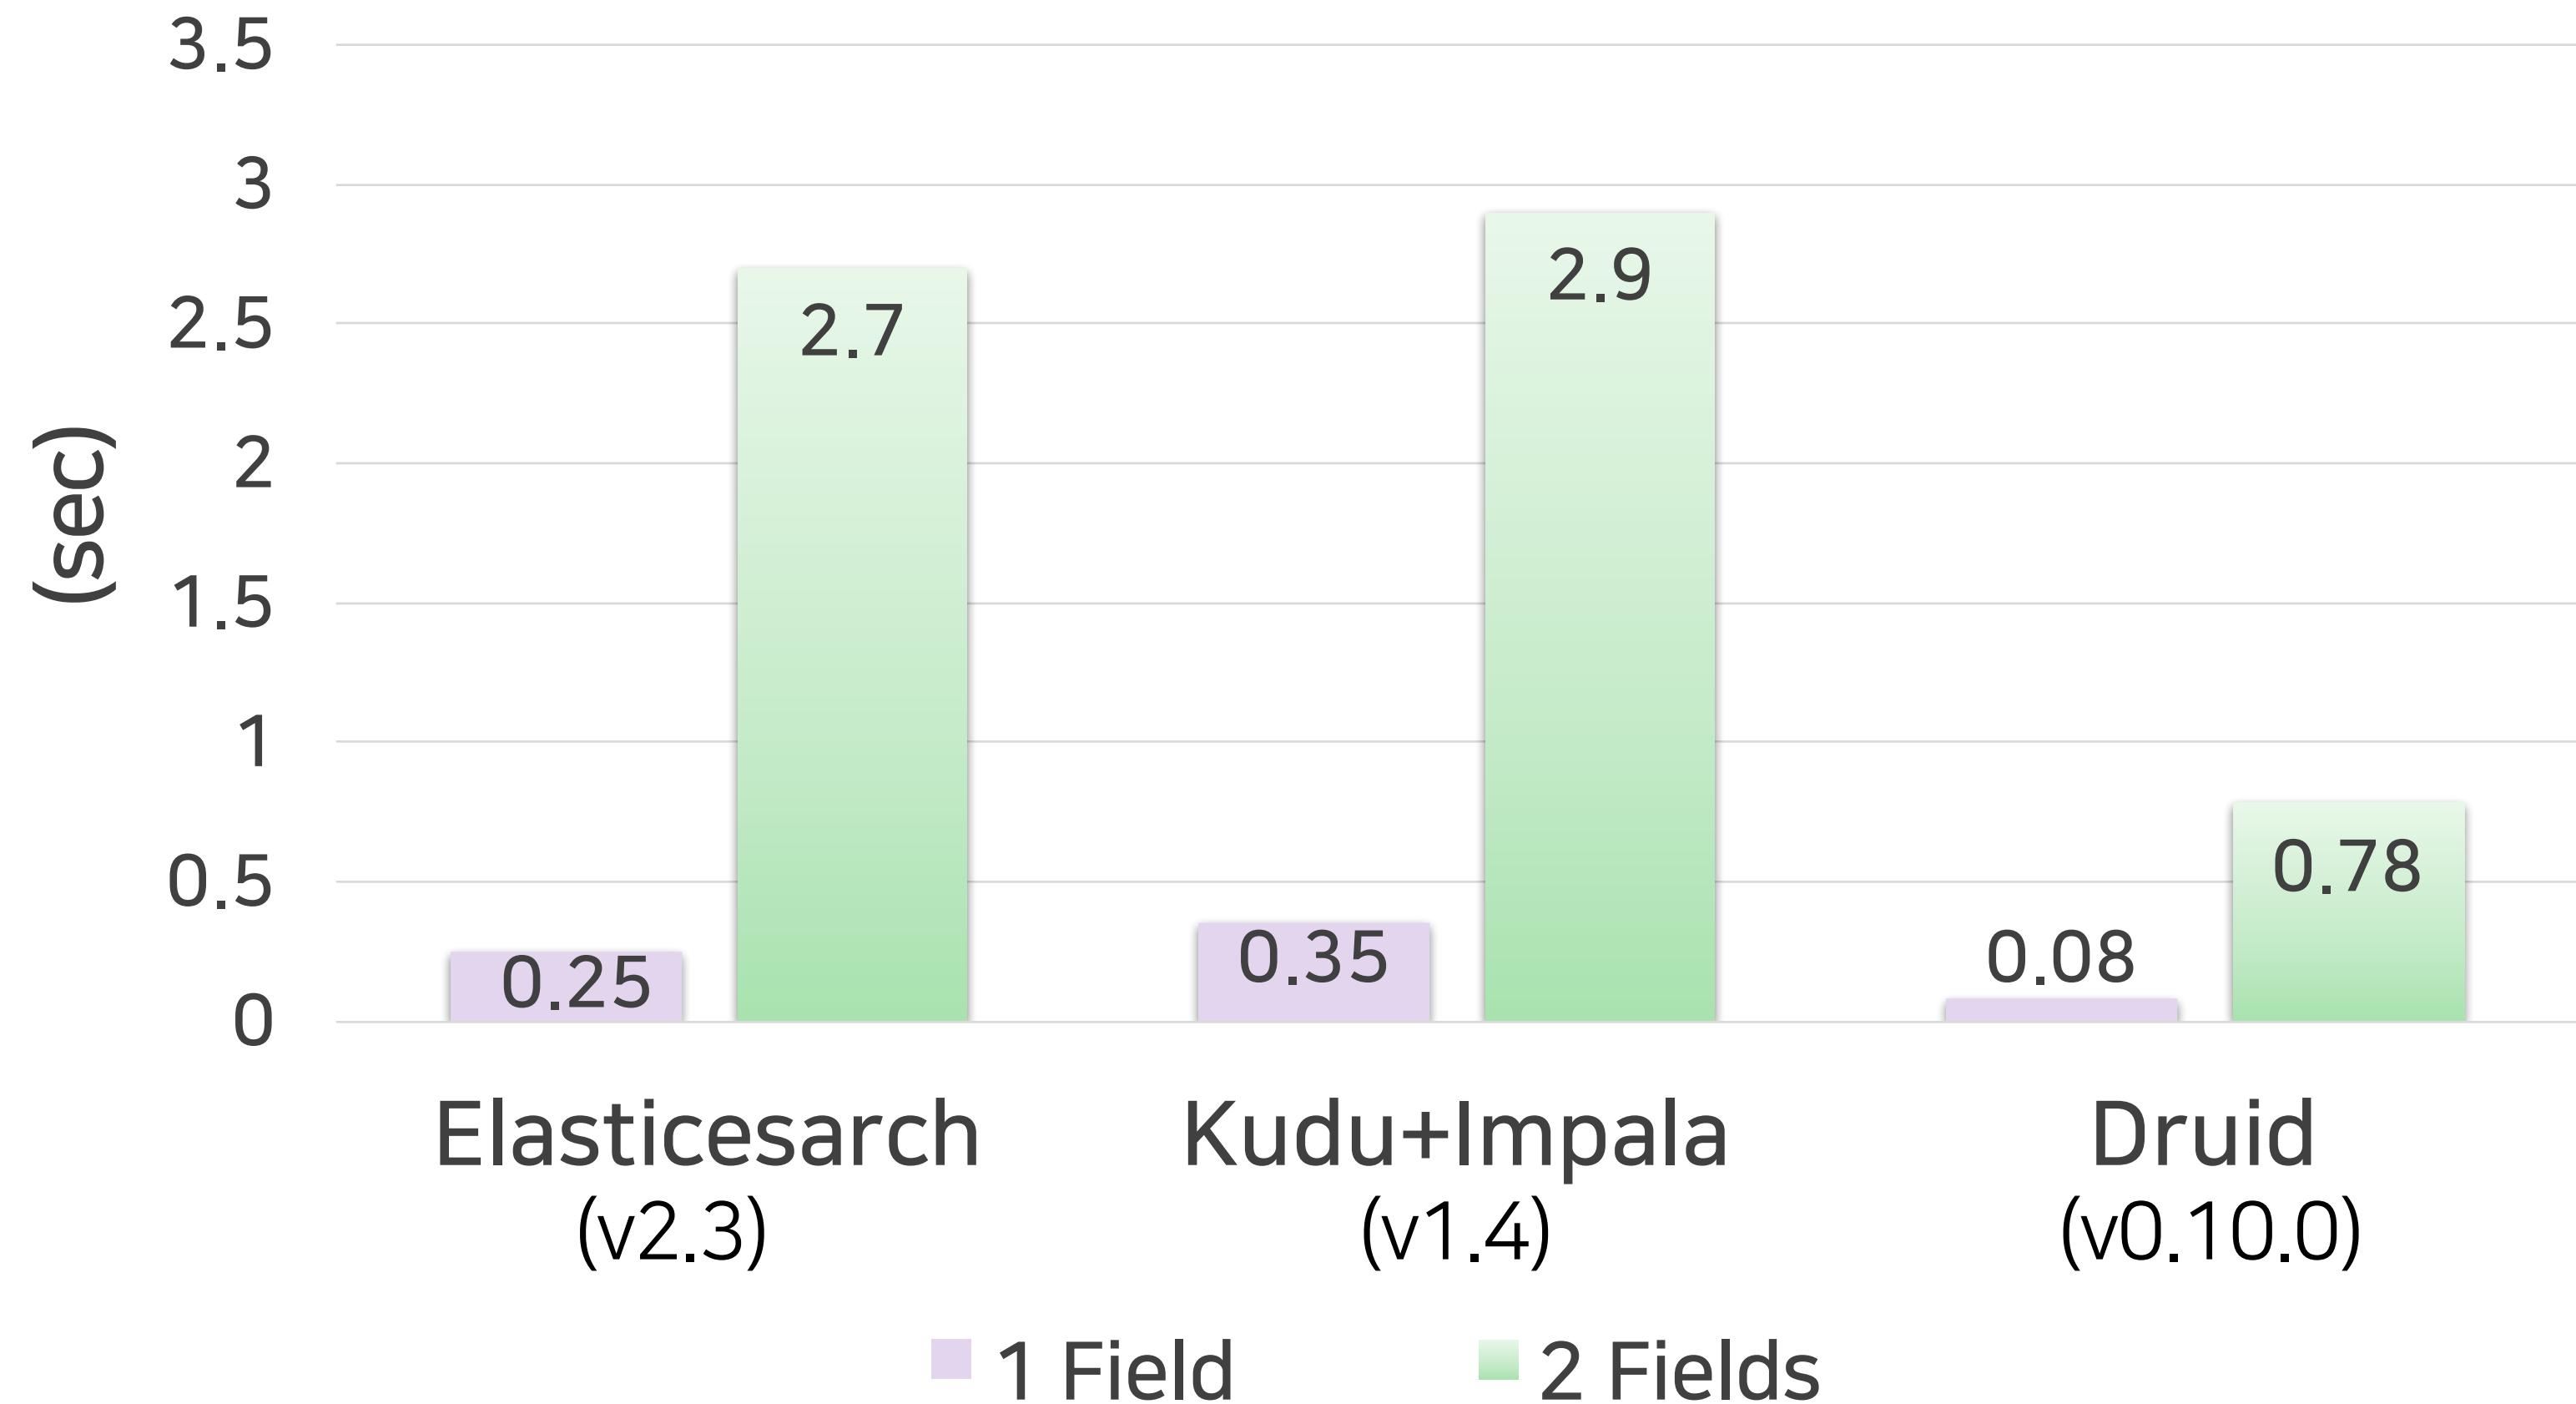
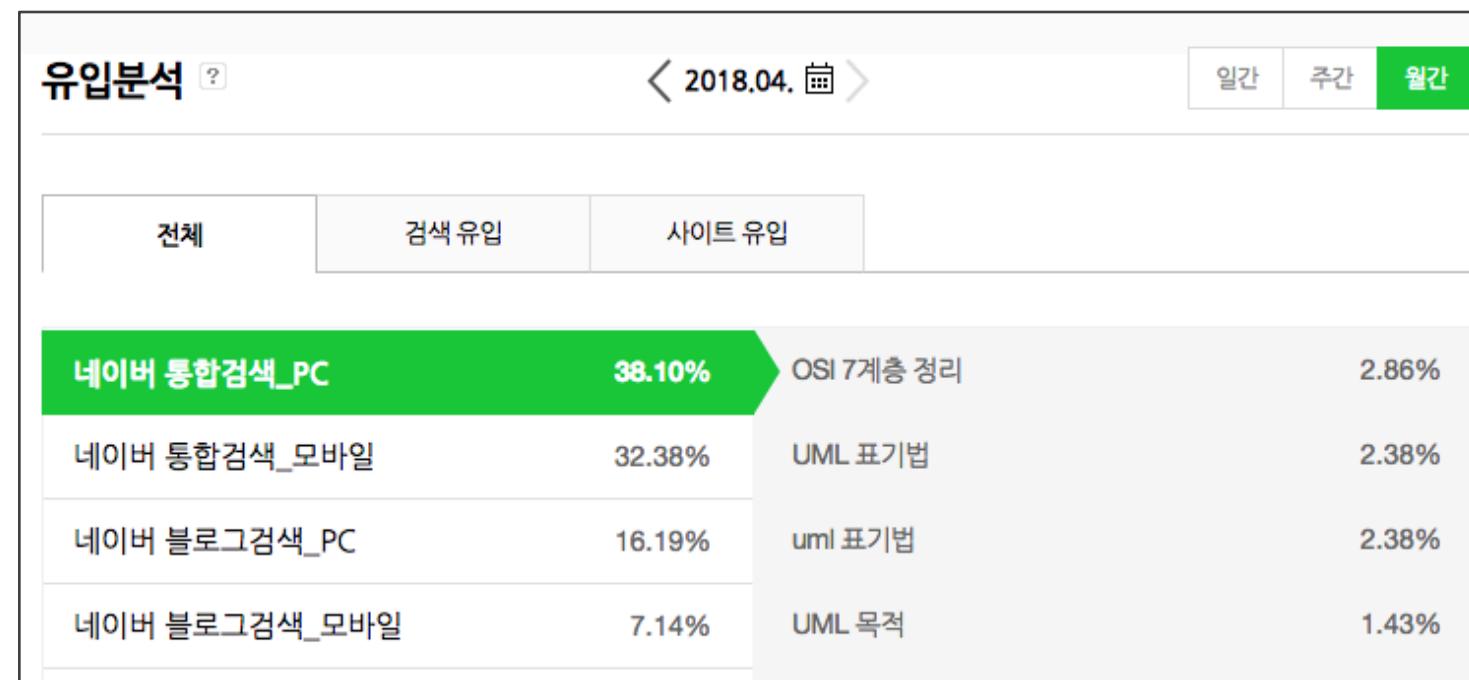
	Random Access	Ranking	Group By	Join
Apache Druid (v0.10.0)	★★★★★☆	★★★★★☆	★★★★★☆	★★★★★☆
Elasticsearch (v2.3)	★★★★★☆	★★★★☆☆	☆☆☆☆☆☆	☆☆☆☆☆☆
Apache Kudu (v1.4) with Impala	★★★★★☆ (pk) ★☆☆☆☆☆(pk외)	☆☆☆☆☆☆	★★★★★☆	★★★★★☆

blueml deview

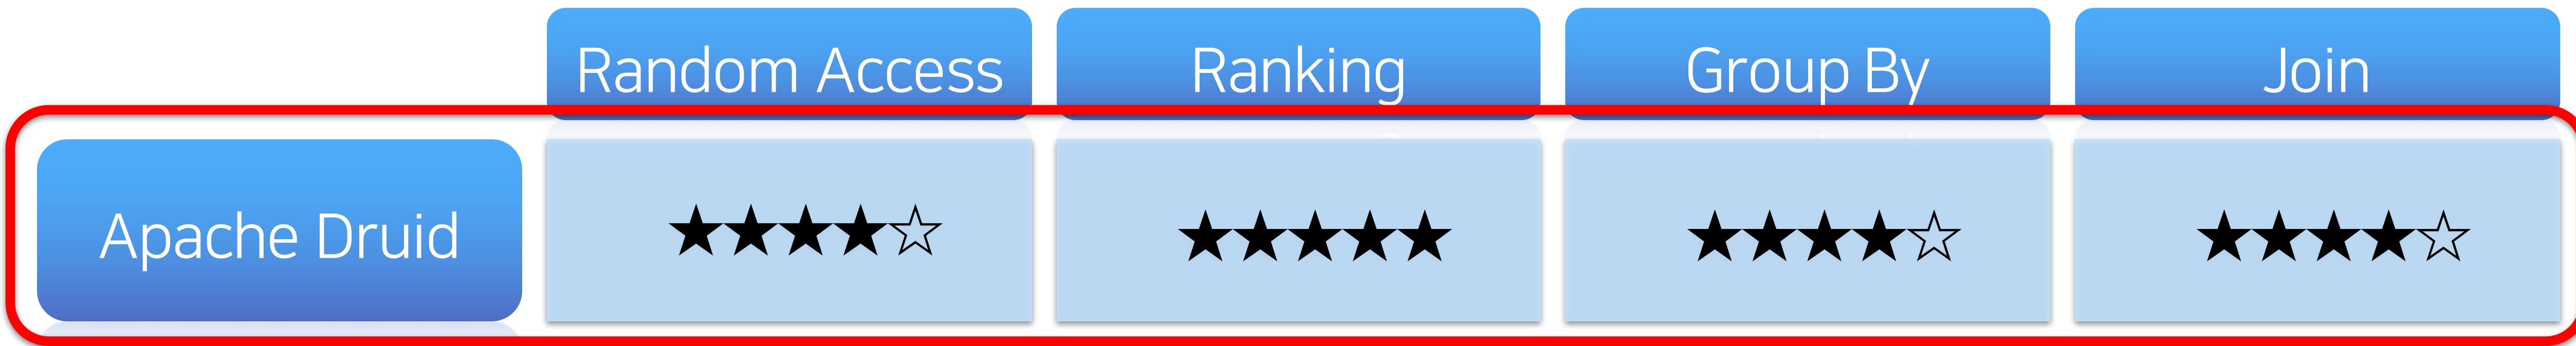
## 2.2 Druid vs. Elasticsearch vs. Kudu

### 응답시간 비교실험 – 다차원 분석

(2.1억건, 동일 병렬환경)



## 2.2 Druid vs. Elasticsearch vs. Kudu



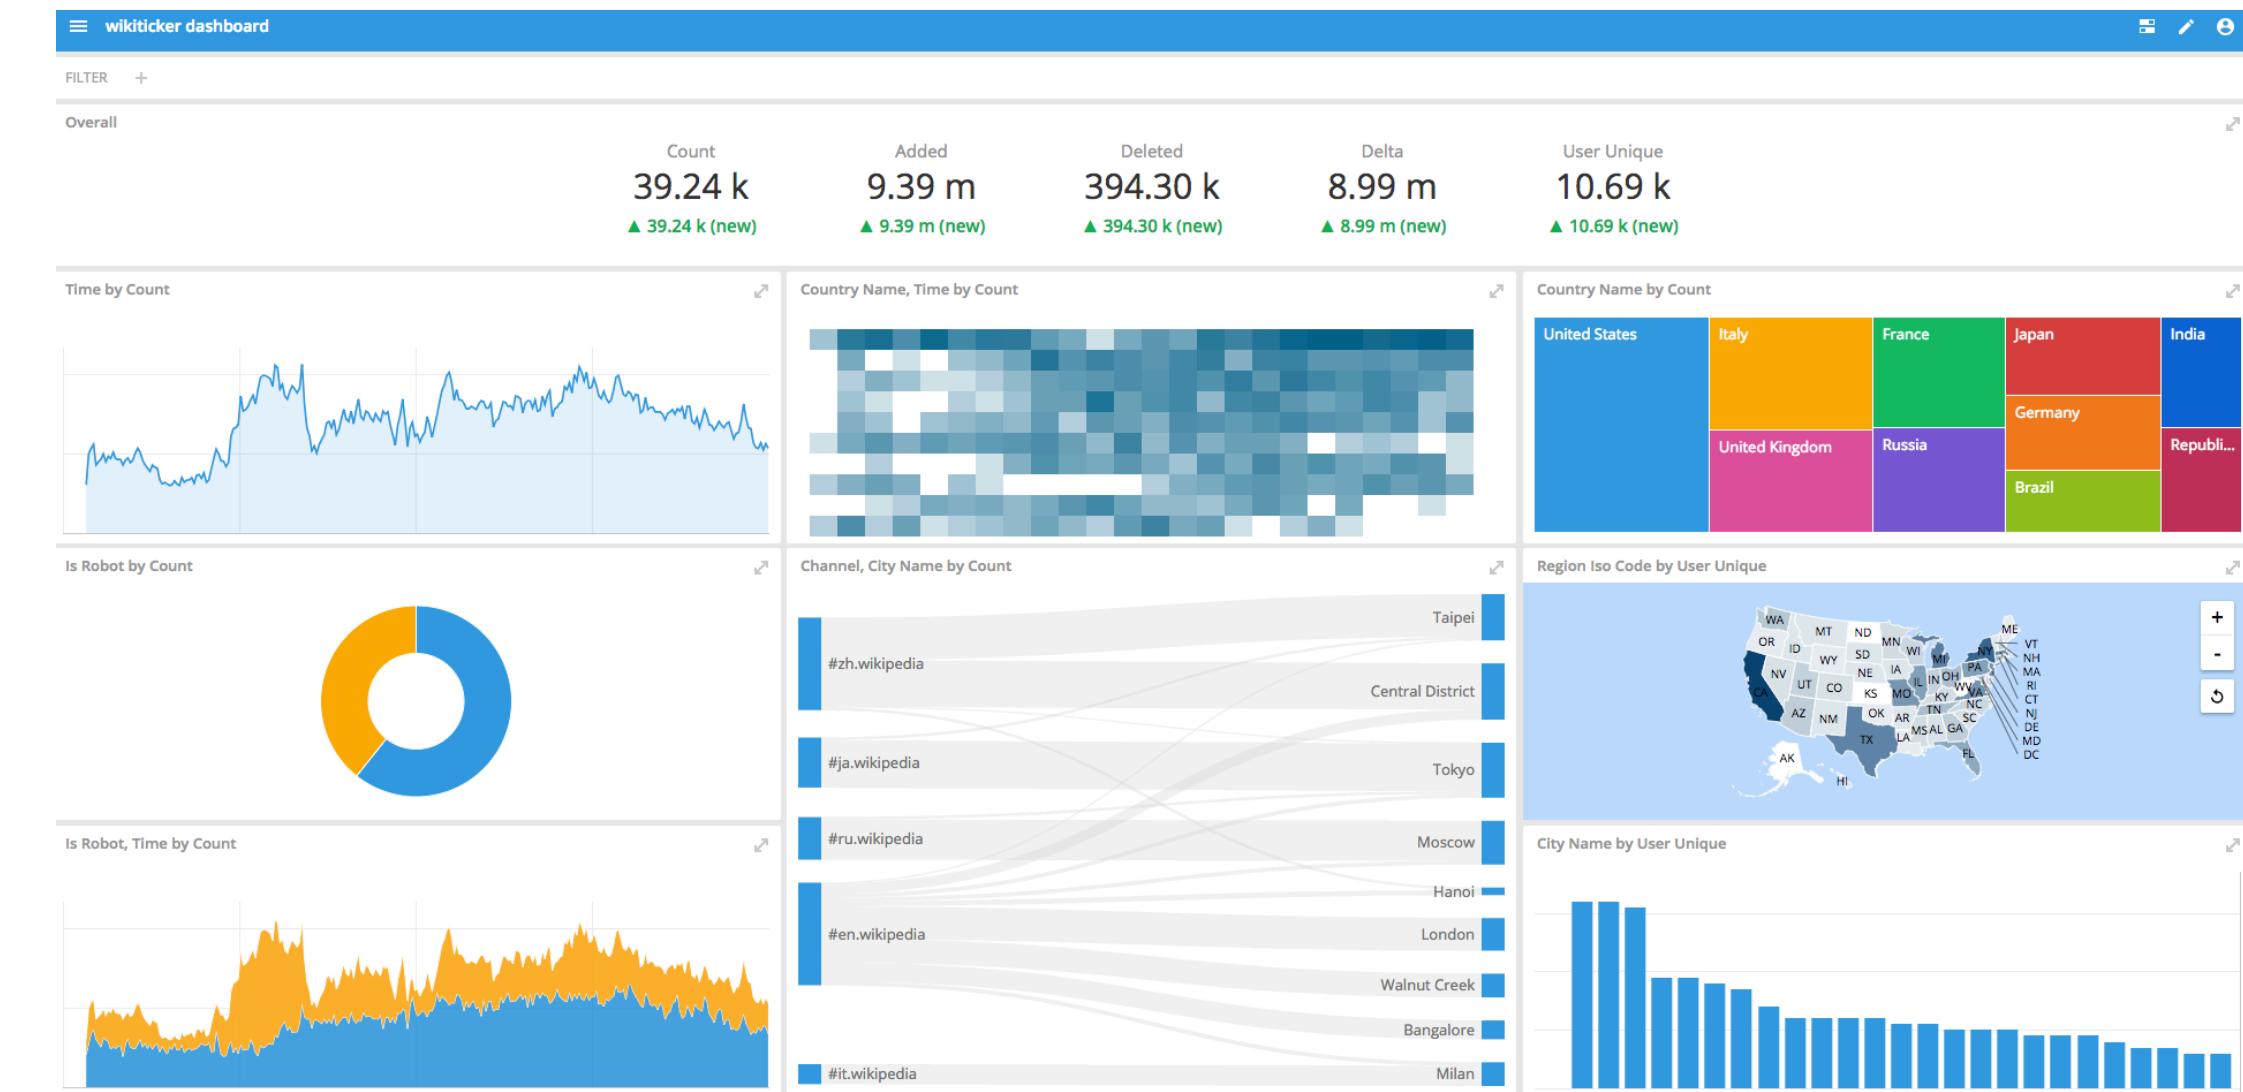
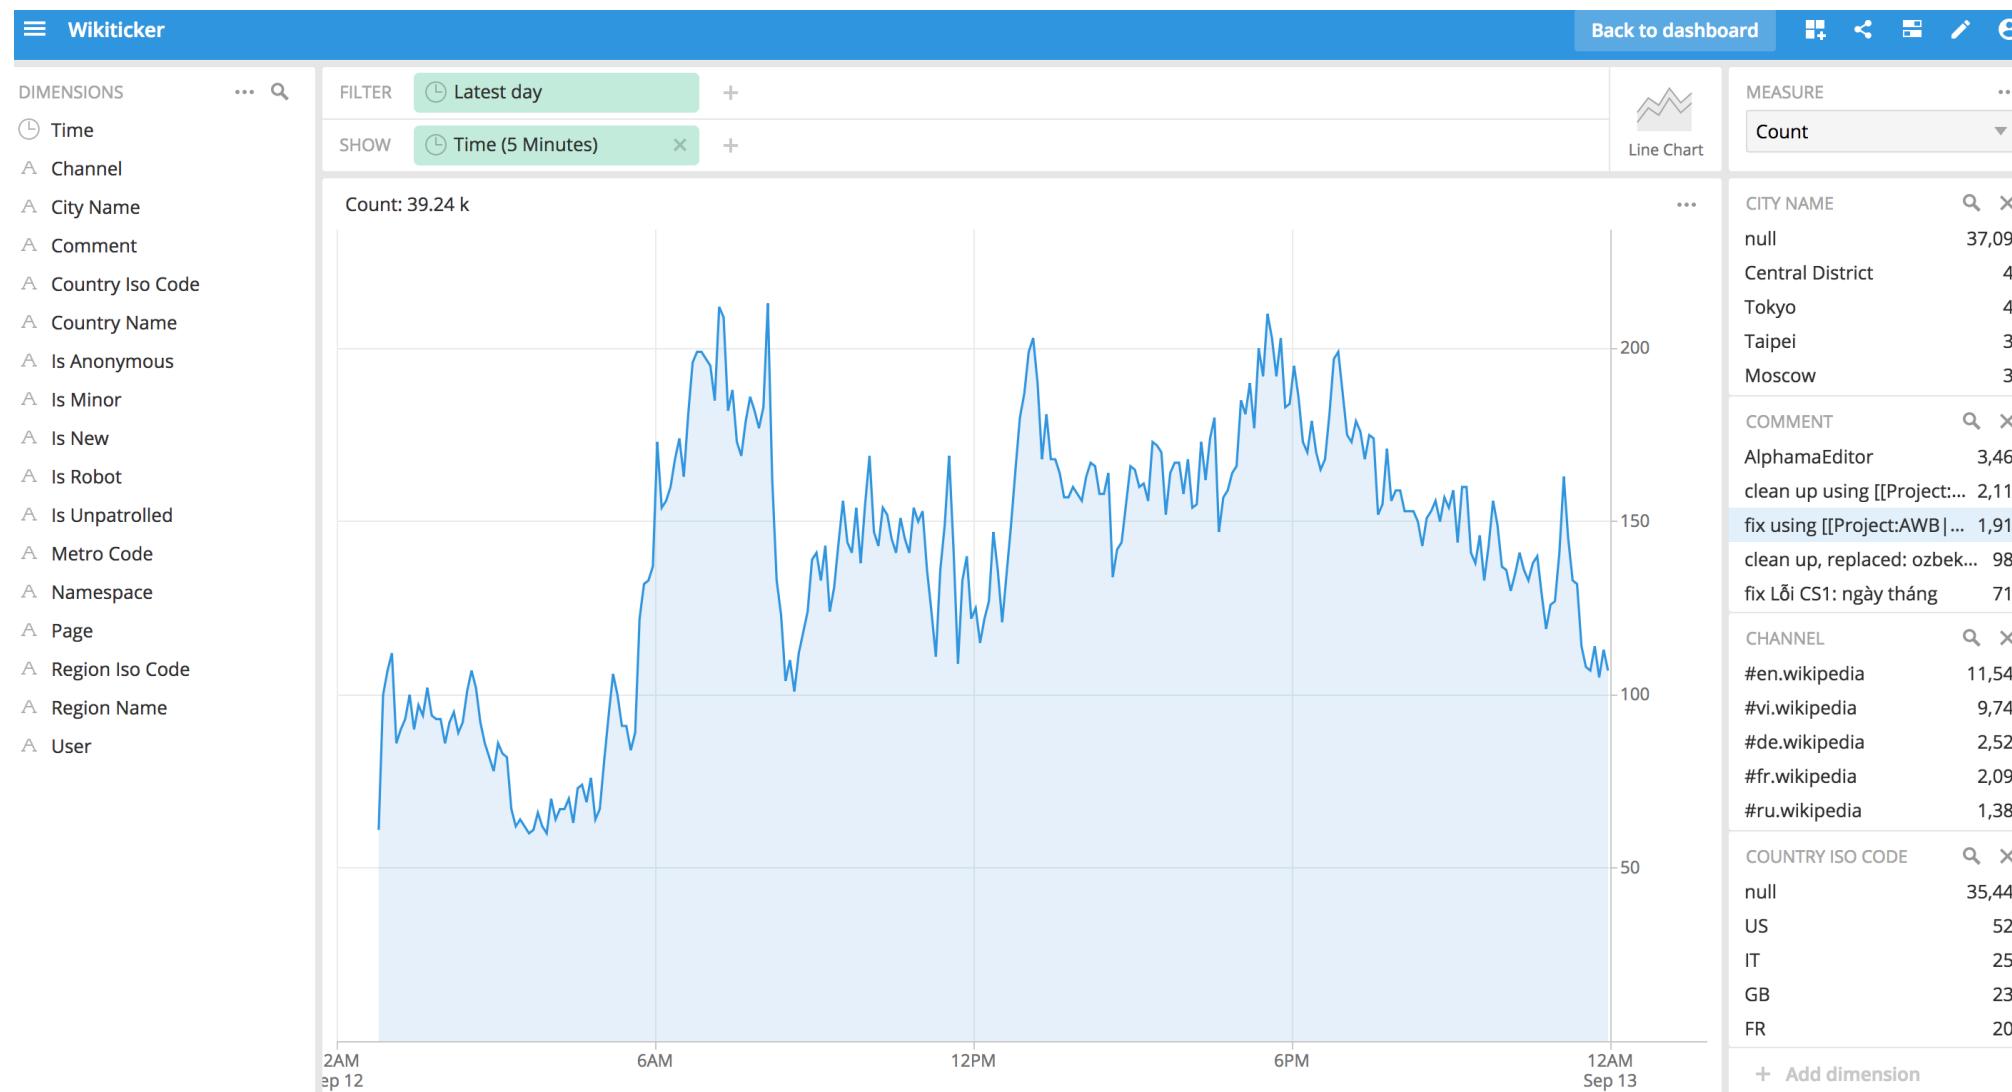
- Secondary Index를 지원하여 PK 외의 key로 조회시에도 빠르게 응답
- Bitmap Index를 사용하여 빠른 필터처리 가능
- 다양한 근사 알고리즘 질의기능 제공
- 복잡한 질의처리 가능 ( Spark Druid Connector )

3.

# Druid로 분석하기 - 기본편

# 3.1 Imply-UI

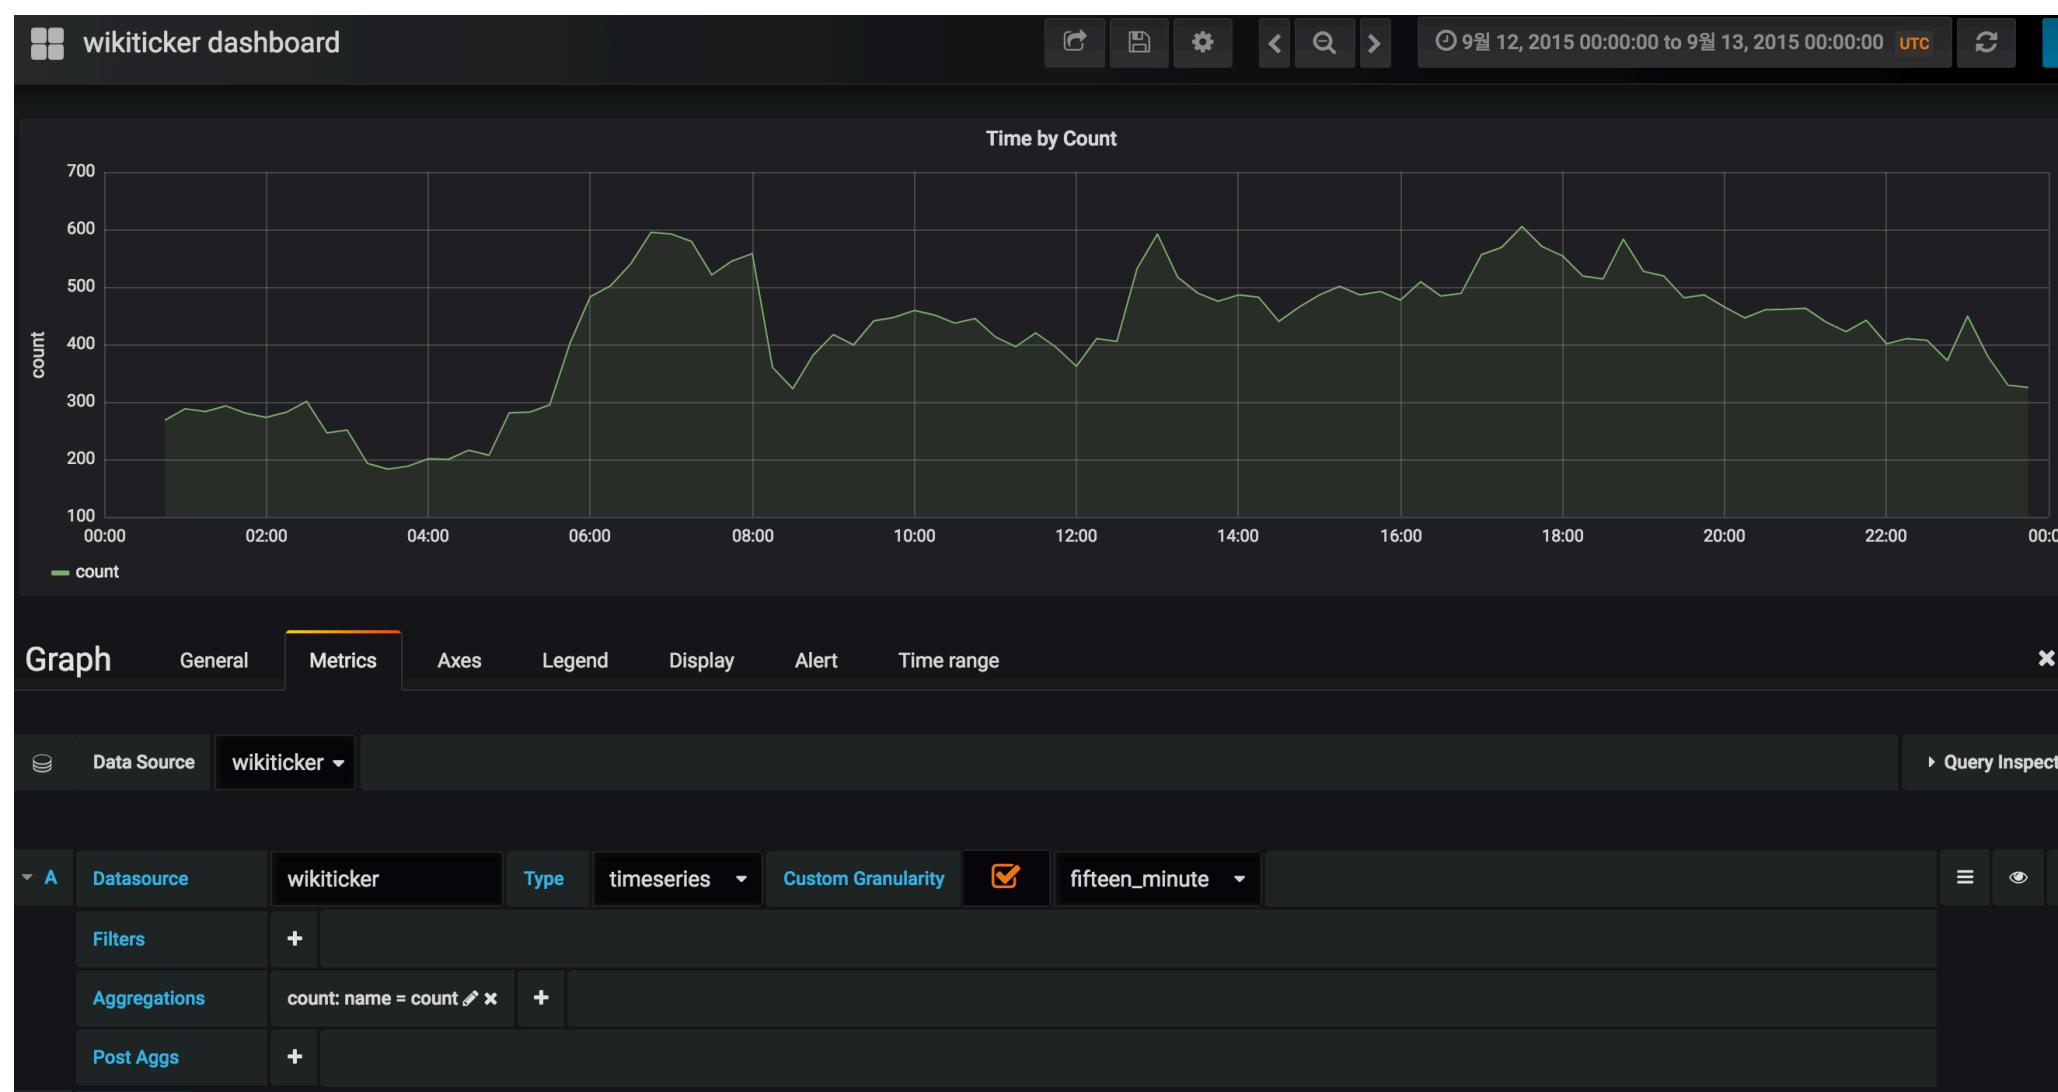
<https://docs.imply.io/on-prem/visualize/>



- Imply에서 개발한 Visualize Tool
- Druid 버전에 따라 빠르게 업그레이드
- Drag & Drop으로 쉽게 질의결과 표현 가능
- 값의 분포, 이전기간과의 비교를 쉽게 해볼 수 있음

## 3.2 Grafana

<https://grafana.com/>

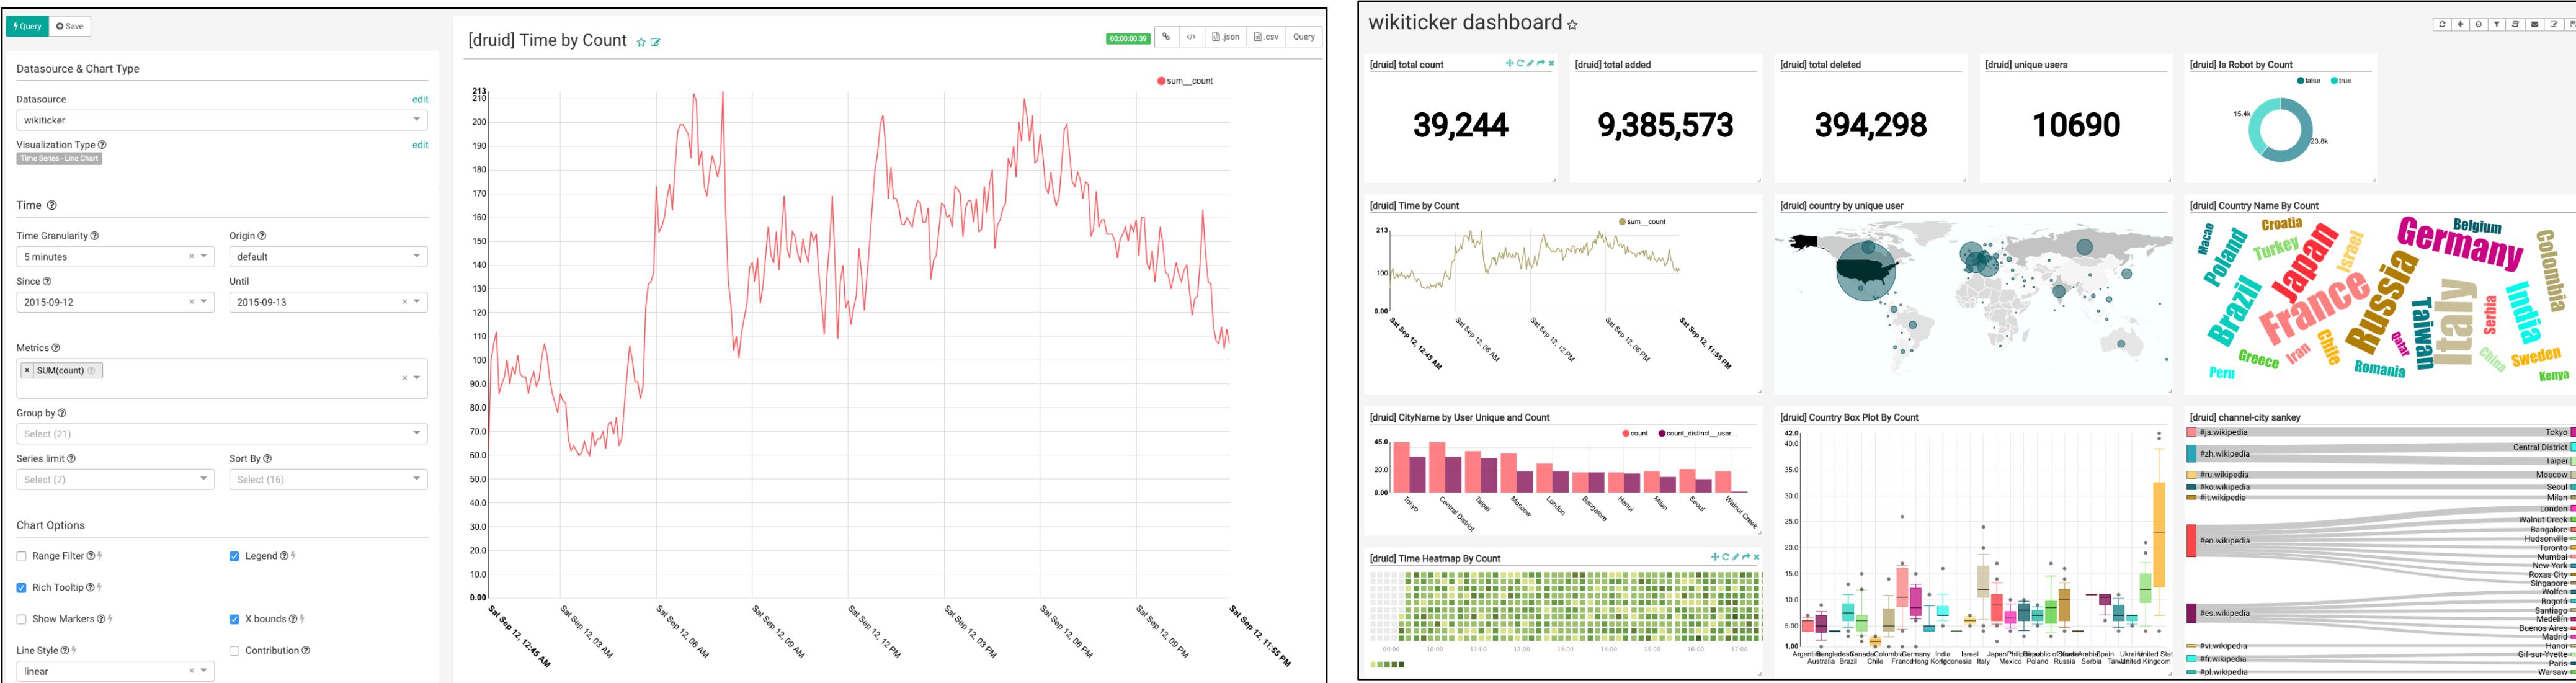


- 다양한 Datasource Visualize Tool
- Druid Plugin을 설치하여 Druid 연동가능

- Plugin을 통해 다양한 그래프 쉽게 표현 가능
- Druid의 DSL 질의방식에 대한 이해가 필요

# 3.3 Superset

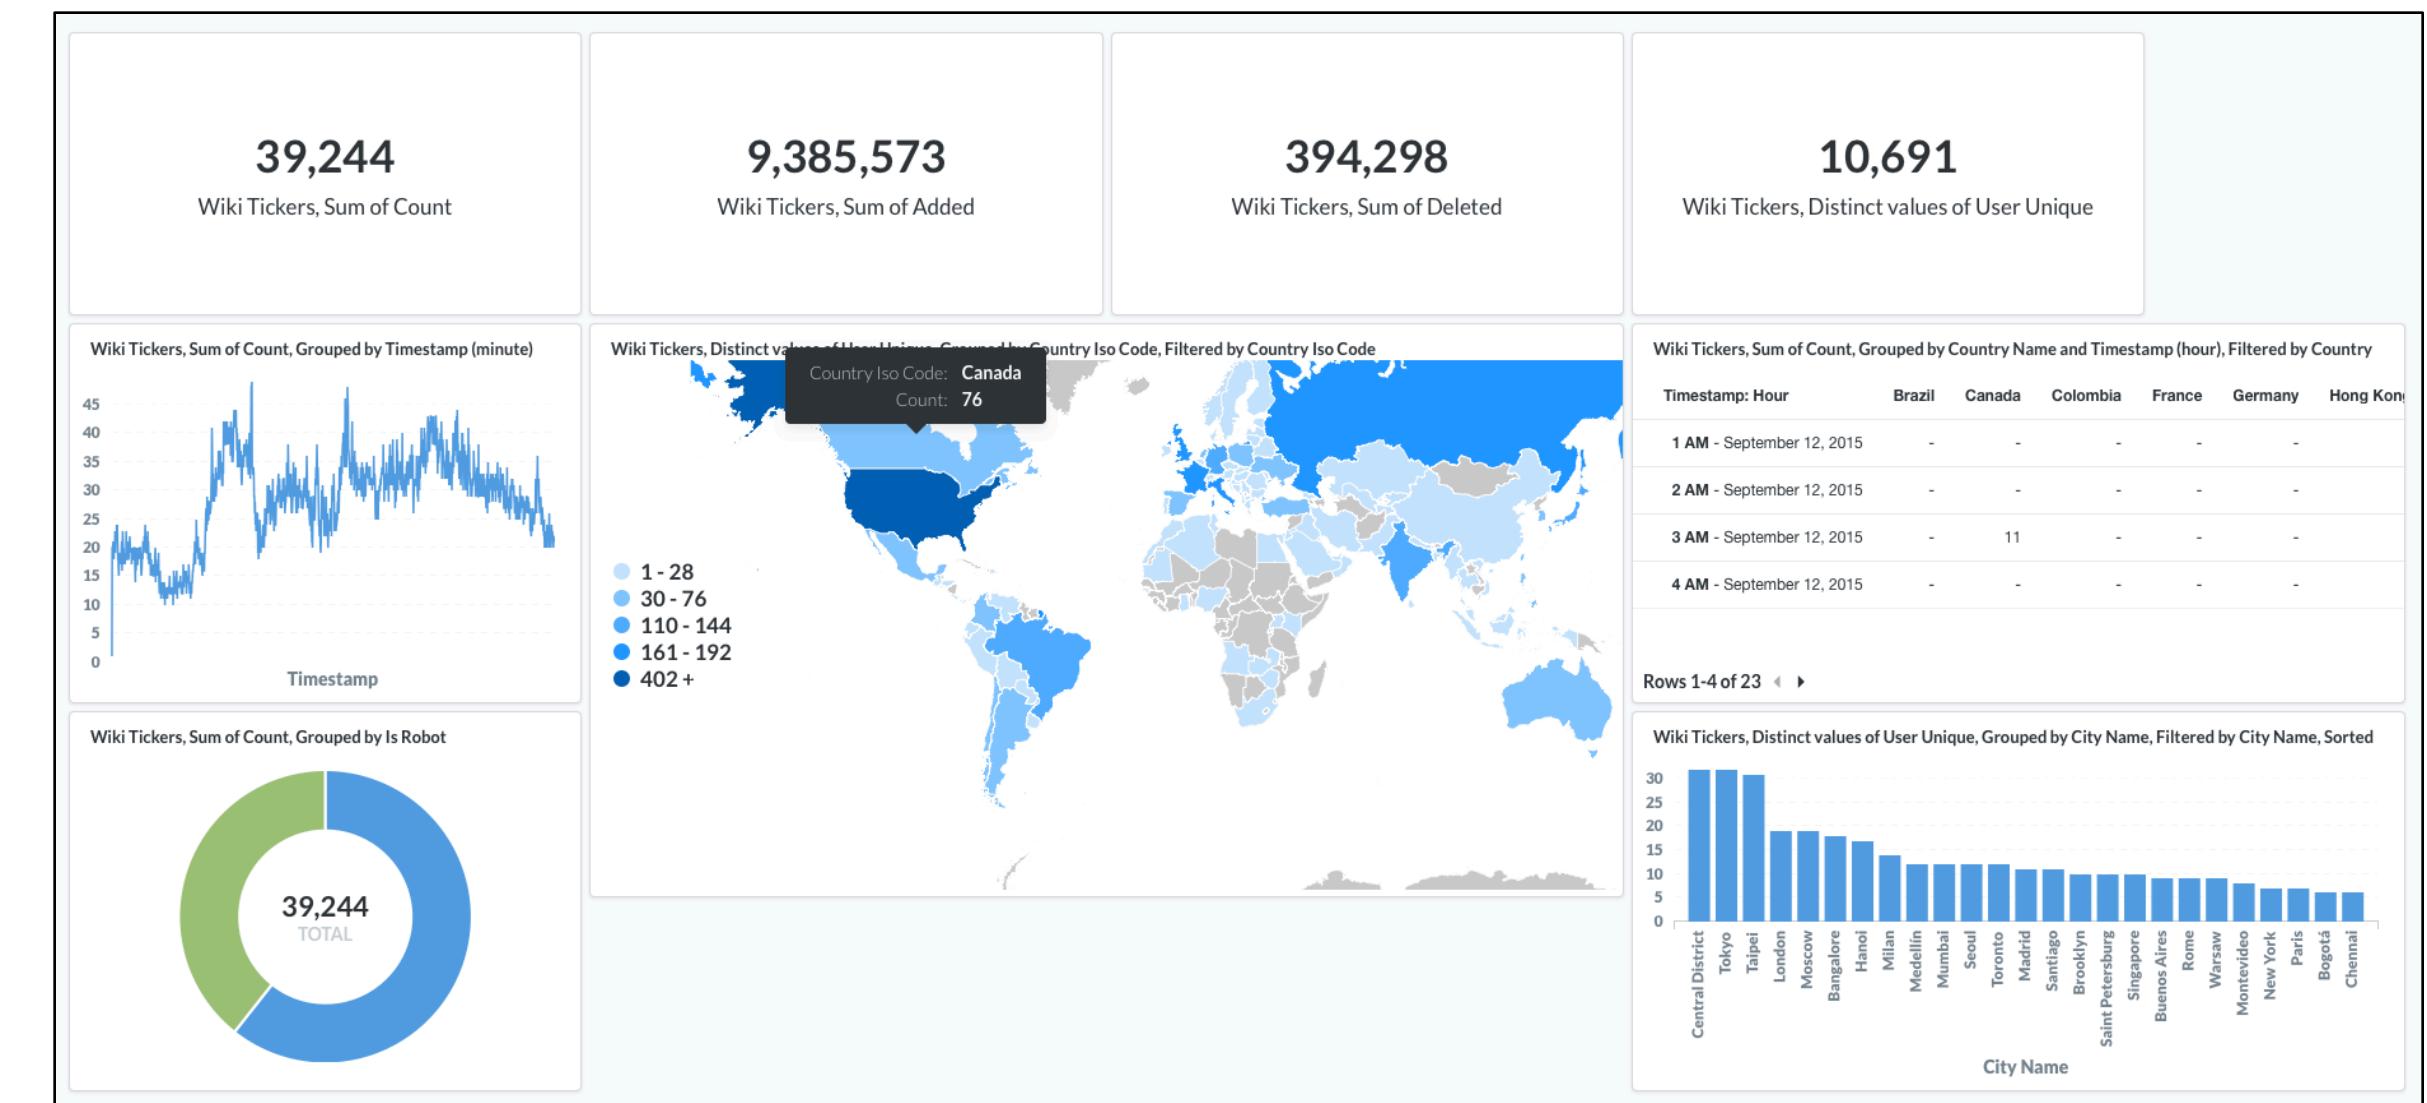
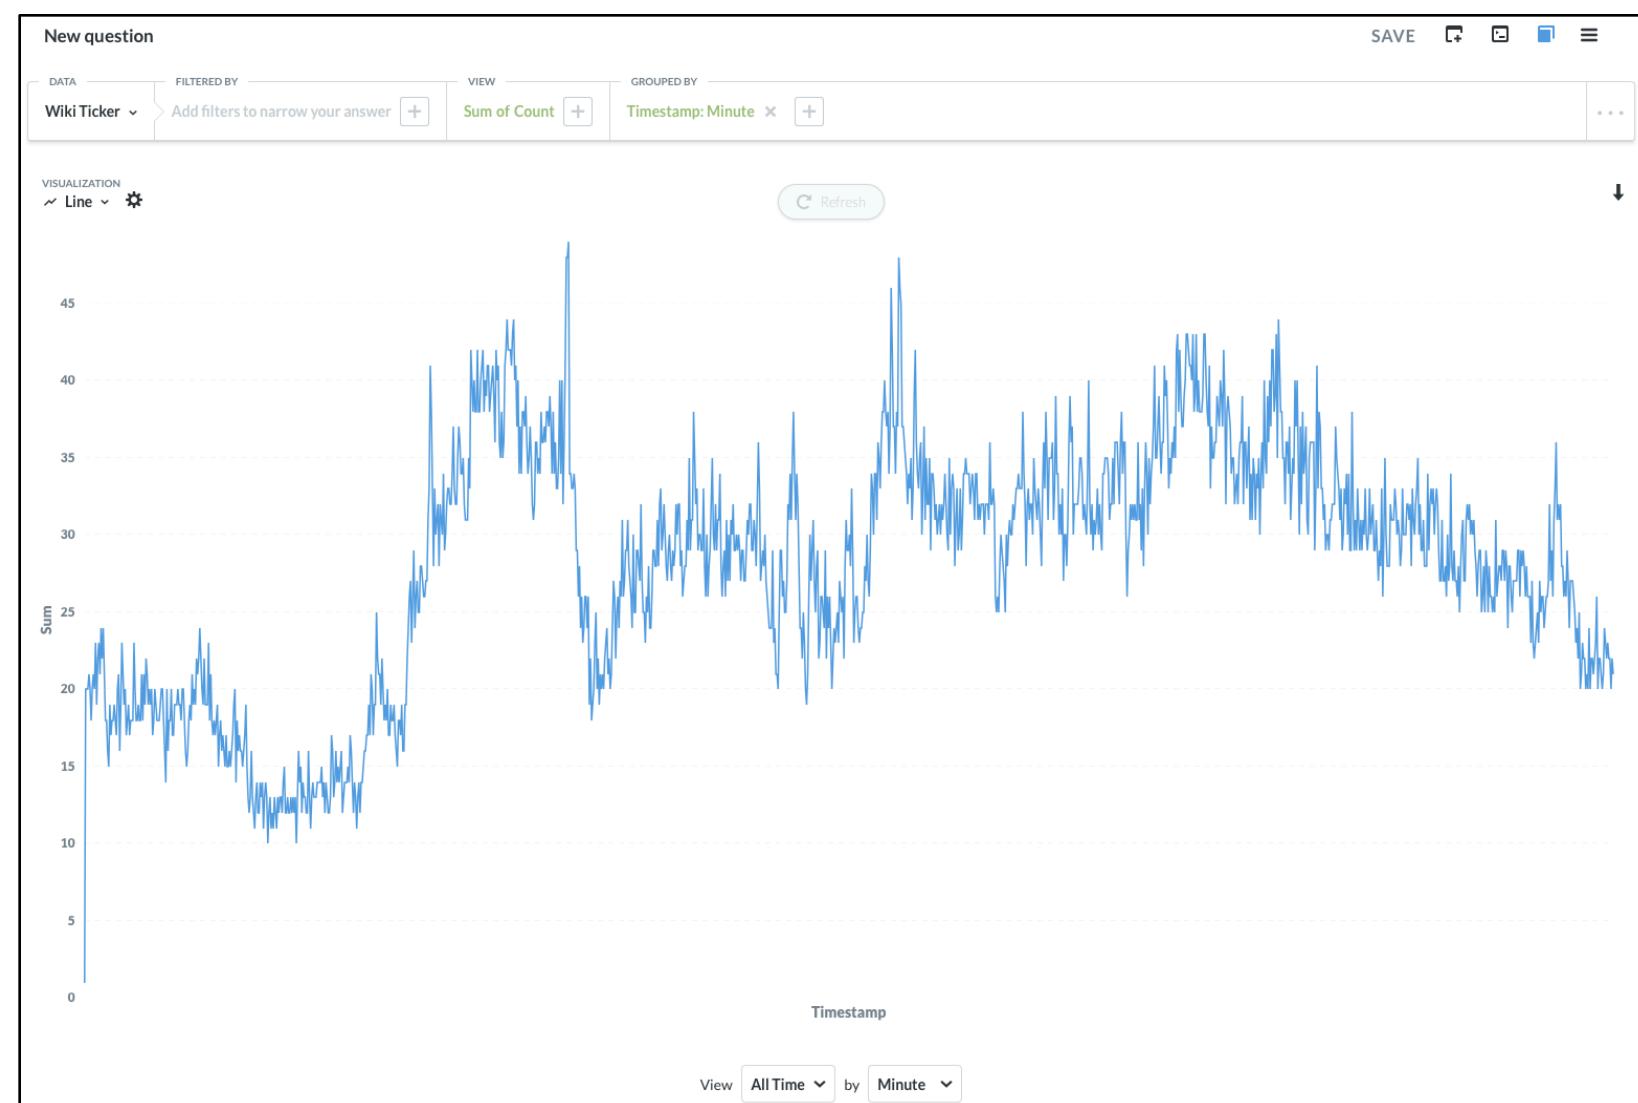
<https://superset.incubator.apache.org/druid.html>



- 기본 제공 그래프 종류가 가장 많음
- 세세한 권한관리 가능
- 표현 자유도가 높음
- 설정할 수 있는게 많아 복잡/어렵게 느껴짐

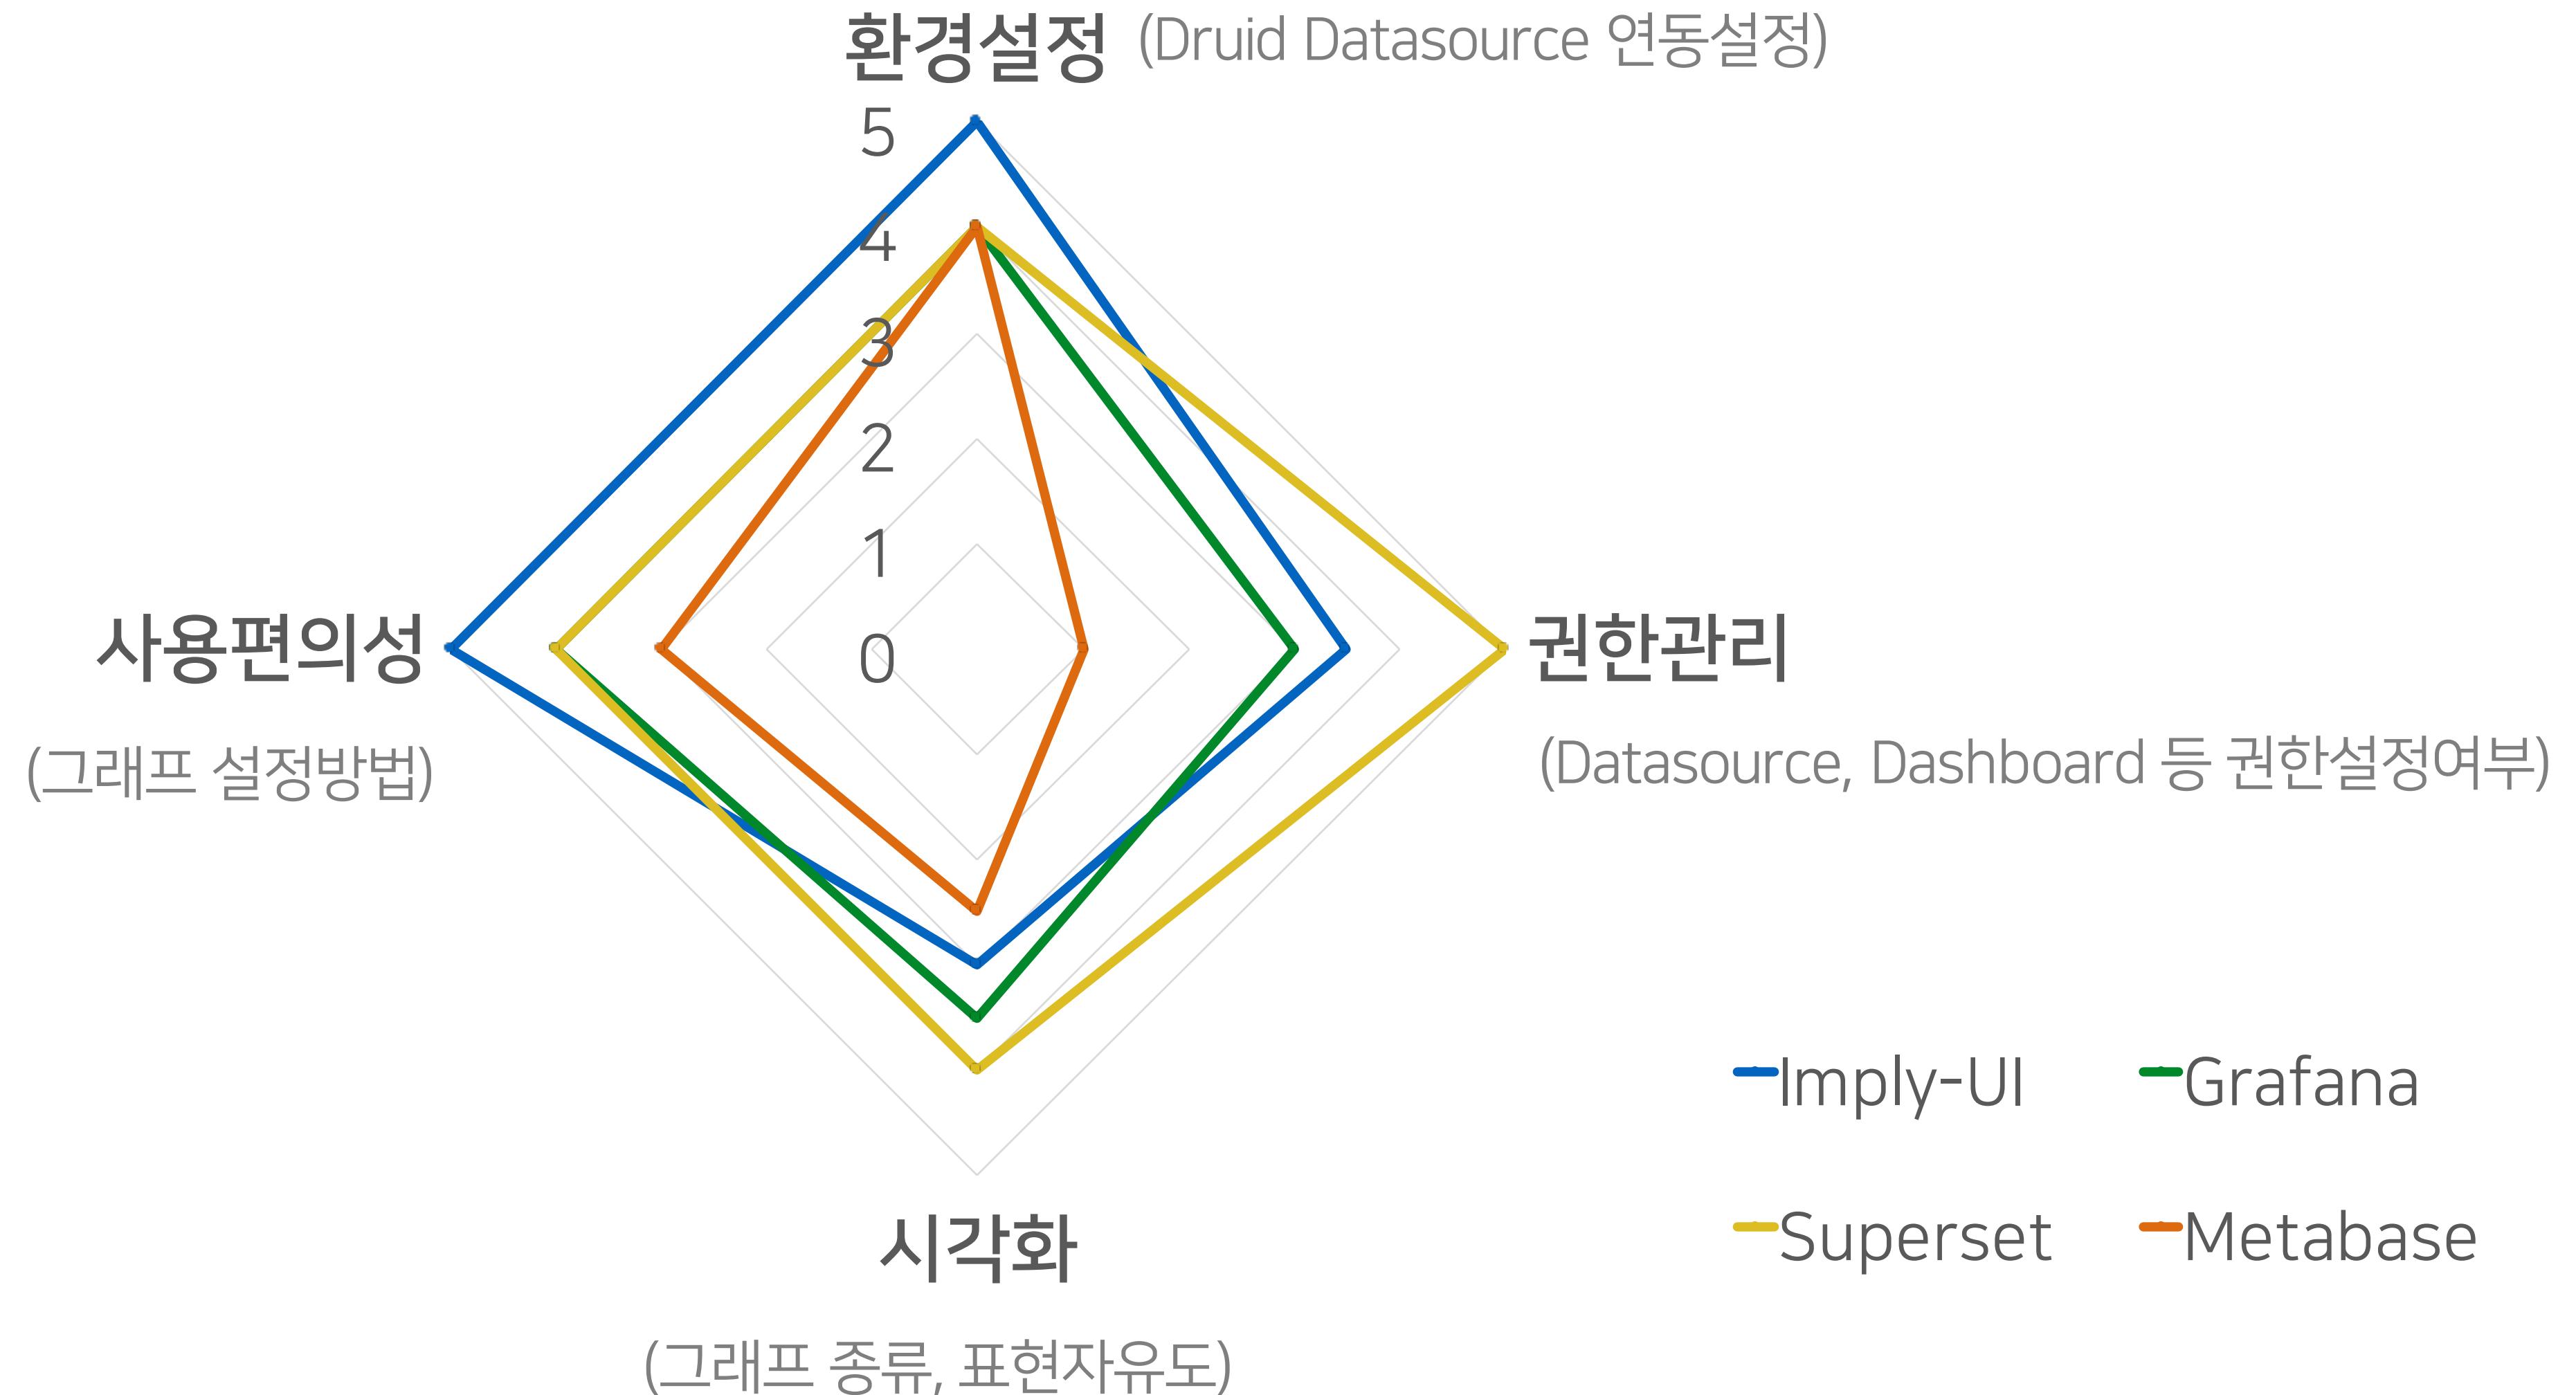
# 3.4 Metabase

<https://www.metabase.com/>



- Application 방식
- 깔끔한 UI
- 그래프의 세부적인 설정 불가

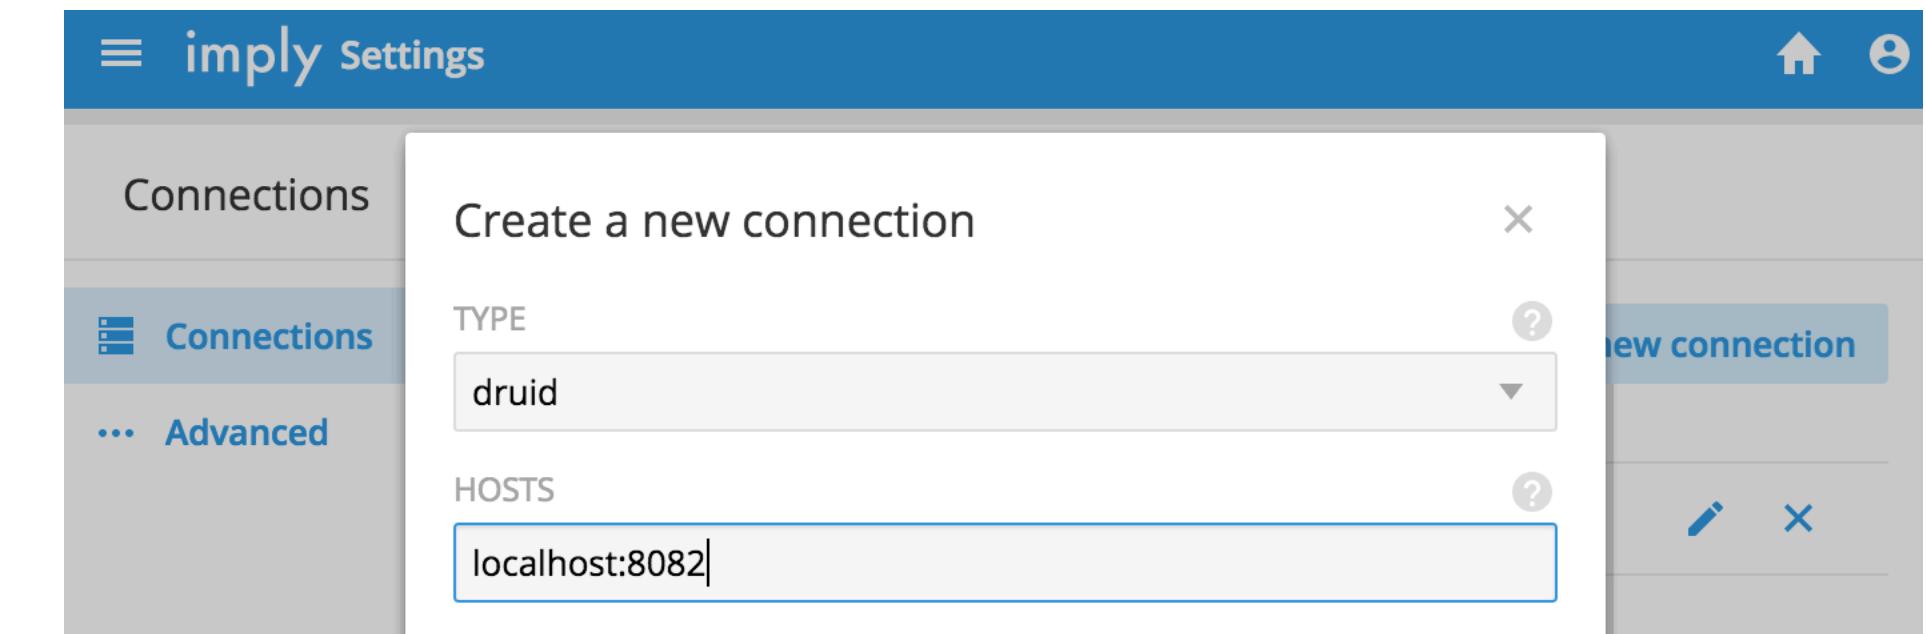
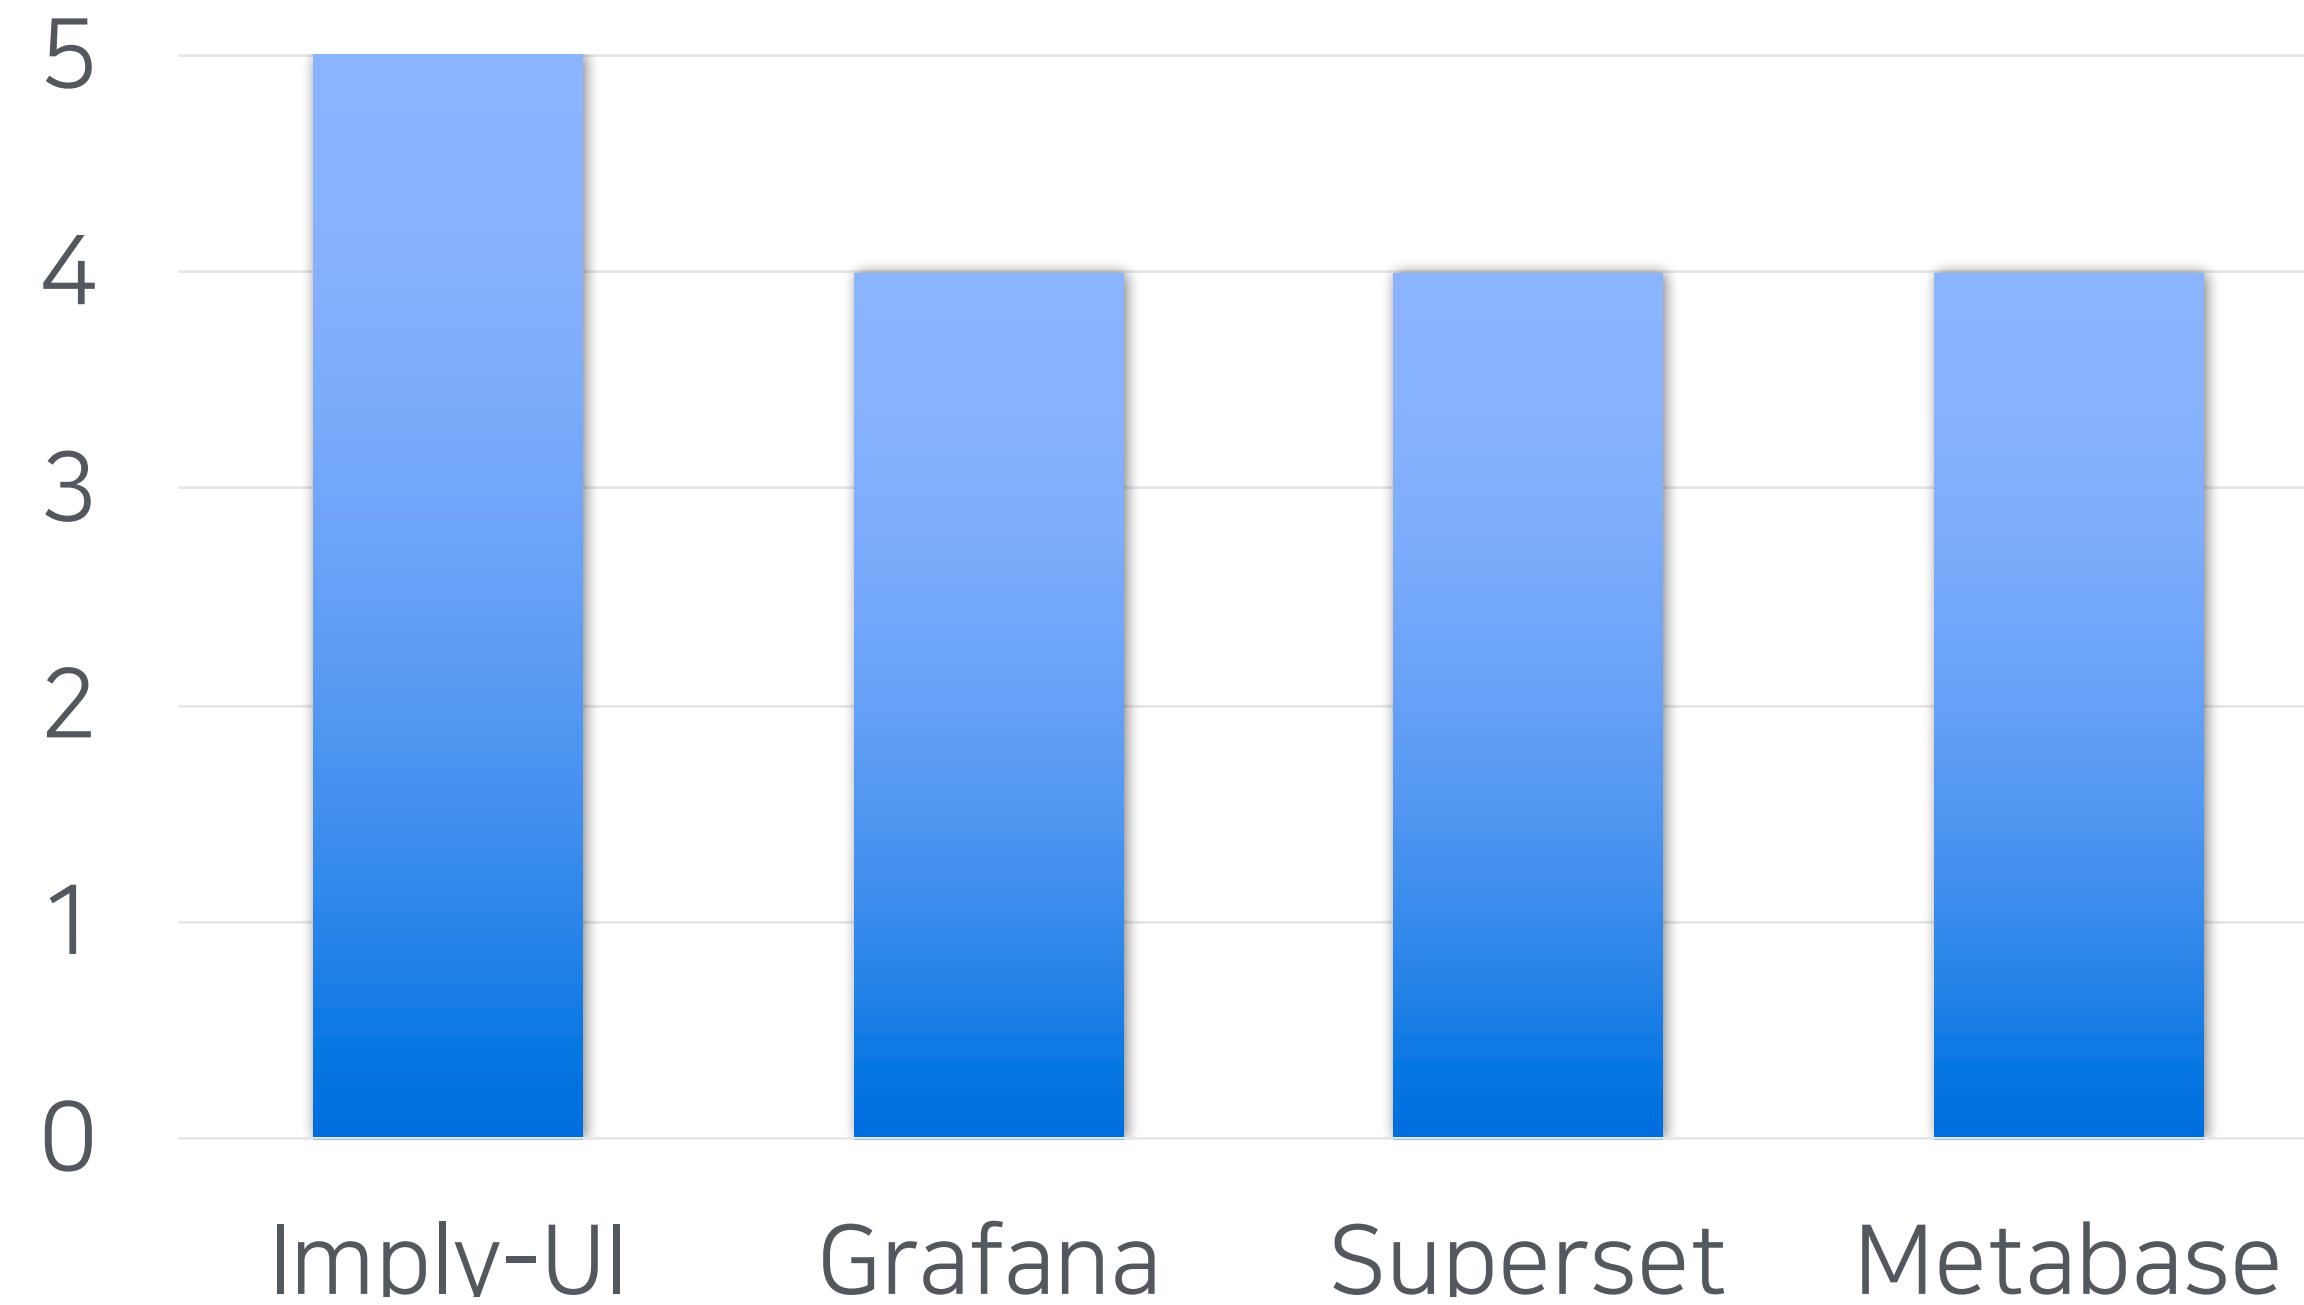
## 3.5 개인적인 총평



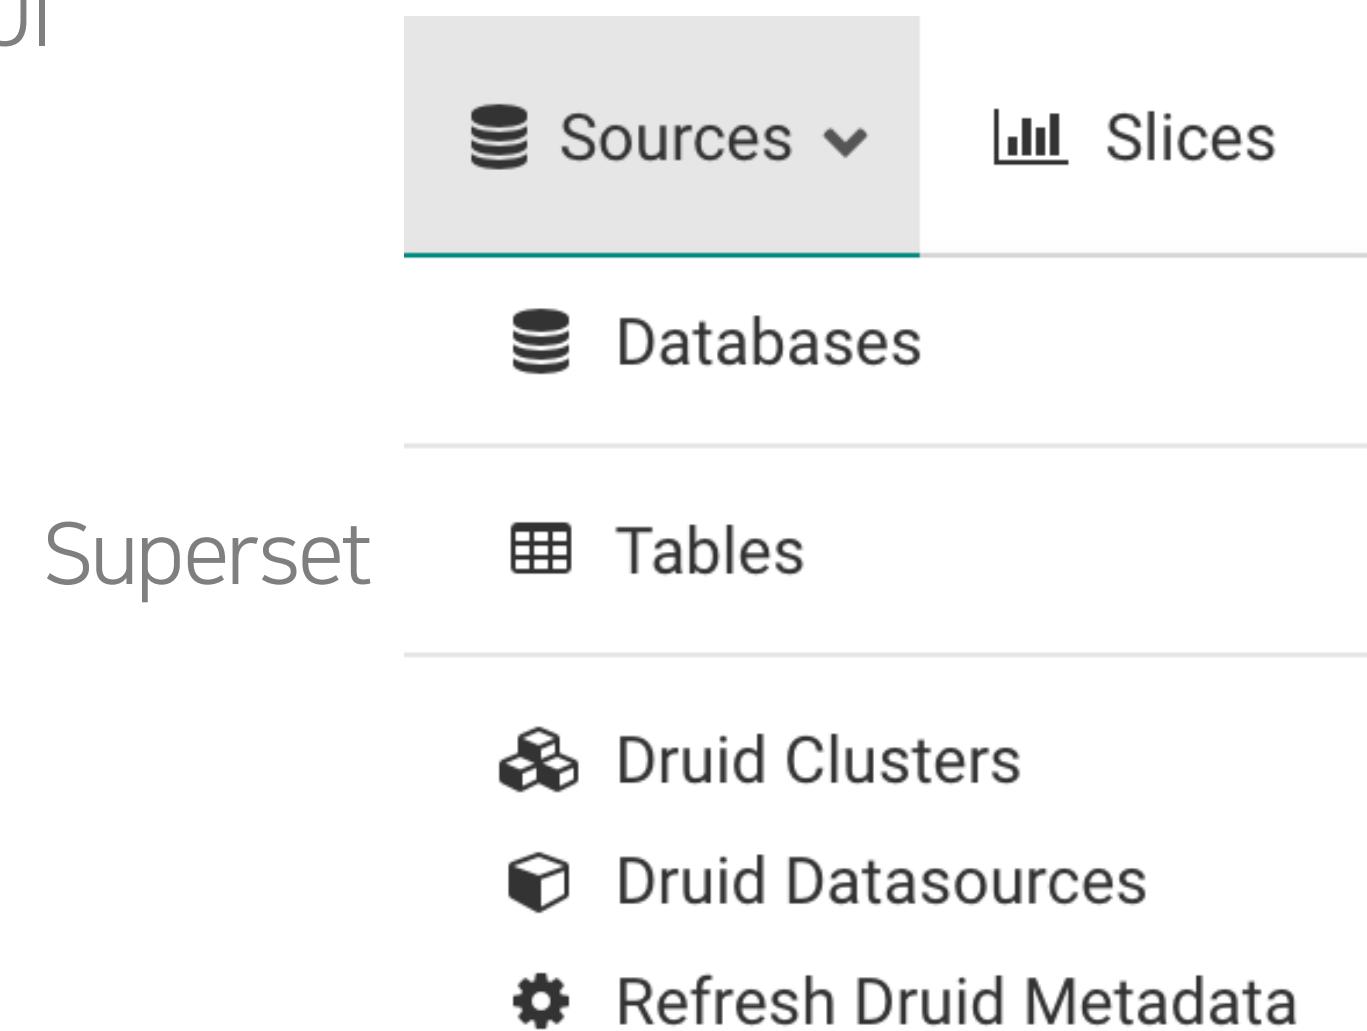
# 3.5 개인적인 총평

## 환경설정

(Druid Datasource 연동)



Imply UI

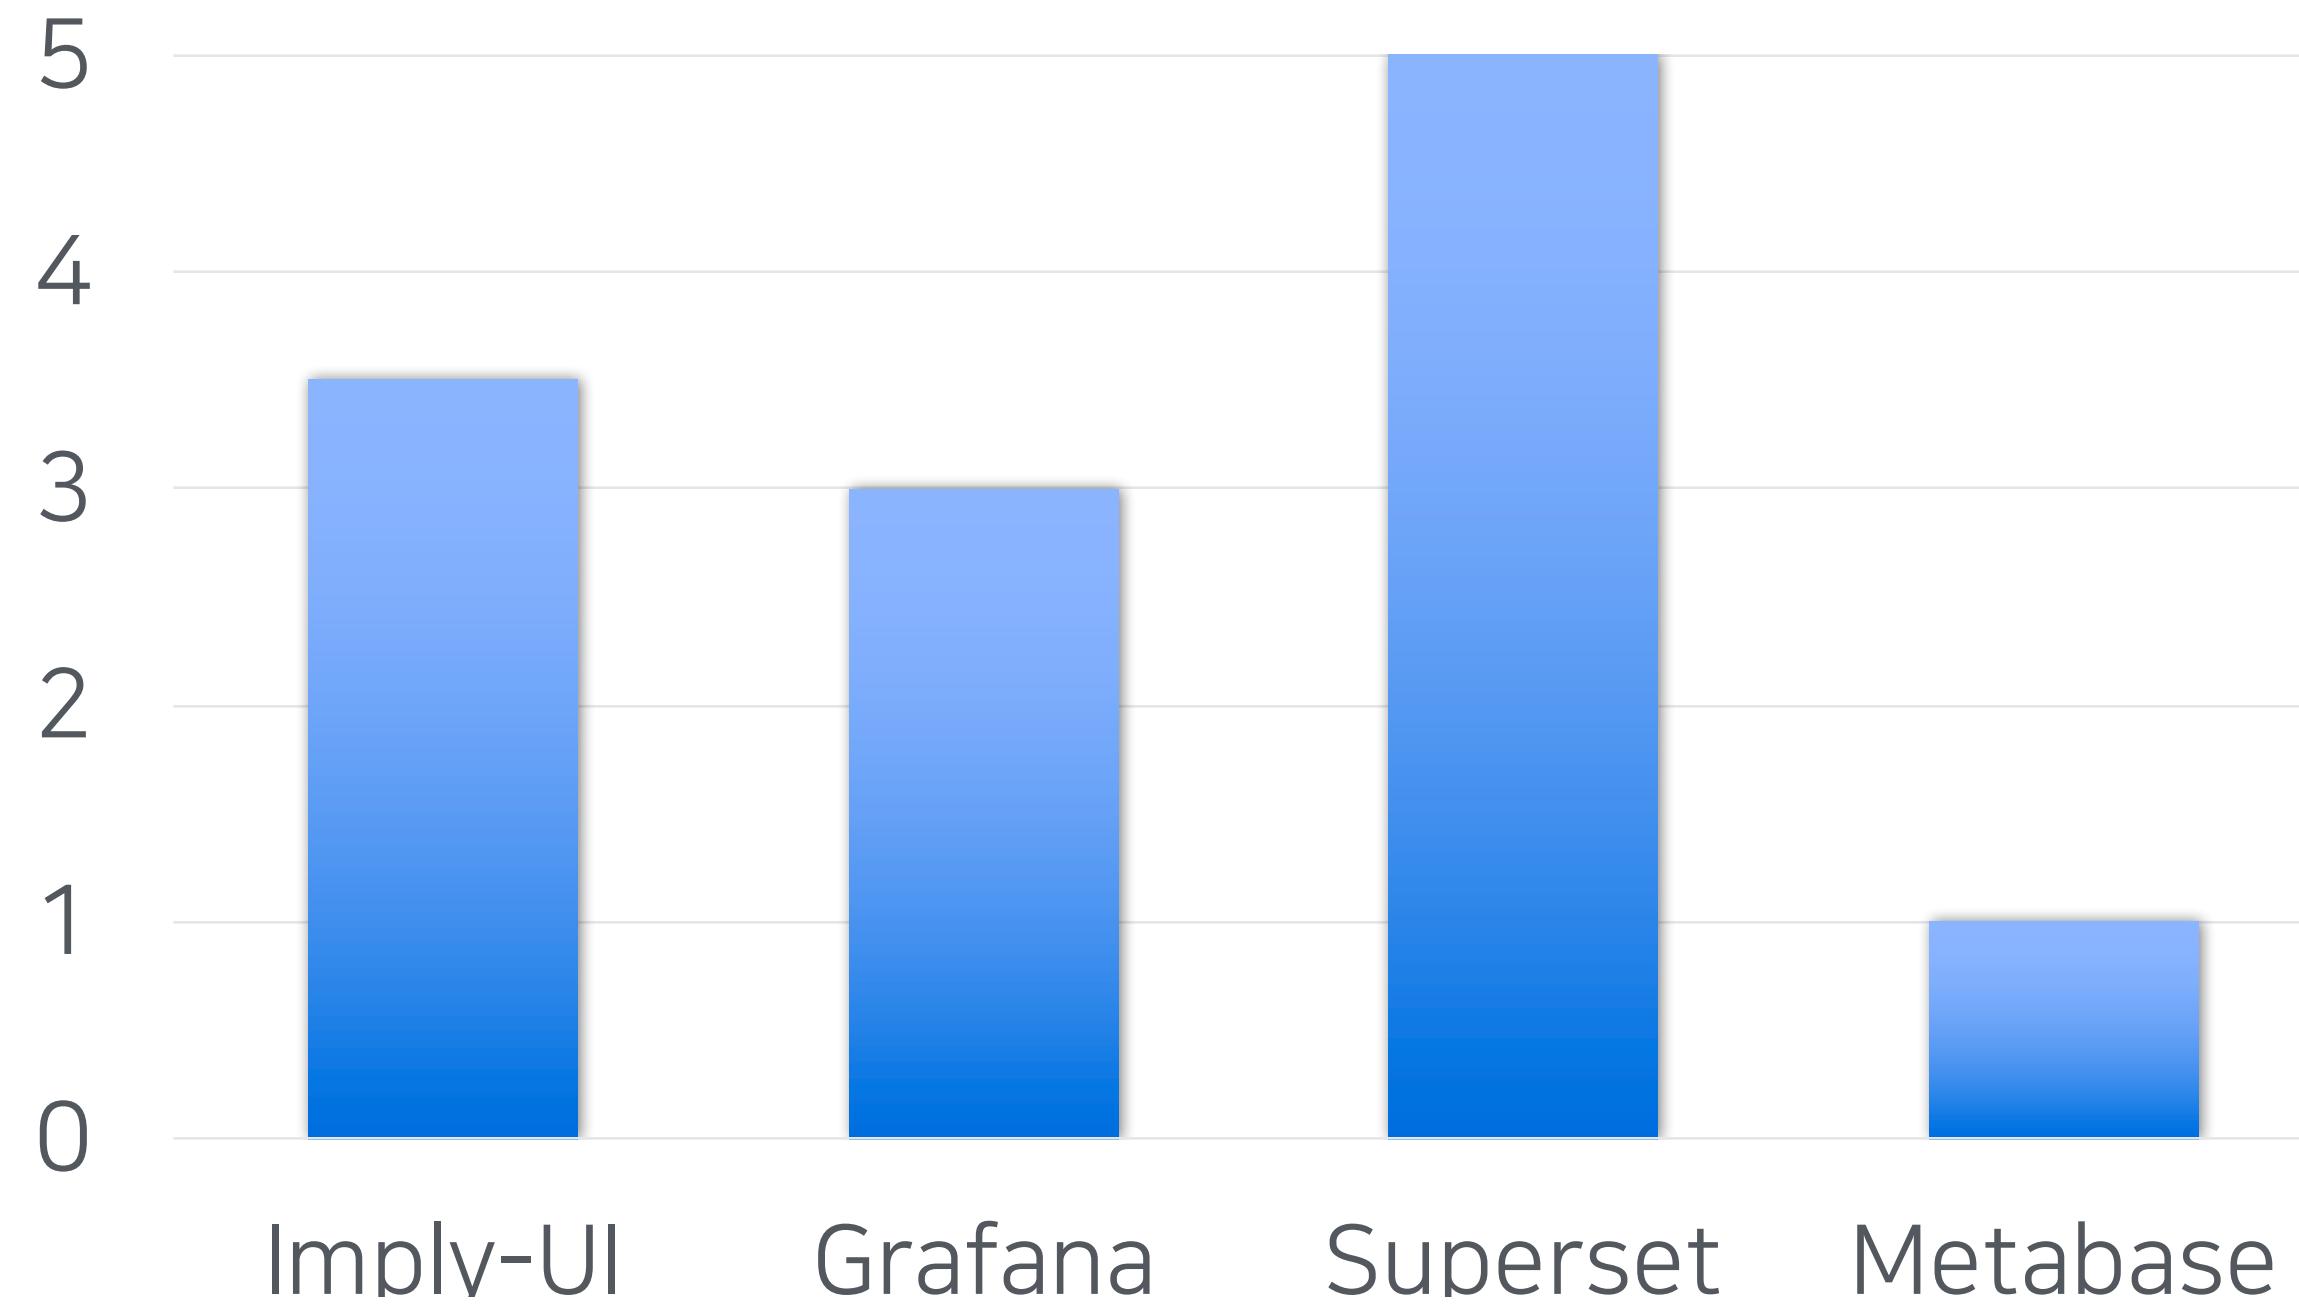


Superset

# 3.5 개인적인 총평

## 권한관리

(Datasource, Dashboard 권한설정)



## Metabase

	Administrators		All Users	
	DATA ACCESS	SQL QUERIES	DATA ACCESS	SQL QUERIES
Sample Dataset	✓	✗	✓	✗
Local-druid	✓	✗	✓	✗

## Superset

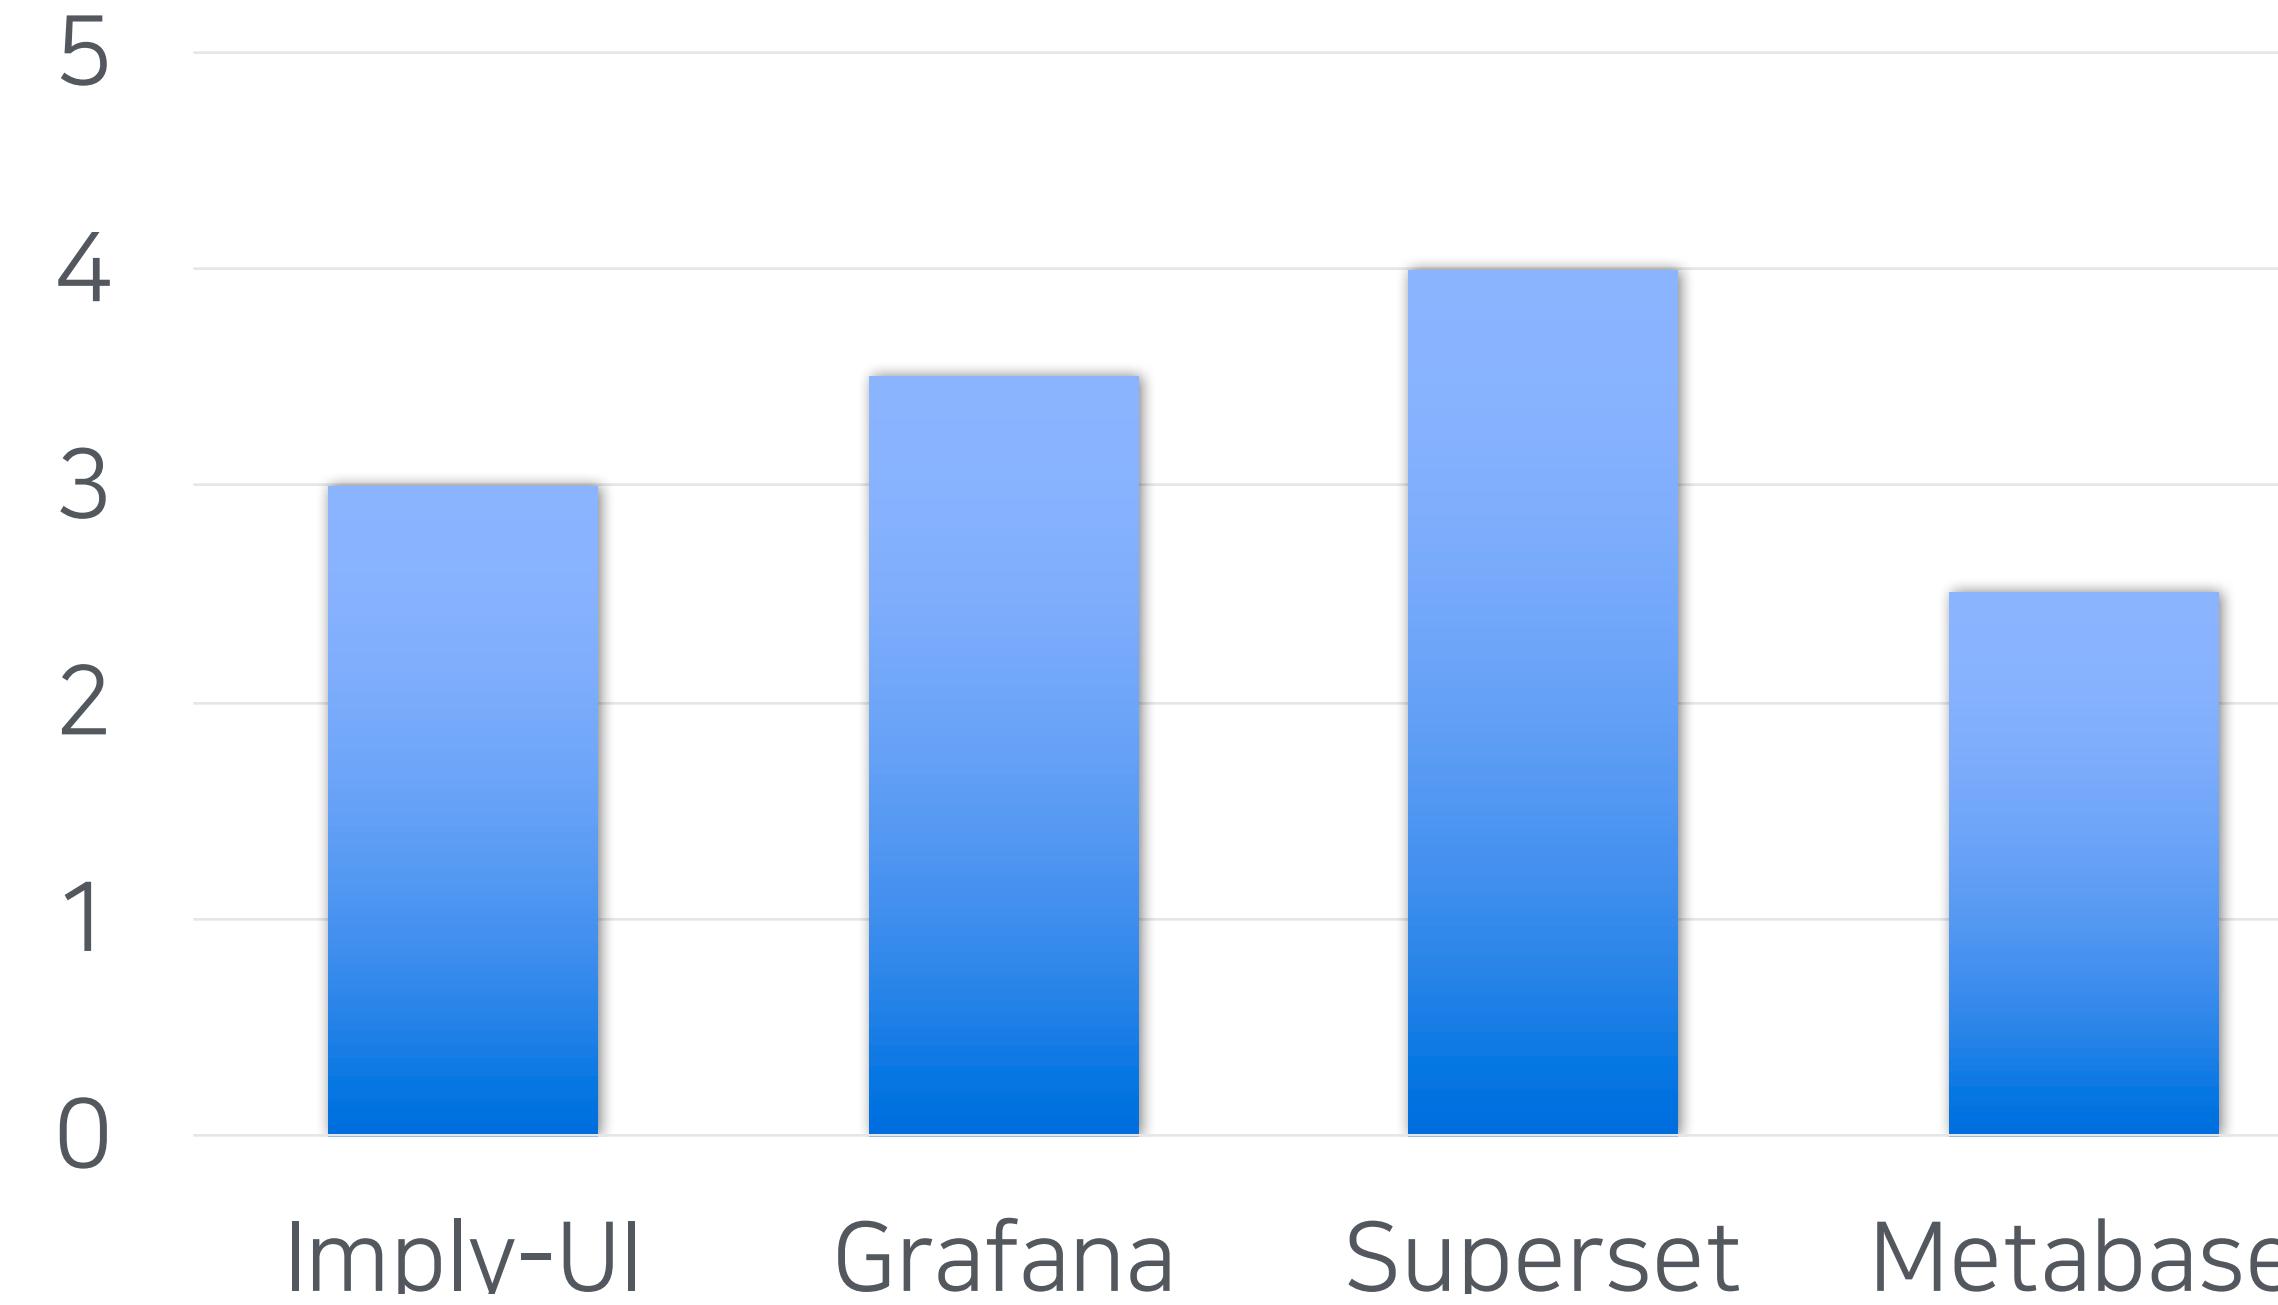
Name	Permissions
Admin	[can this form post on ResetPasswordView, can this form get on ResetMyPasswordView, can this form get on UserInfoEditView, can this form get on UserDBModelView, can userinfo on UserDBModelView, can delete on UserDBModelView, can download on UserDBModelView, can list on UserDBModelView, can add on UserDBModelView, can show on UserDBModelView, resetmypassword on UserDBModelView, resetpasswords on UserDBModelView, userinfoedit on UserDBModelView, menu access on List Users, menu access on Security, can edit on RoleModelView, can delete on RoleModelView, can download on RoleModelView, can list on RoleModelView, can add on RoleModelView, can show on RoleModelView, Copy Role on RoleModelView, menu access on List Roles, can chart on UserStatsChartView, menu access on User's Statistics, can list on PermissionModelView, menu access on Base Permissions, can list on ViewMenuModelView, menu access on Views/Menus, can list on PermissionViewModelView, menu access on Permission on Views/Menus, menu access on Import Dashboards, menu access on Manage, can edit on DatabaseView, can delete on DatabaseView, can download on DatabaseView, can list on DatabaseView, can add on DatabaseView, can show on DatabaseView, can delete on DatabaseView, can download on DatabaseView, can list on DatabaseAsync, can add on DatabaseAsync, can edit on DatabaseAsync, can delete on DatabaseAsync, can download on DatabaseAsync, can list on DatabaseAsync, can add on DatabaseTablesAsync, can show on DatabaseAsync, can delete on DatabaseTablesAsync, can edit on DatabaseTablesAsync, can delete on DatabaseTablesAsync]

- Security
- List Users
- List Roles
- User's Statistics
- Base Permissions
- Views/Menus
- Permission on Views/Menus
- Access requests
- Action Log

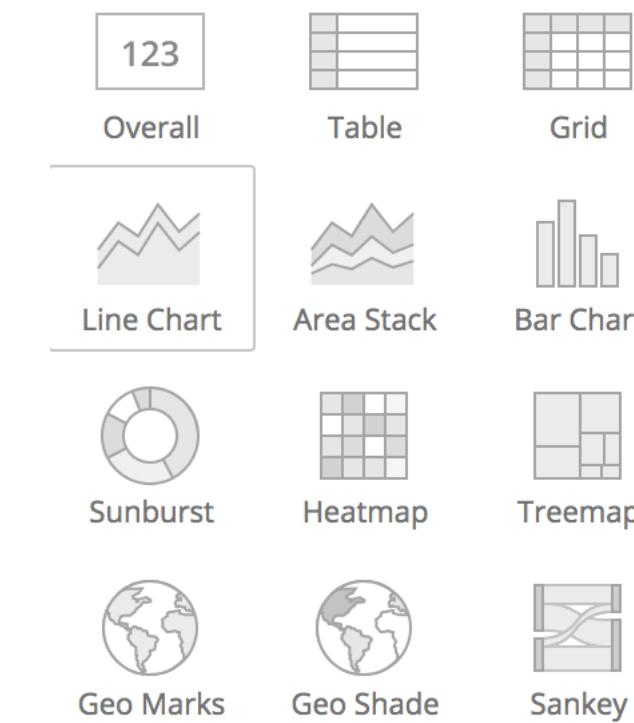
# 3.5 개인적인 총평

## 시각화

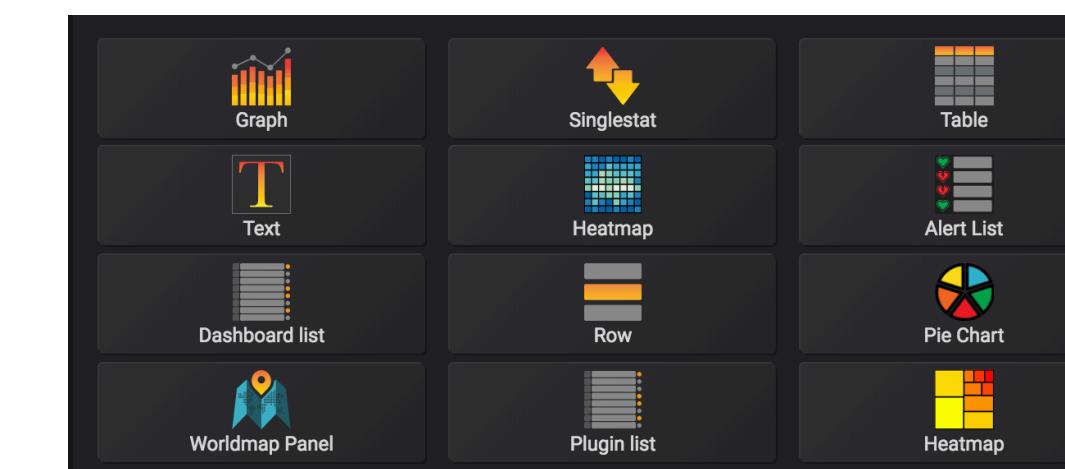
(표현가능 그래프, 표현 자유도)



## Implv UI



## Grafana

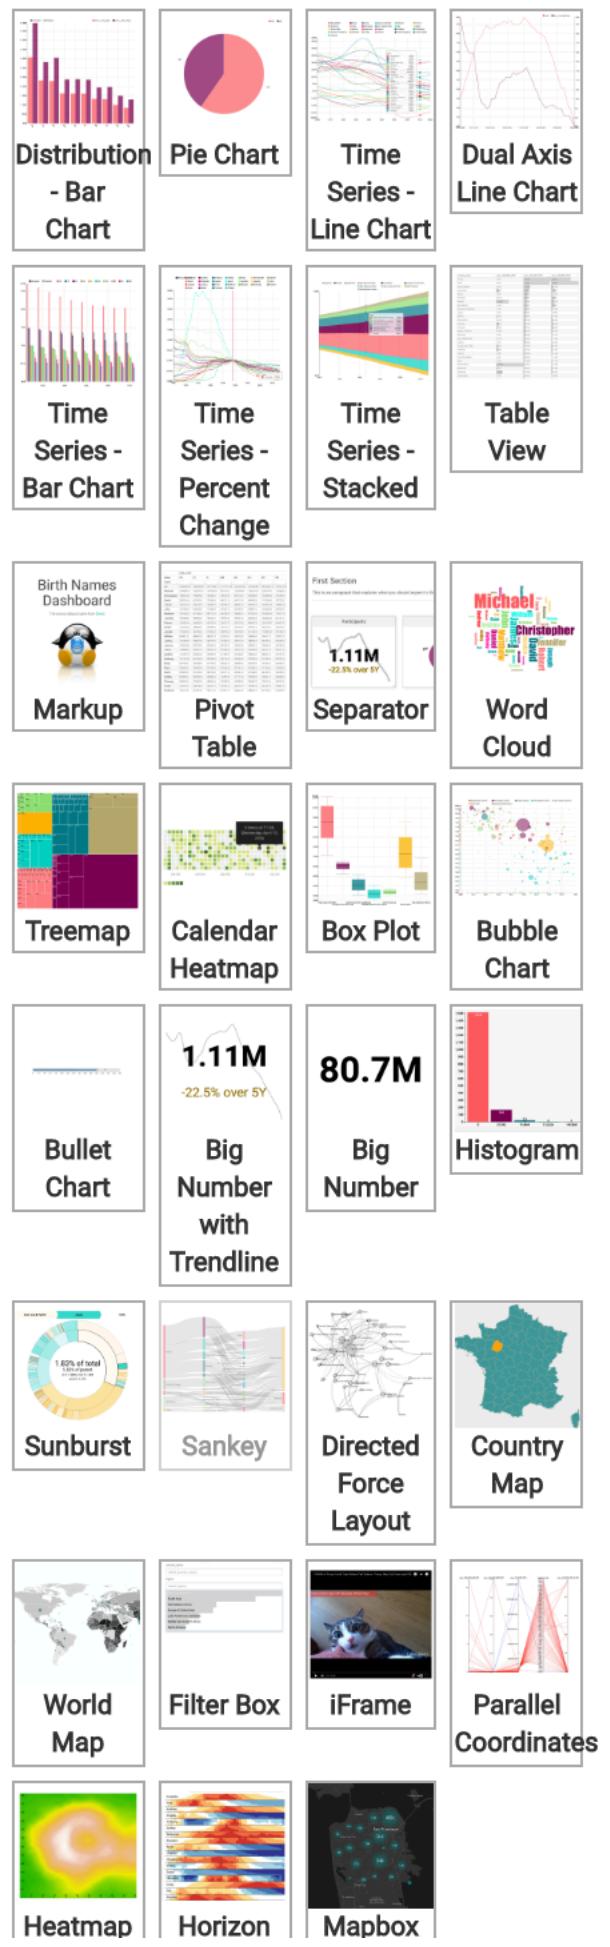


## Metabase



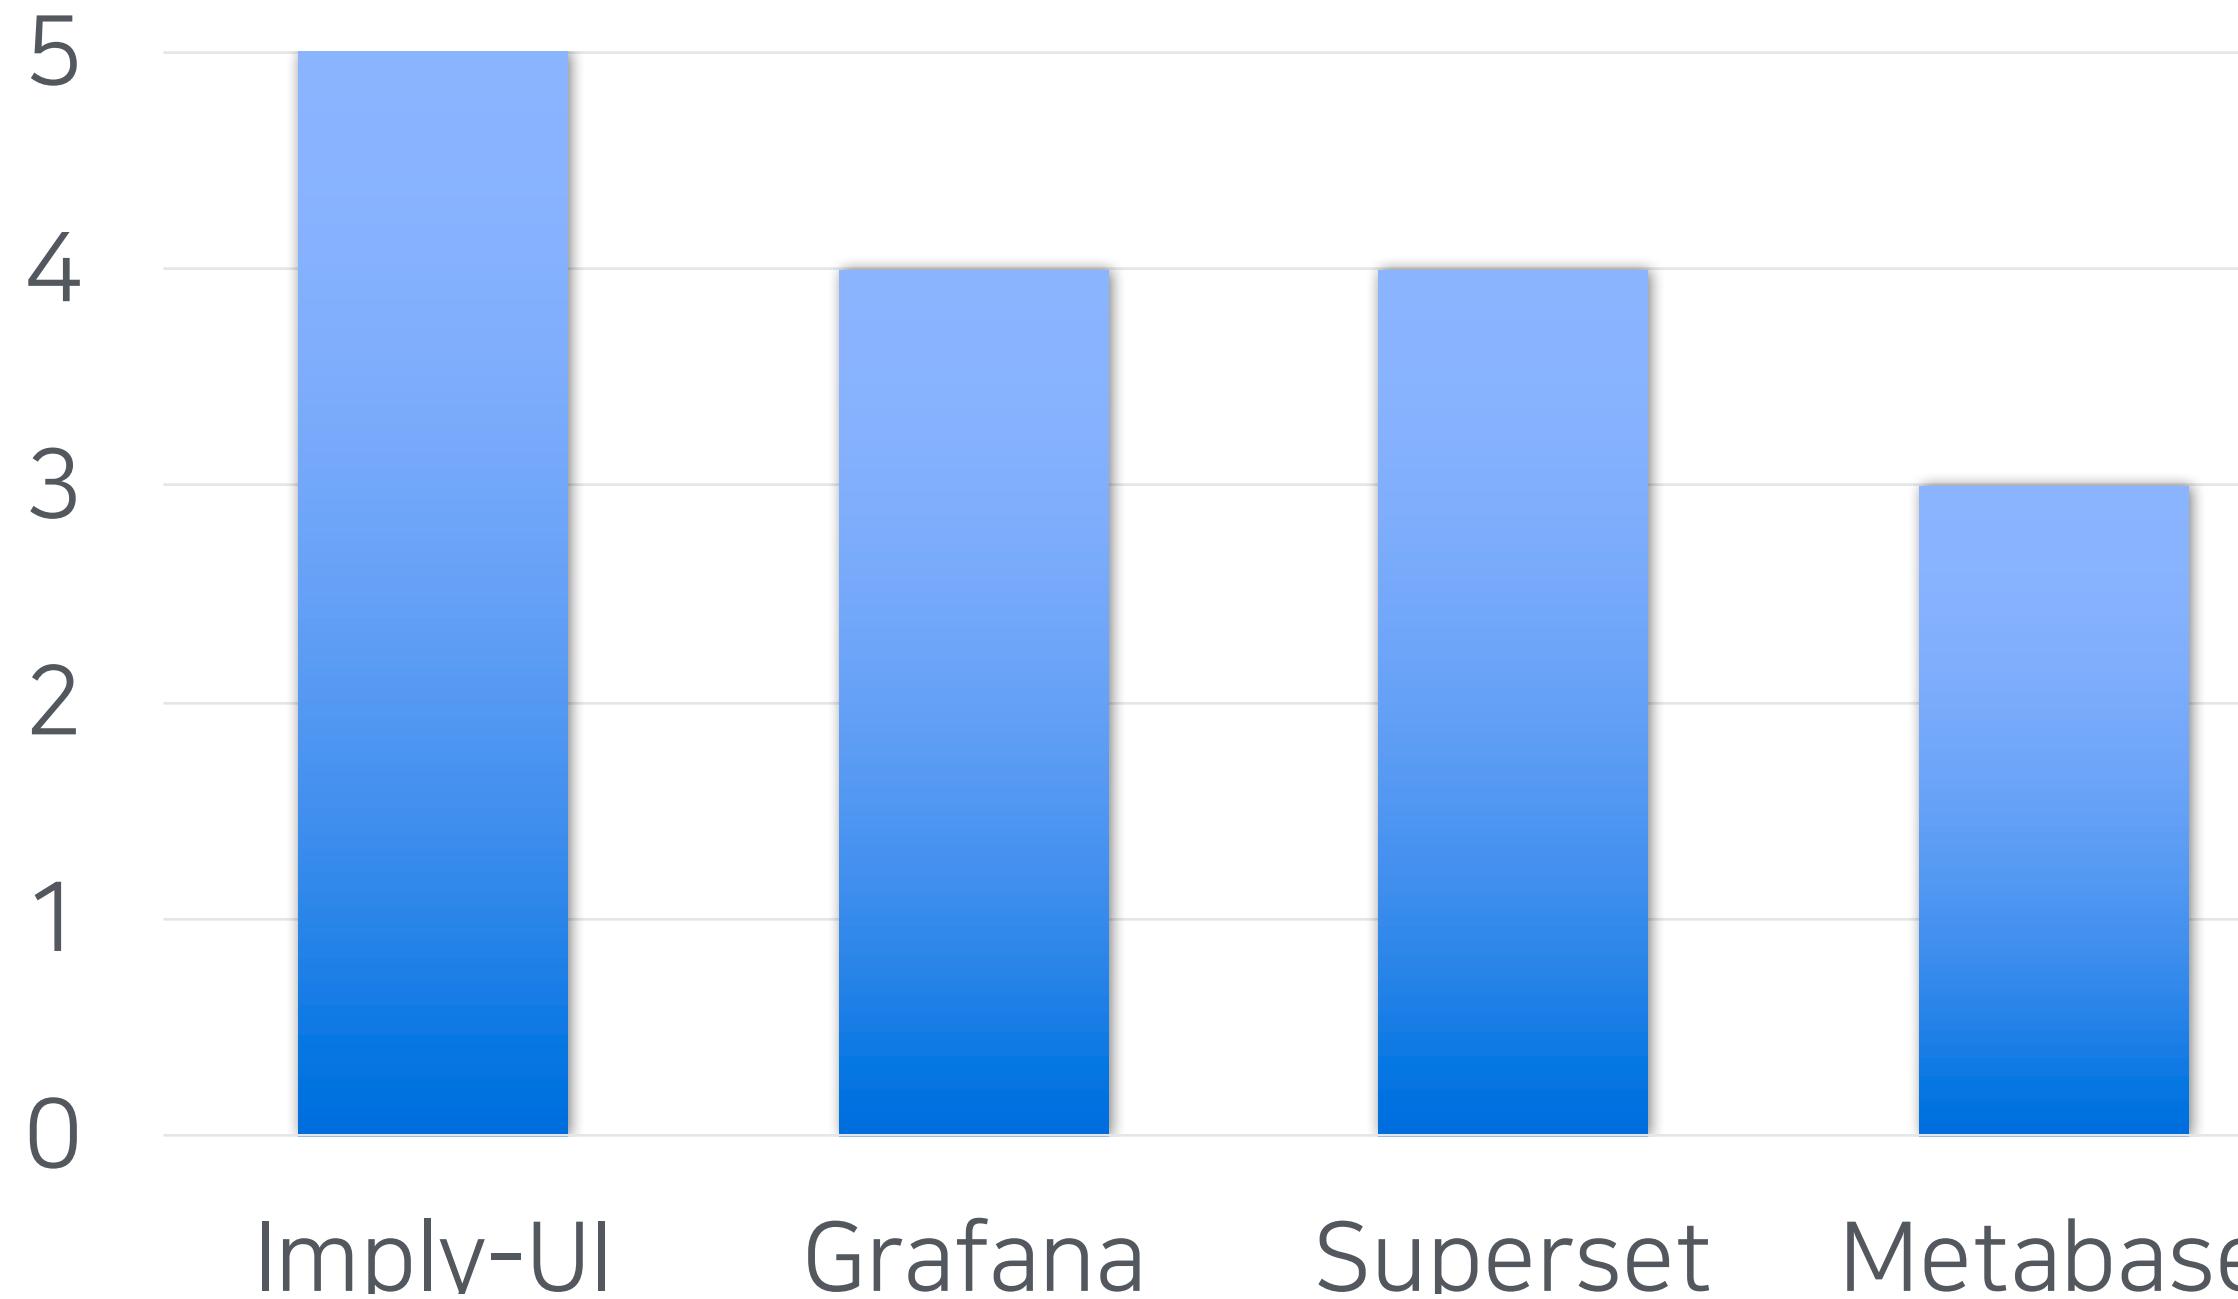
- Progress
- Table
- ✓ Line
- ▲ Area
- . Bar
- Row Chart
- Scatter
- Pie
- Map
- ▼ Funnel

## Superset

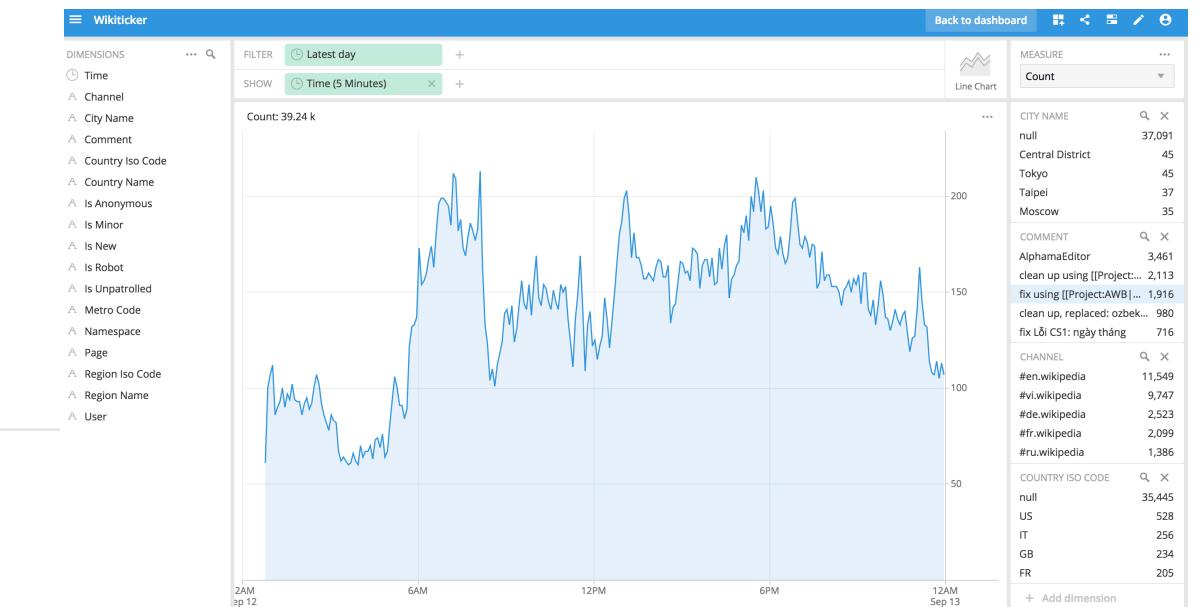


# 3.5 개인적인 총평

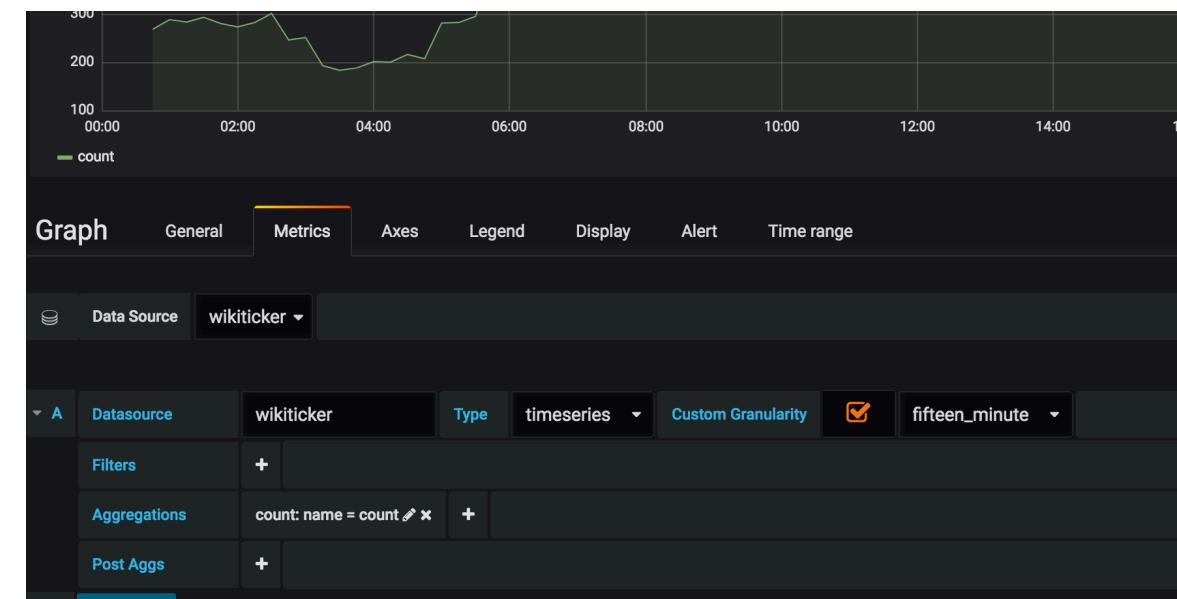
사용편의성  
(그래프설정방법)



Imply UI



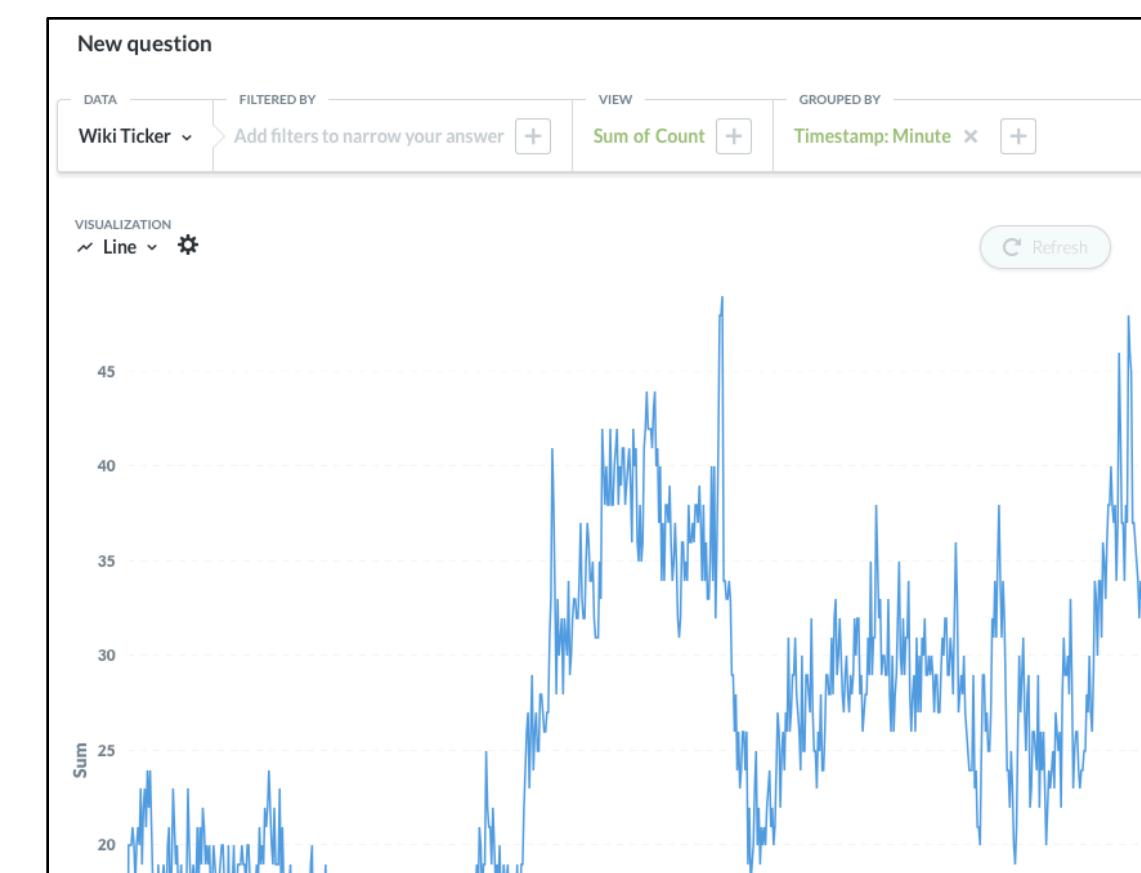
Grafana



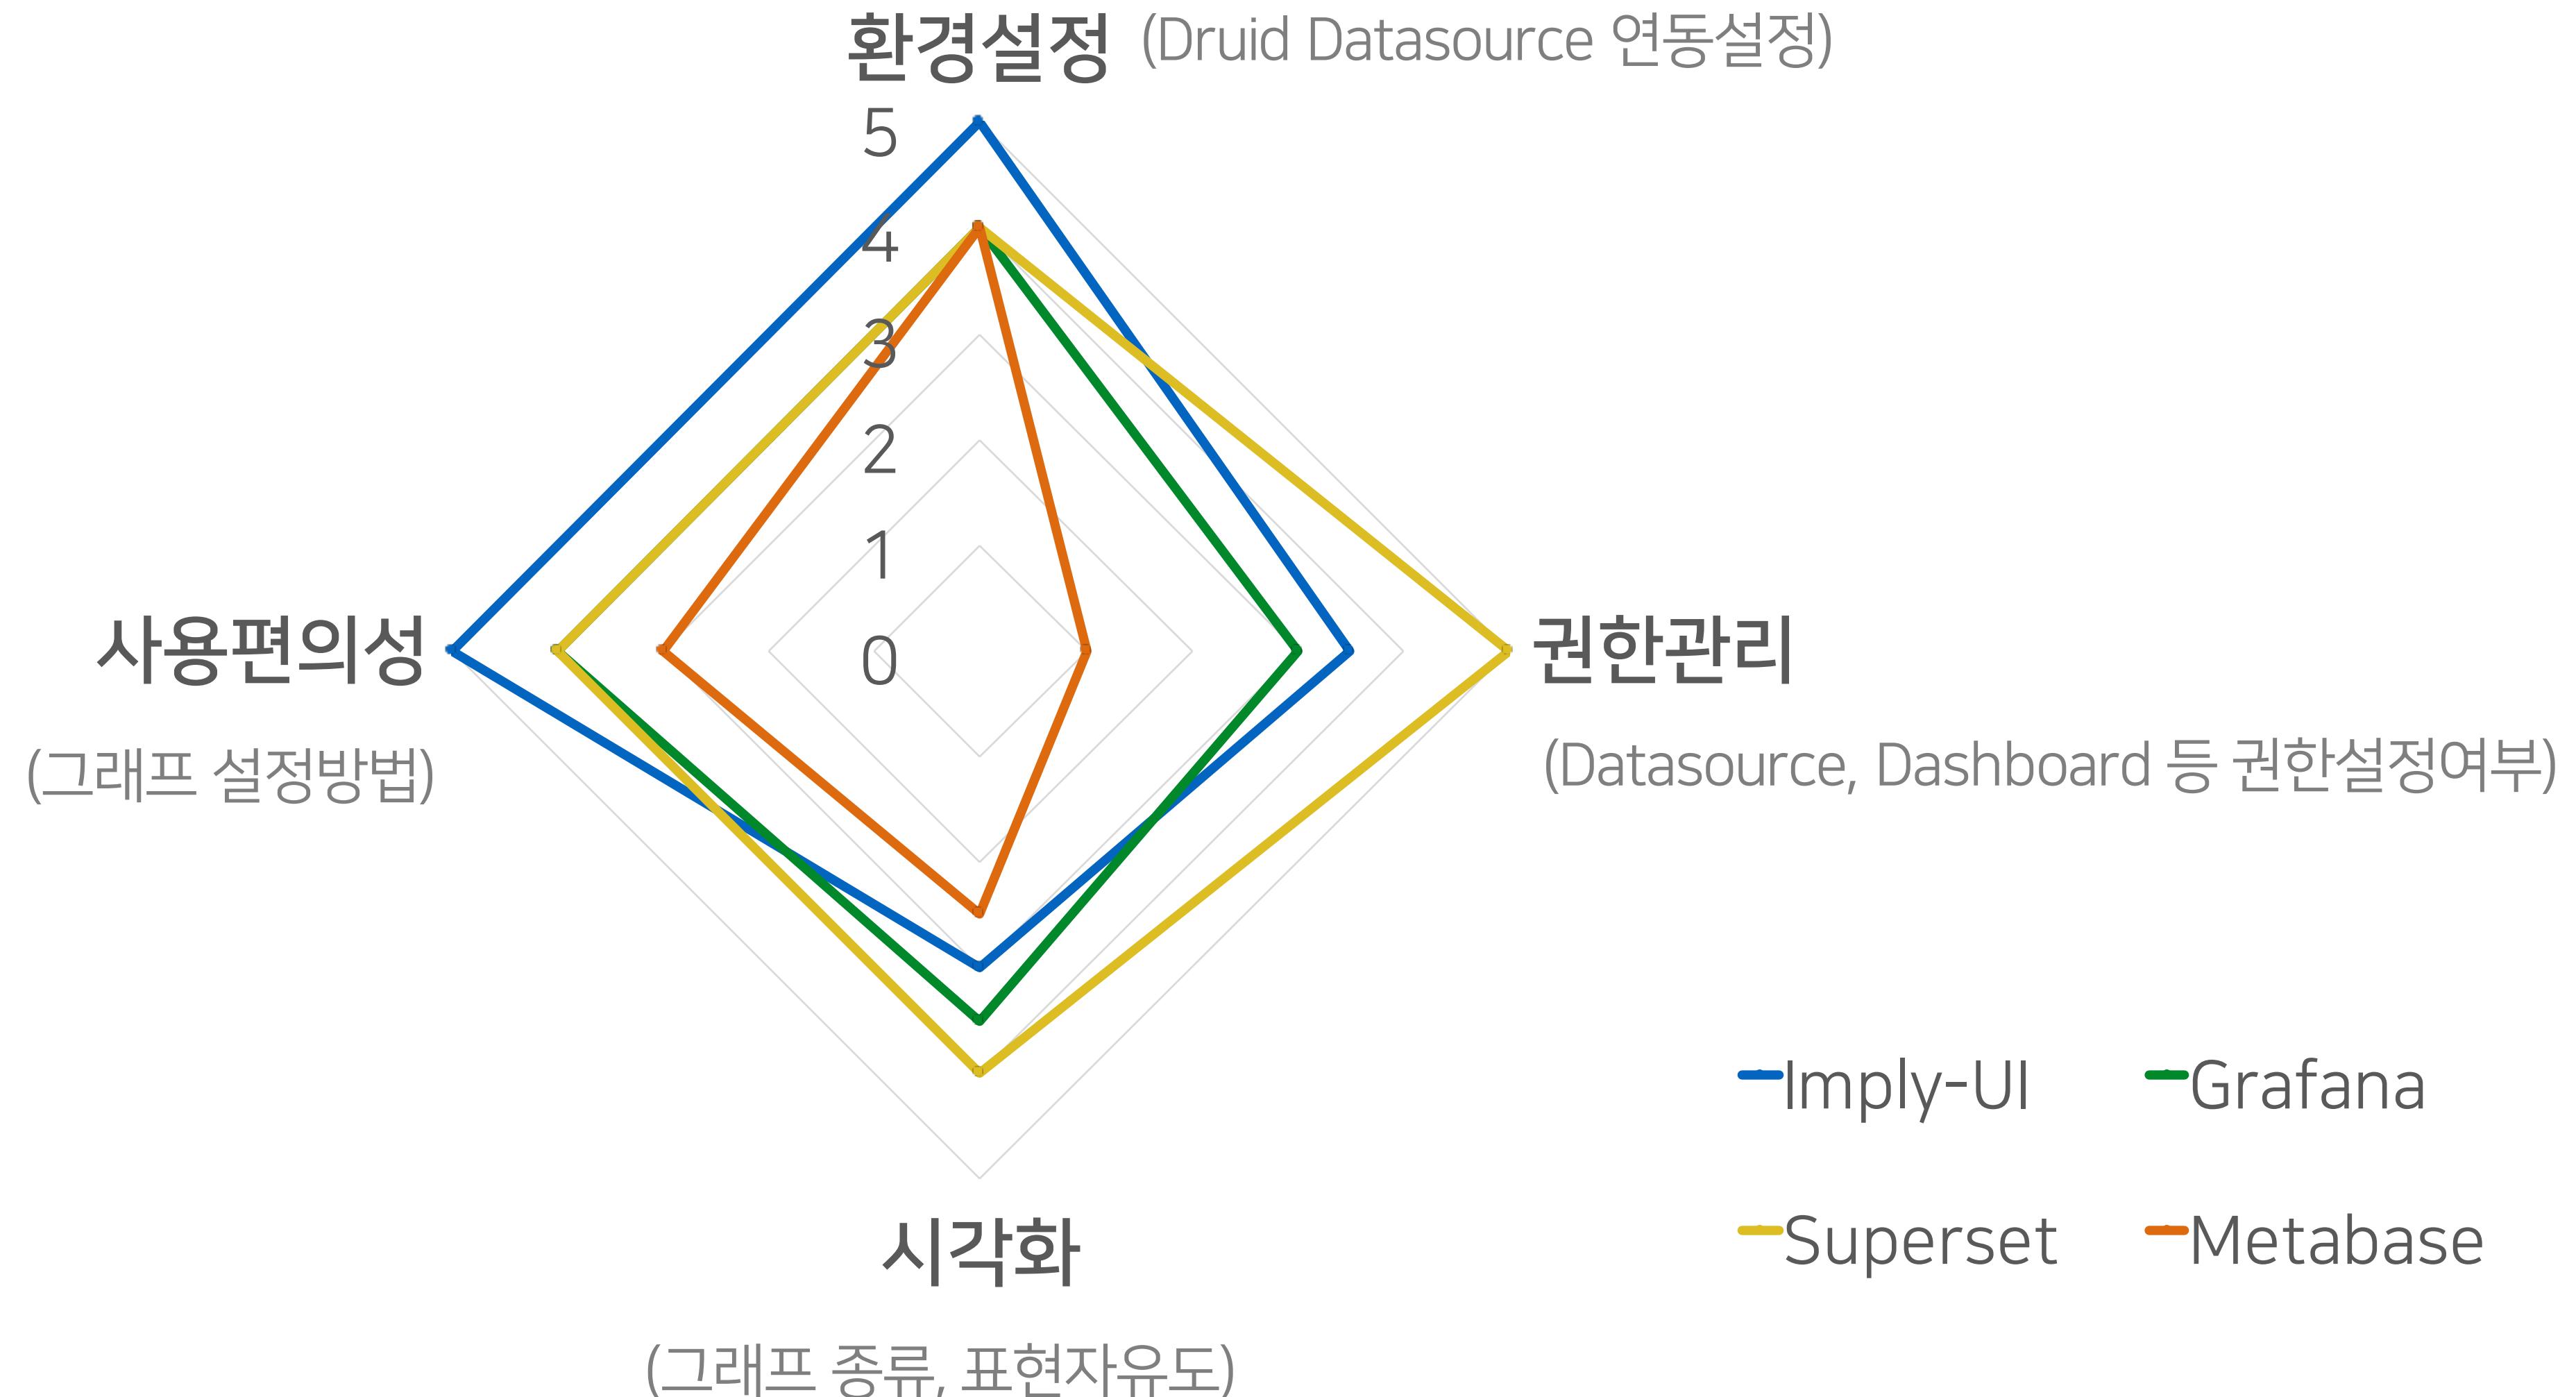
Superset



Metabase



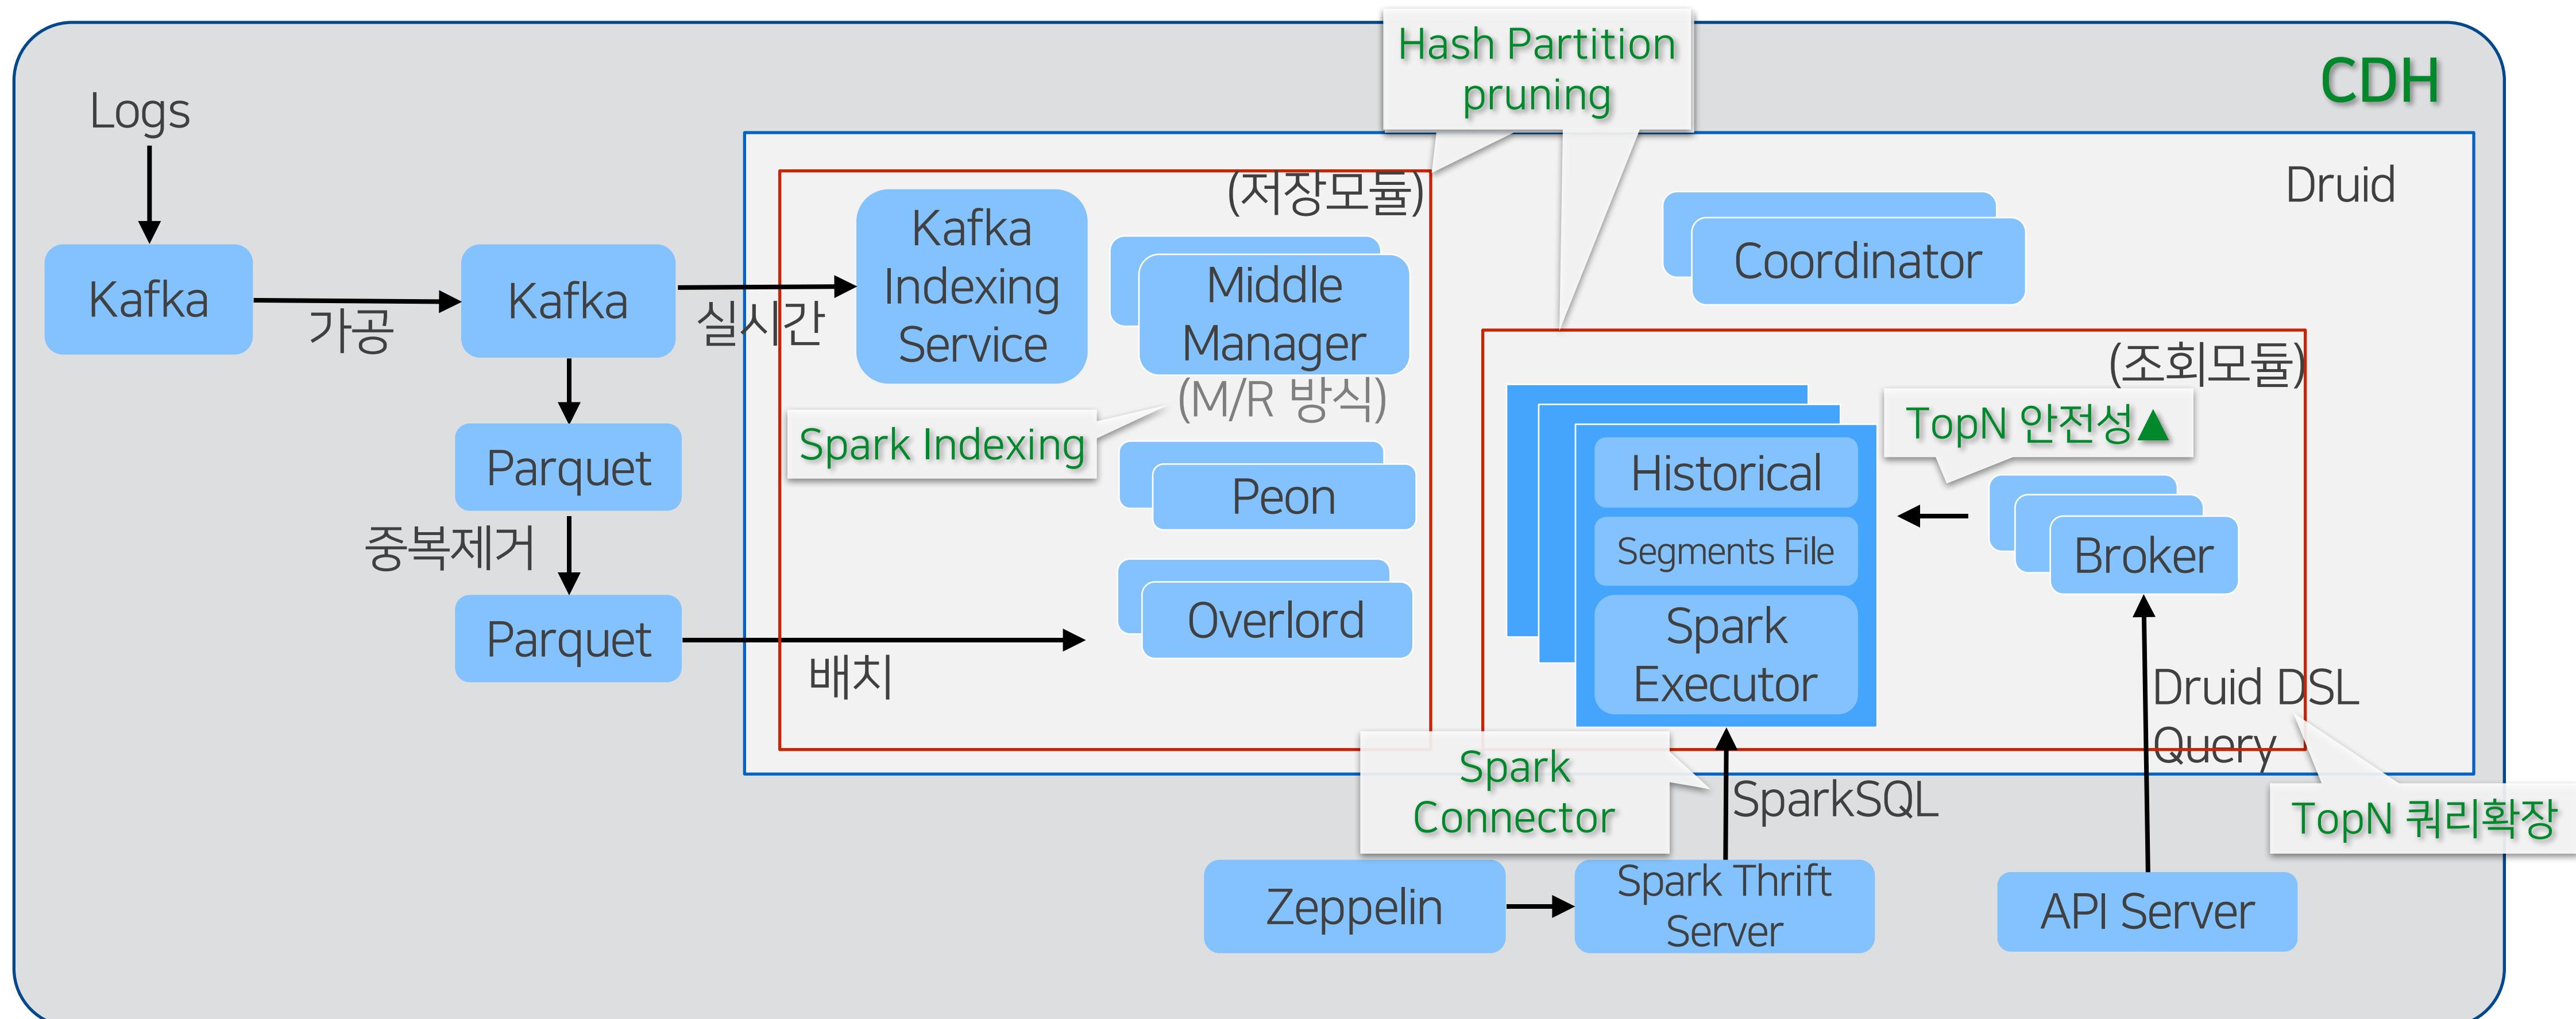
## 3.5 개인적인 총평



4.

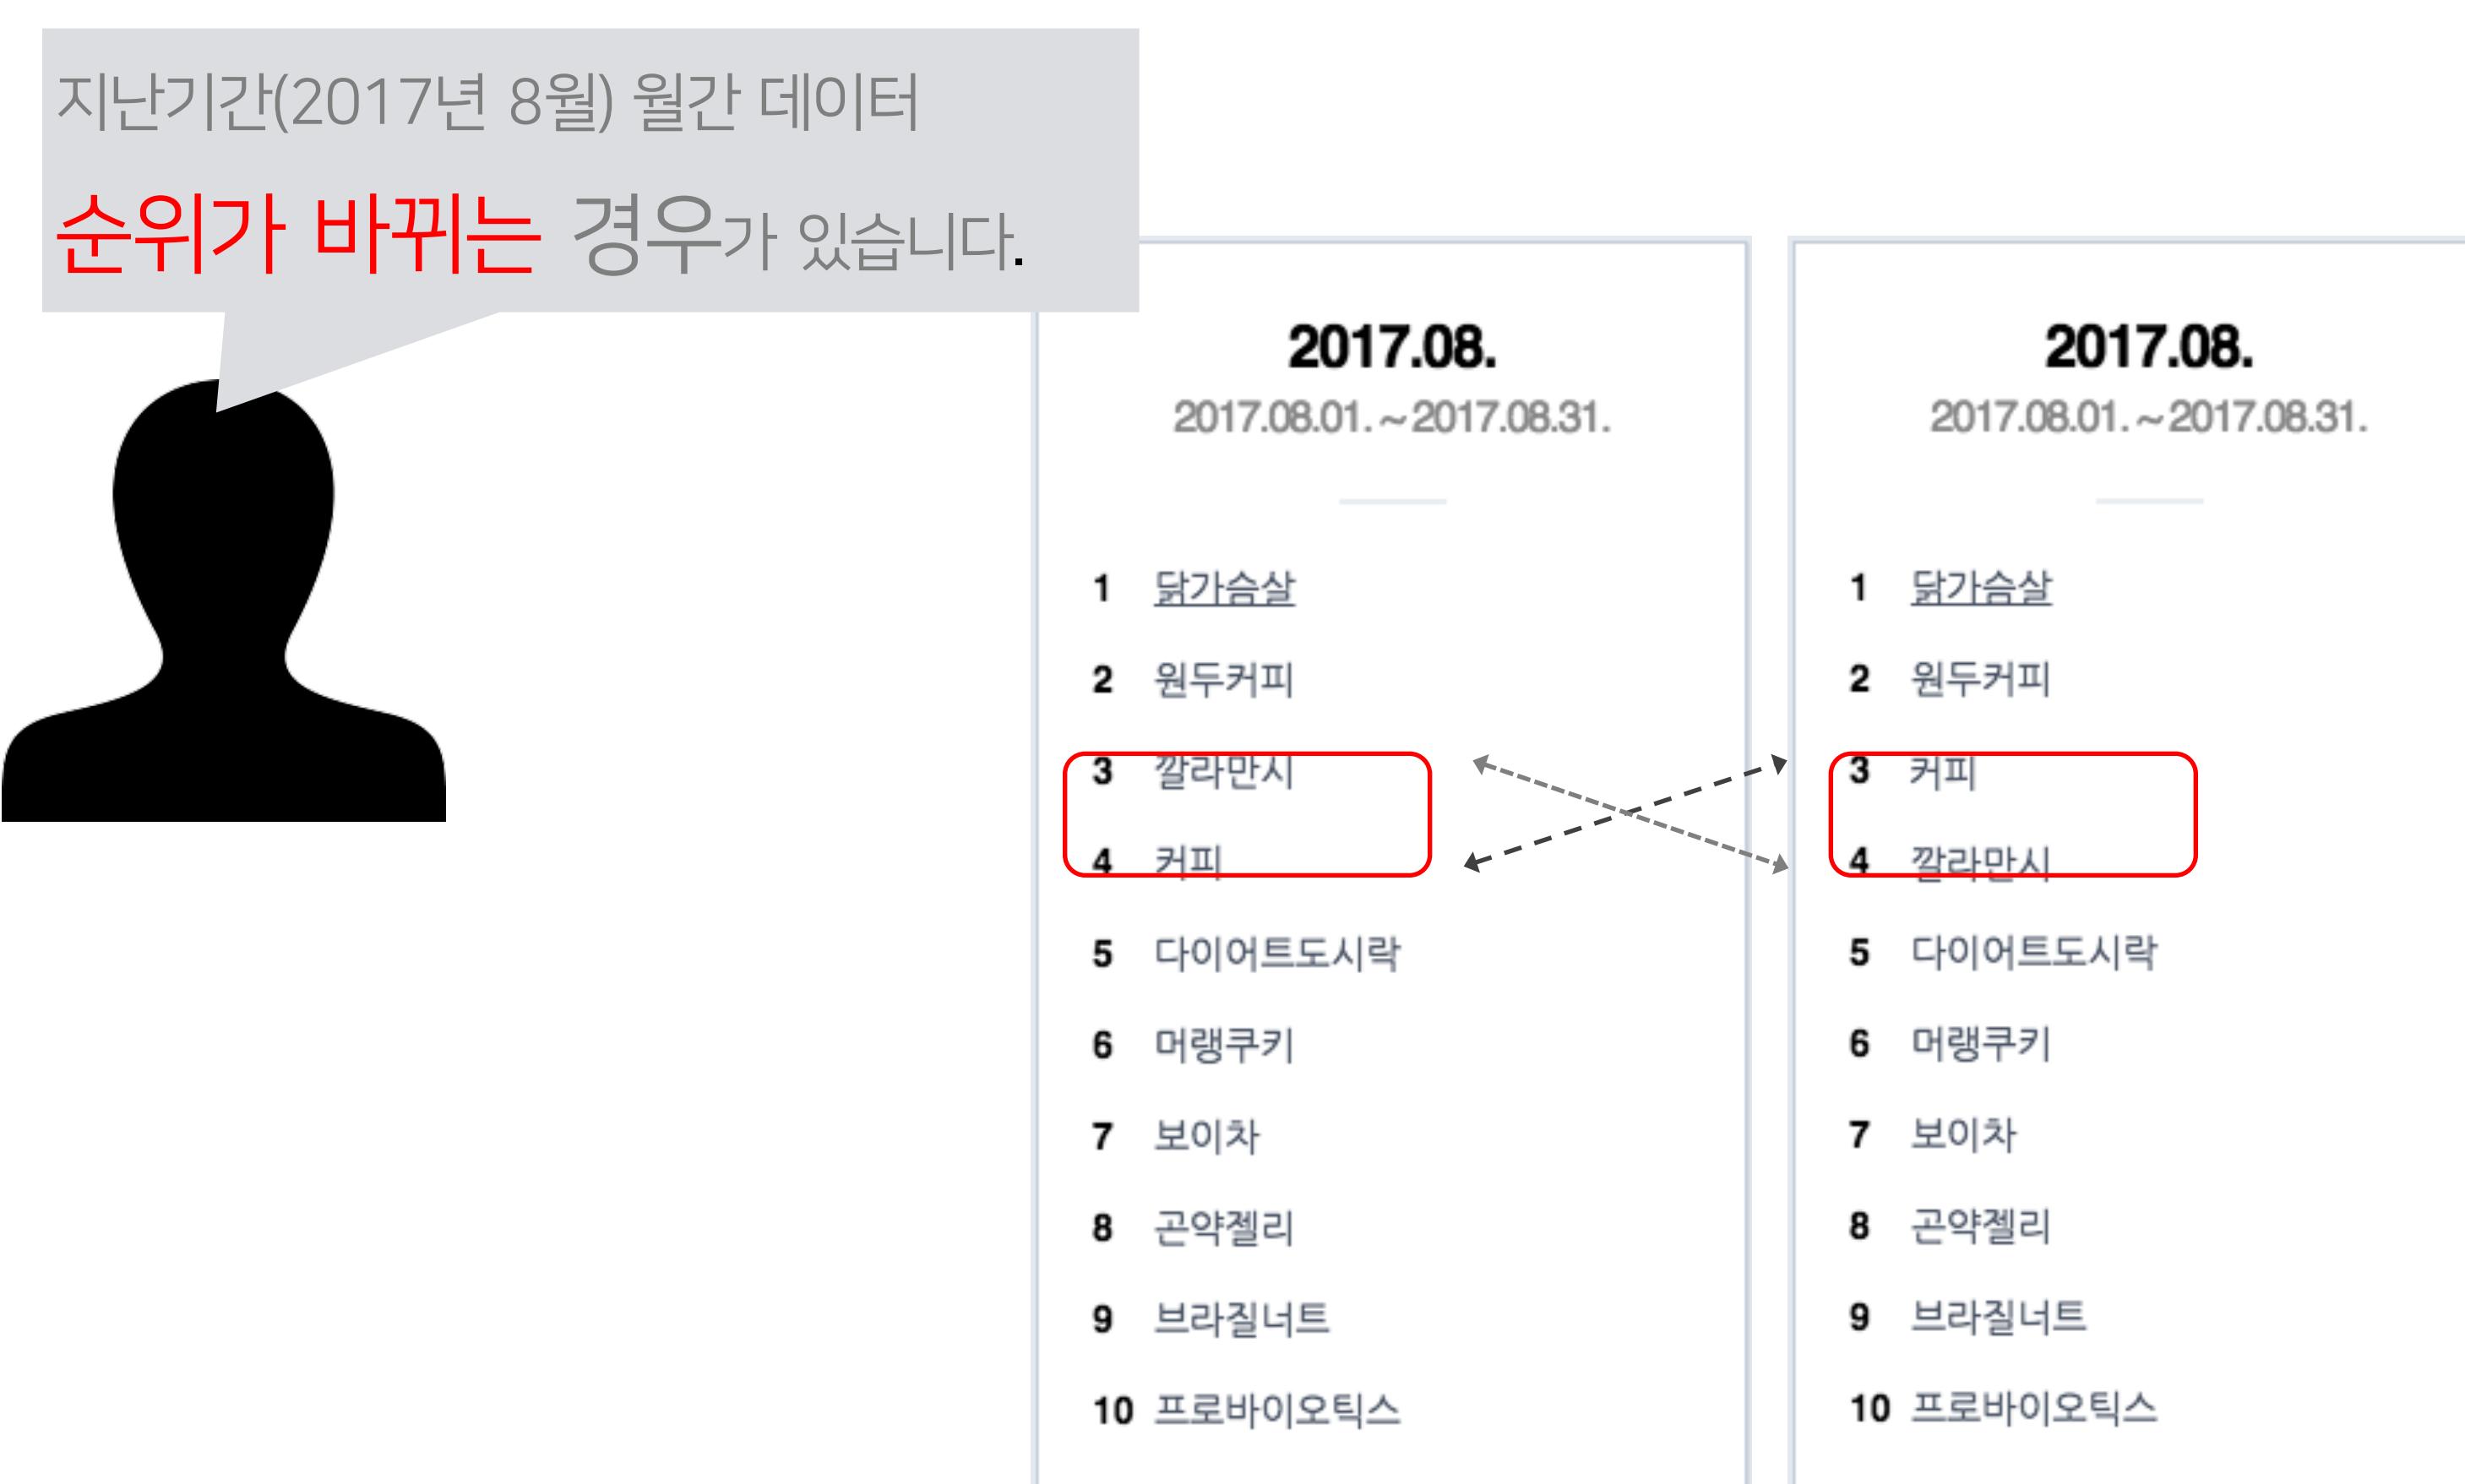
# Druid로 분석하기 - 고급편

# 4.1 Service Architecture



# 4.2 질의성능 향상시키기

## 1) TopN 결과 안정성 향상



## 4.2 질의성능 향상시키기

country	SUM(duration)
korea	114
uk	47
usa	21

Seg. A Ranking

country	SUM(duration)
uk	67
usa	25
korea	17

Seg. B Ranking

country	SUM(duration)
china	33
korea	23
japan	27

Seg. C Ranking

$$(A + B) + C \neq A + (B + C)$$

country	SUM(duration)
korea	131
uk	114
usa	46

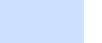
A + B

country	SUM(duration)
korea	154
uk	114
usa	46

 $(A + B) + C$ 

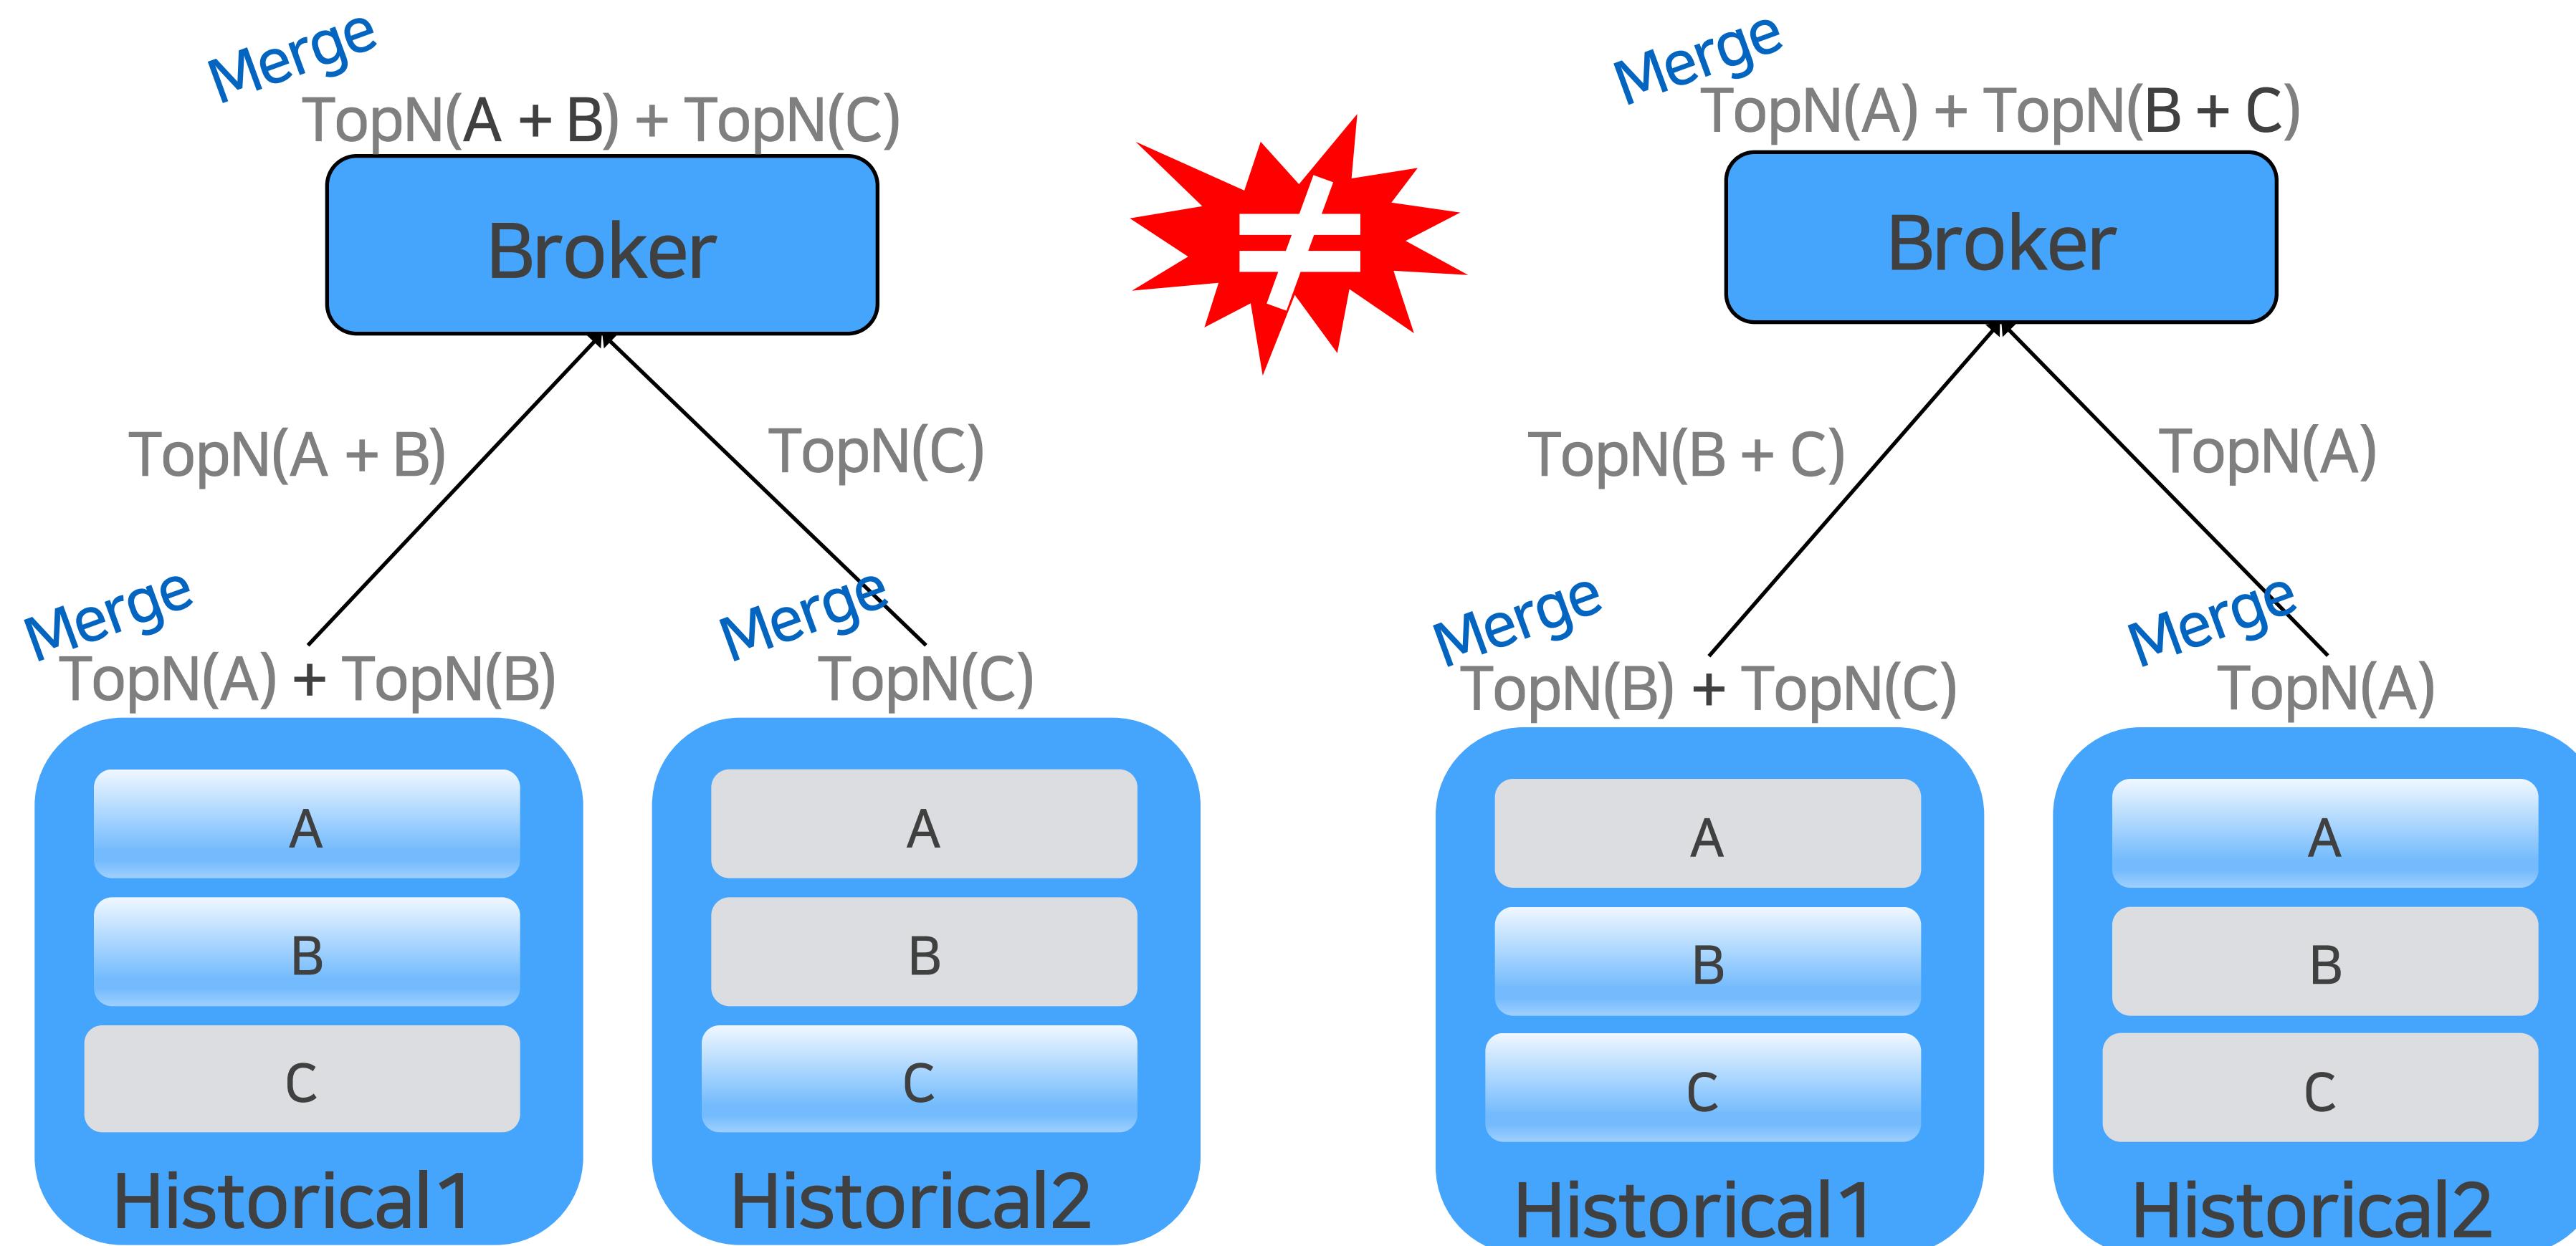
B + C

country	SUM(duration)
uk	67
korea	40
china	33
japan	27
usa	25

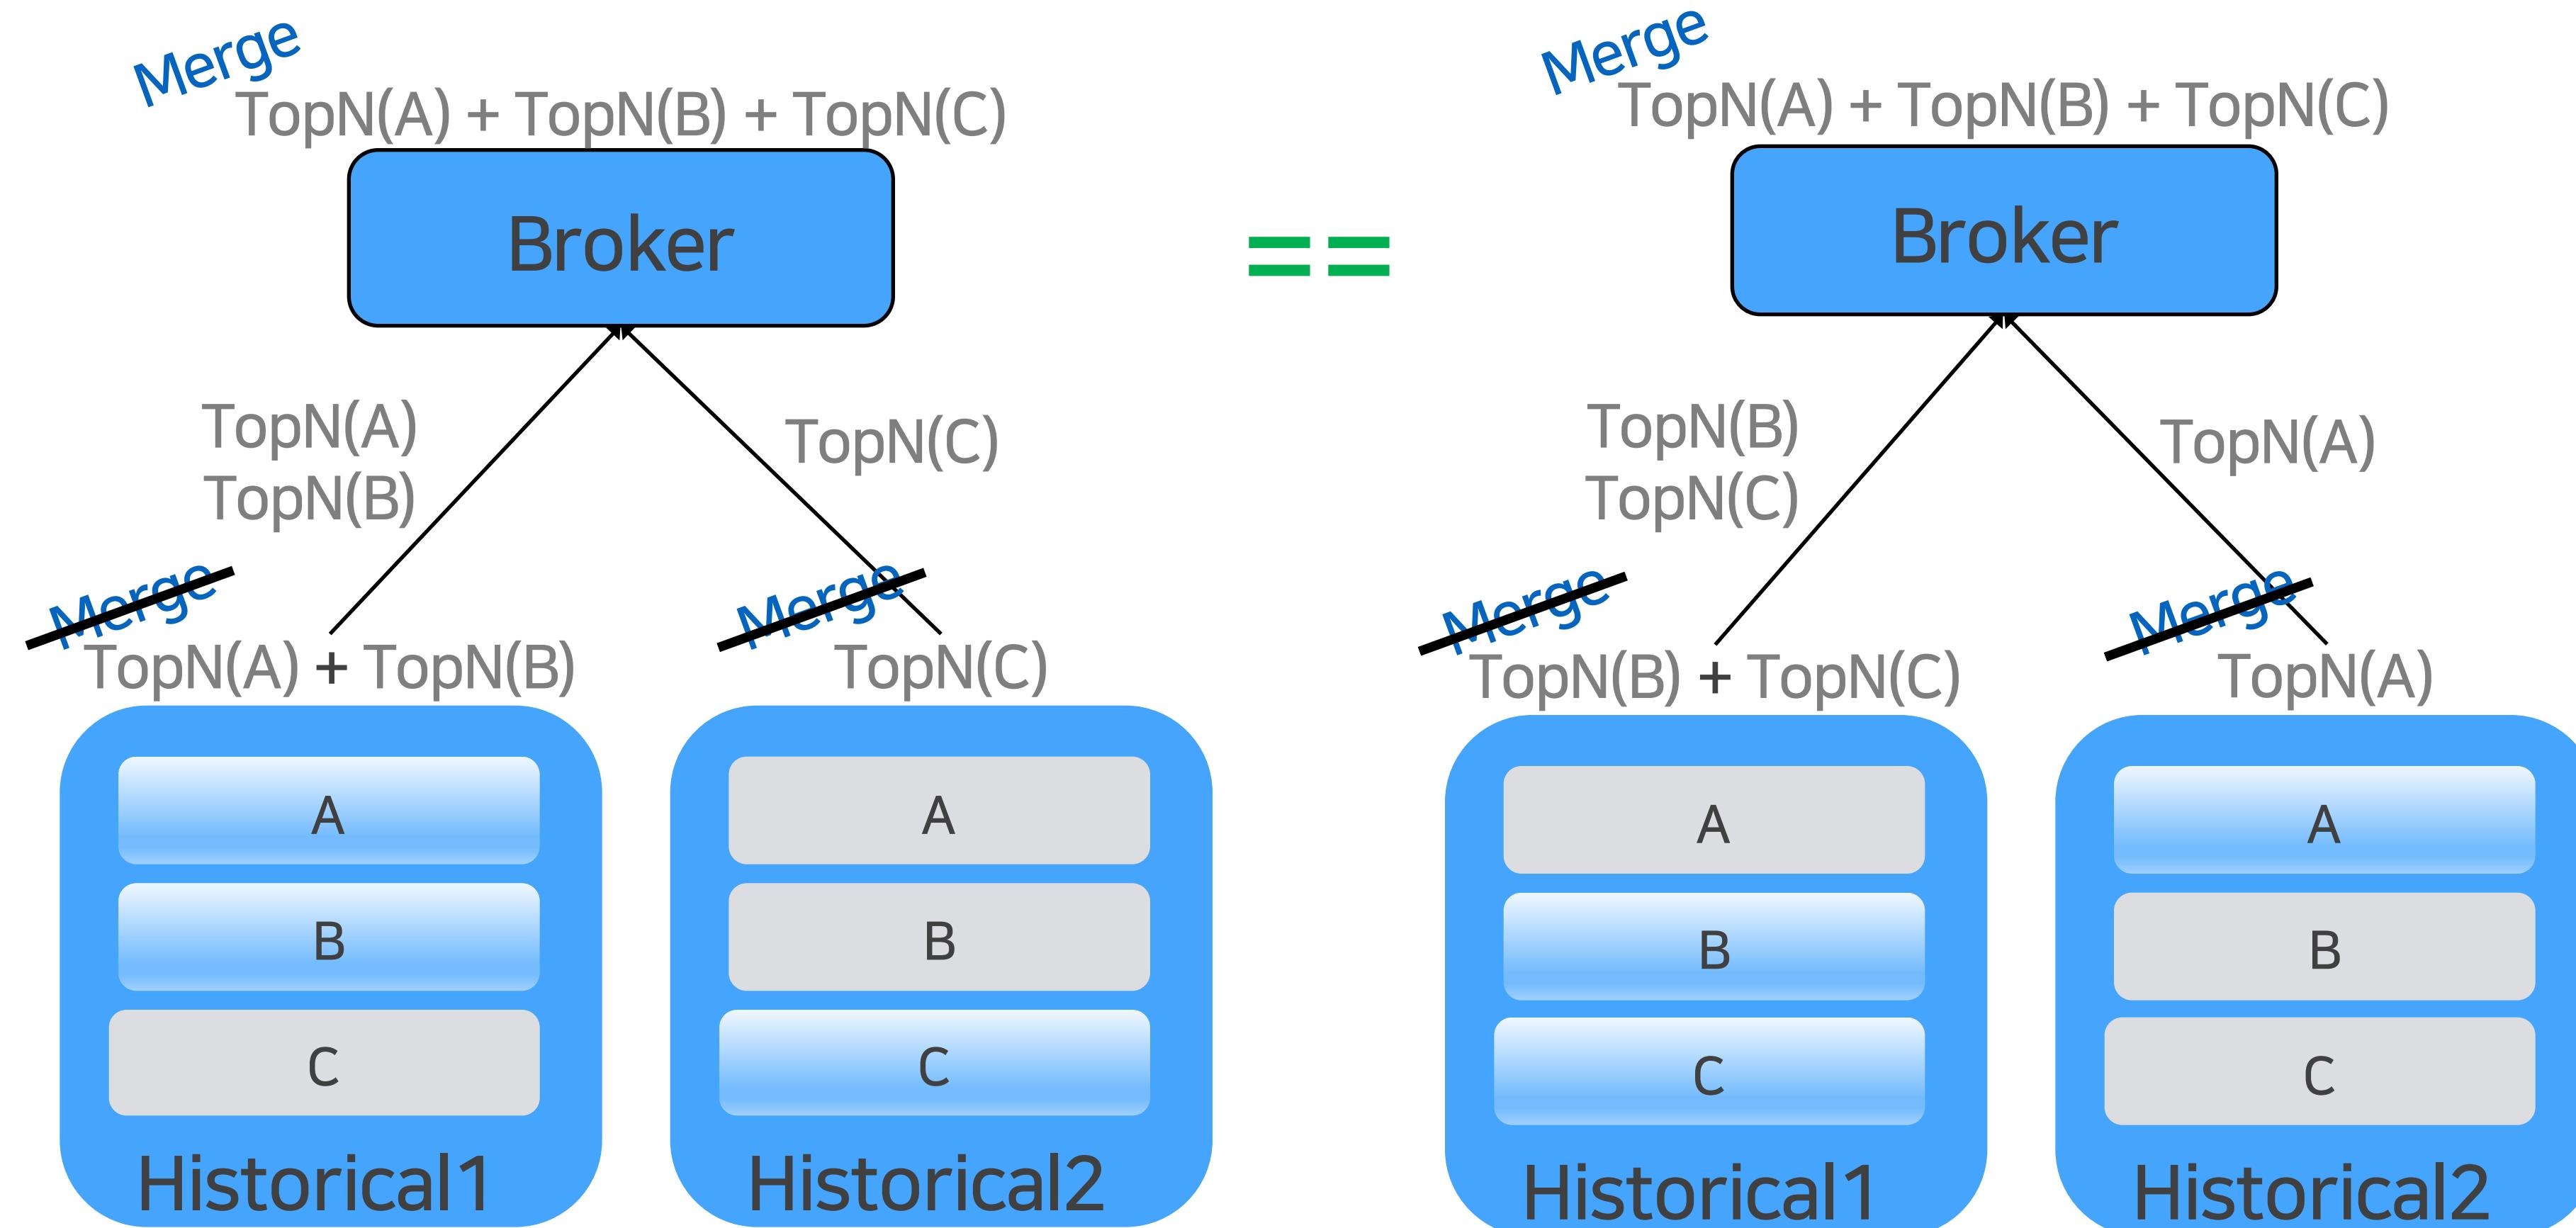
 $A + (B + C)$ 

Final Ranking  
Result

## 4.2 질의성능 향상시키기



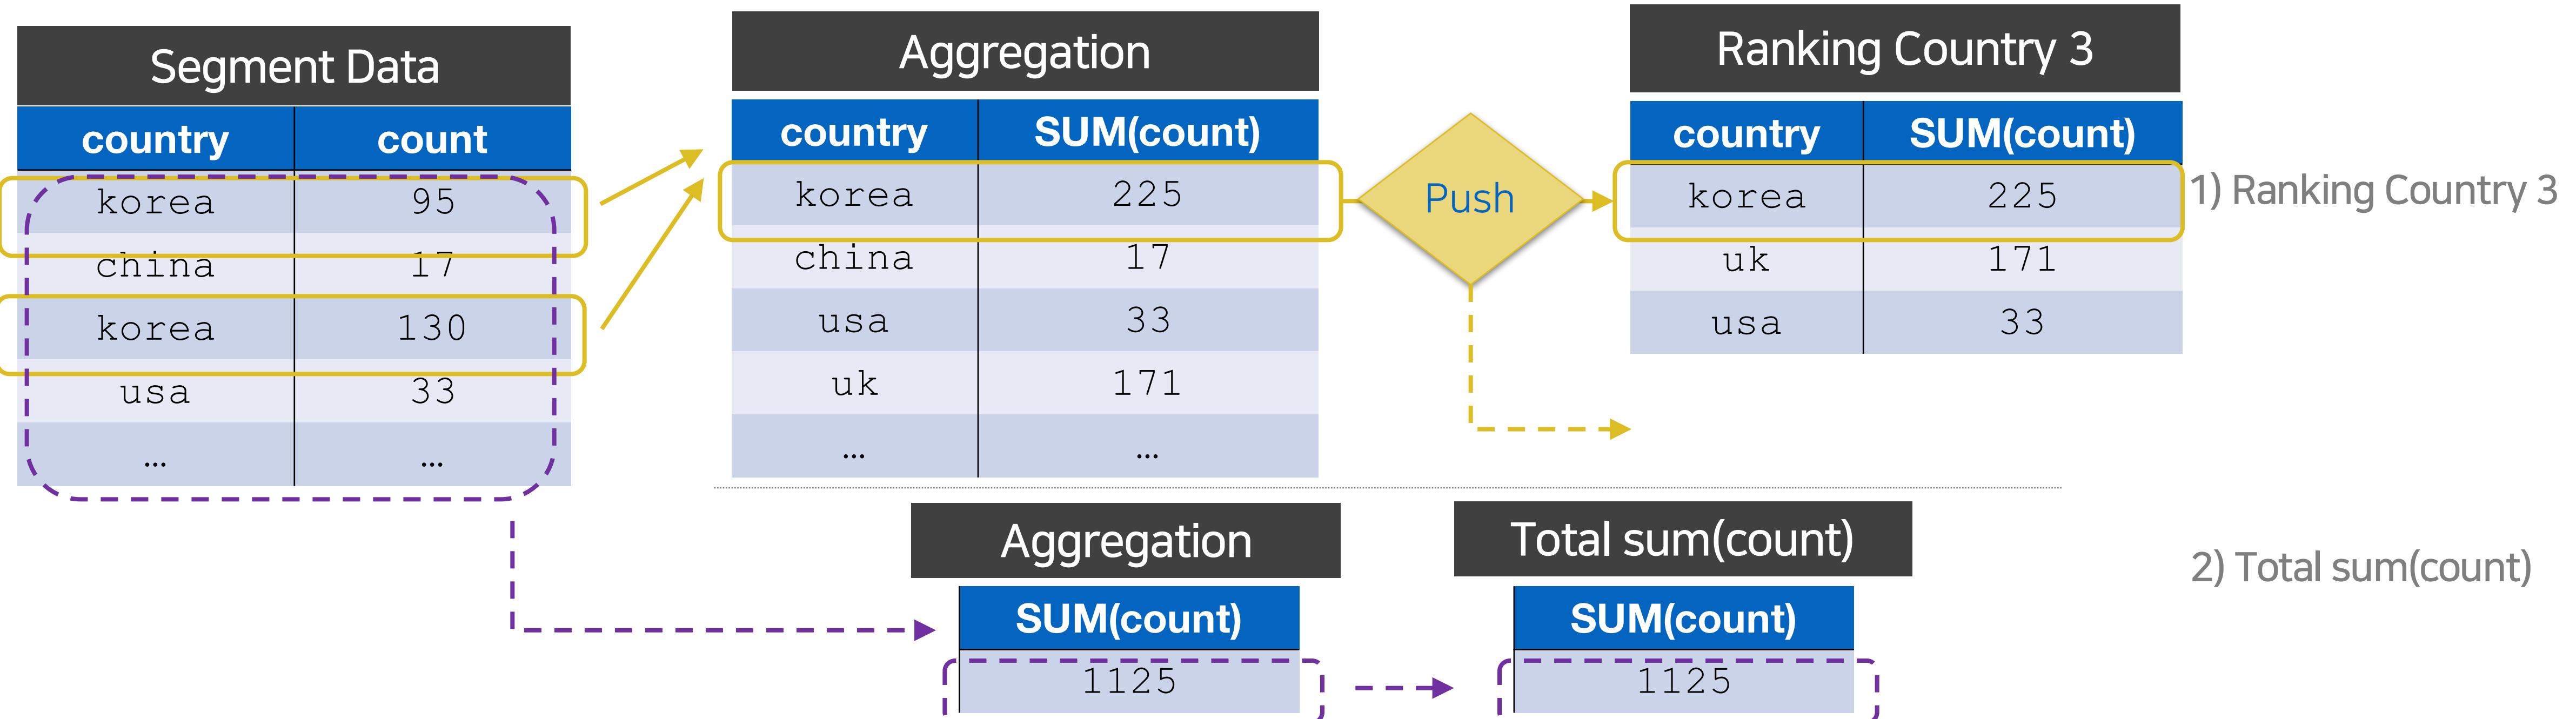
## 4.2 질의성능 향상시키기



## 4.2 질의성능 향상시키기

Country	Sum(count)	Ratio over total count
korea	225	20%
uk	171	15.2%
usa	33	2.9%

처리과정이 비슷한데  
한번에  
처리할 수 없을까?

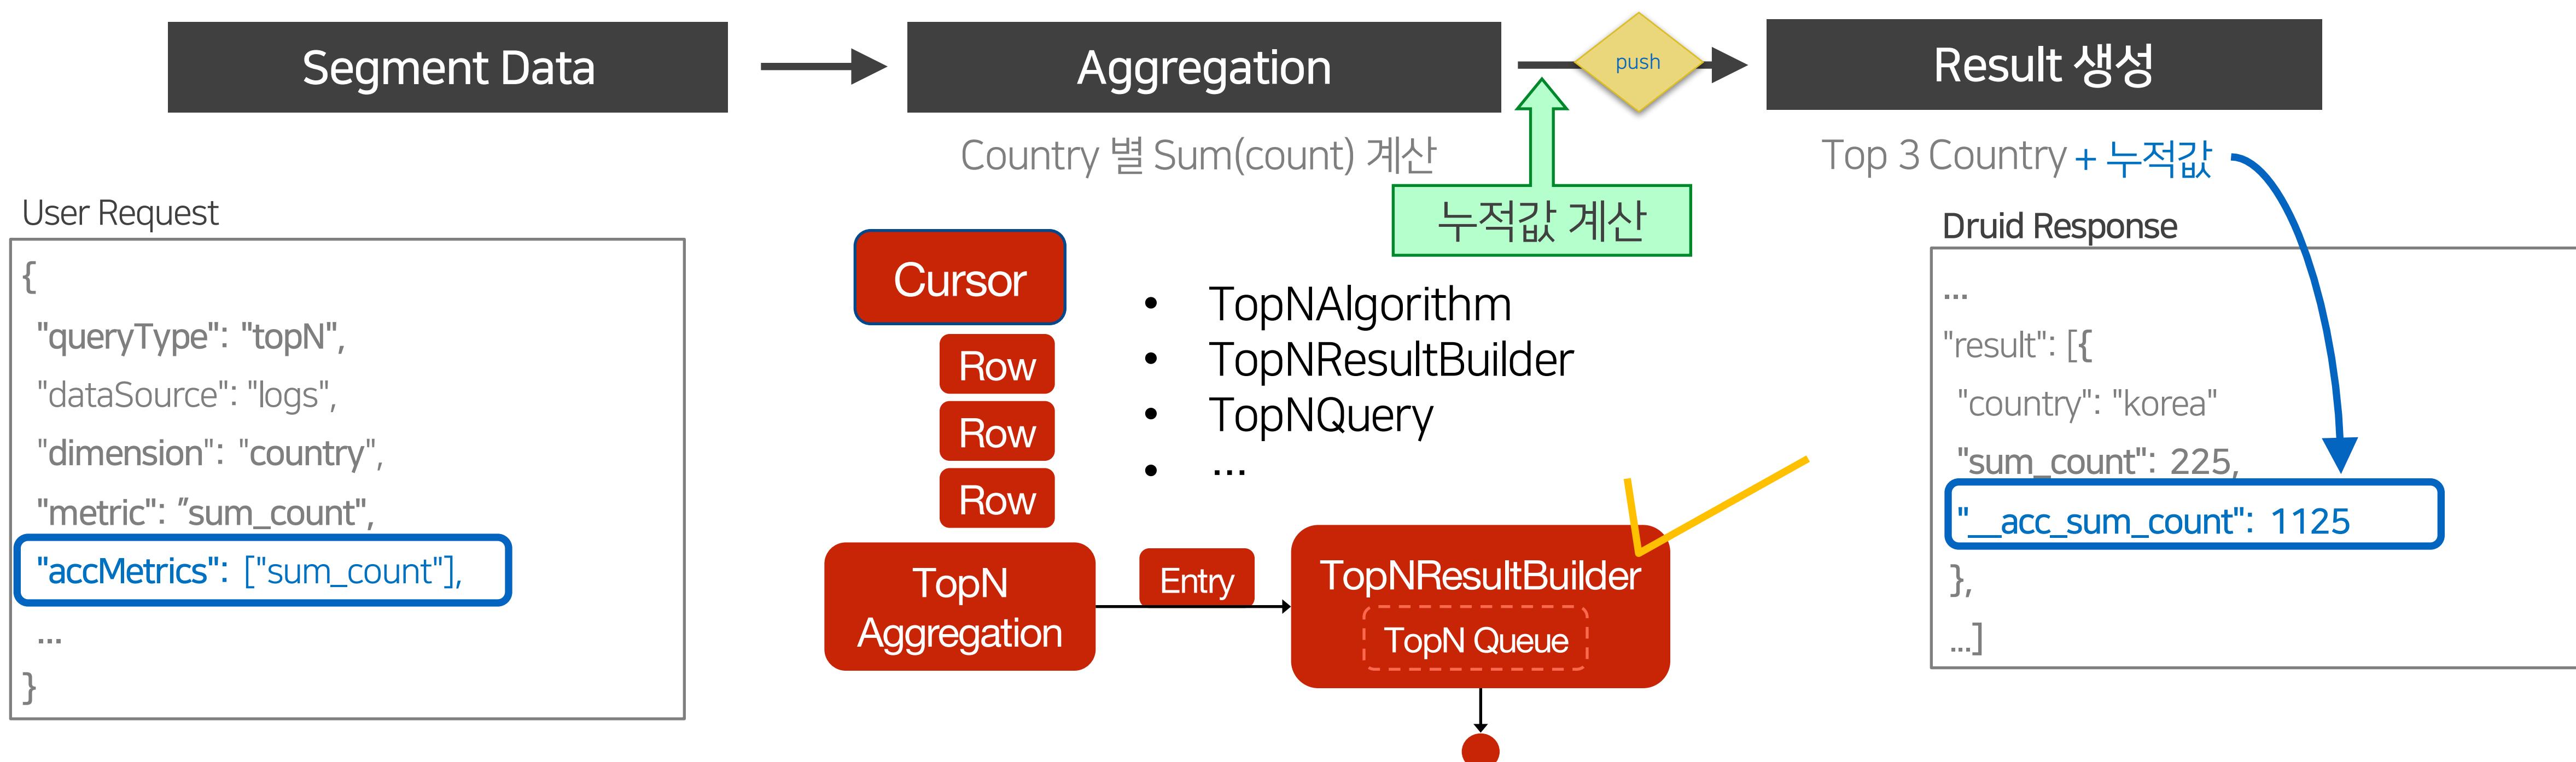


## 4.2 질의성능 향상시키기

### 2) TopN 쿼리 확장

Country	Sum(count)	Ratio over total count
korea	225	20%
uk	171	15.2%
usa	33	2.9%

1) Ranking Country 3  
+ Total sum(count)

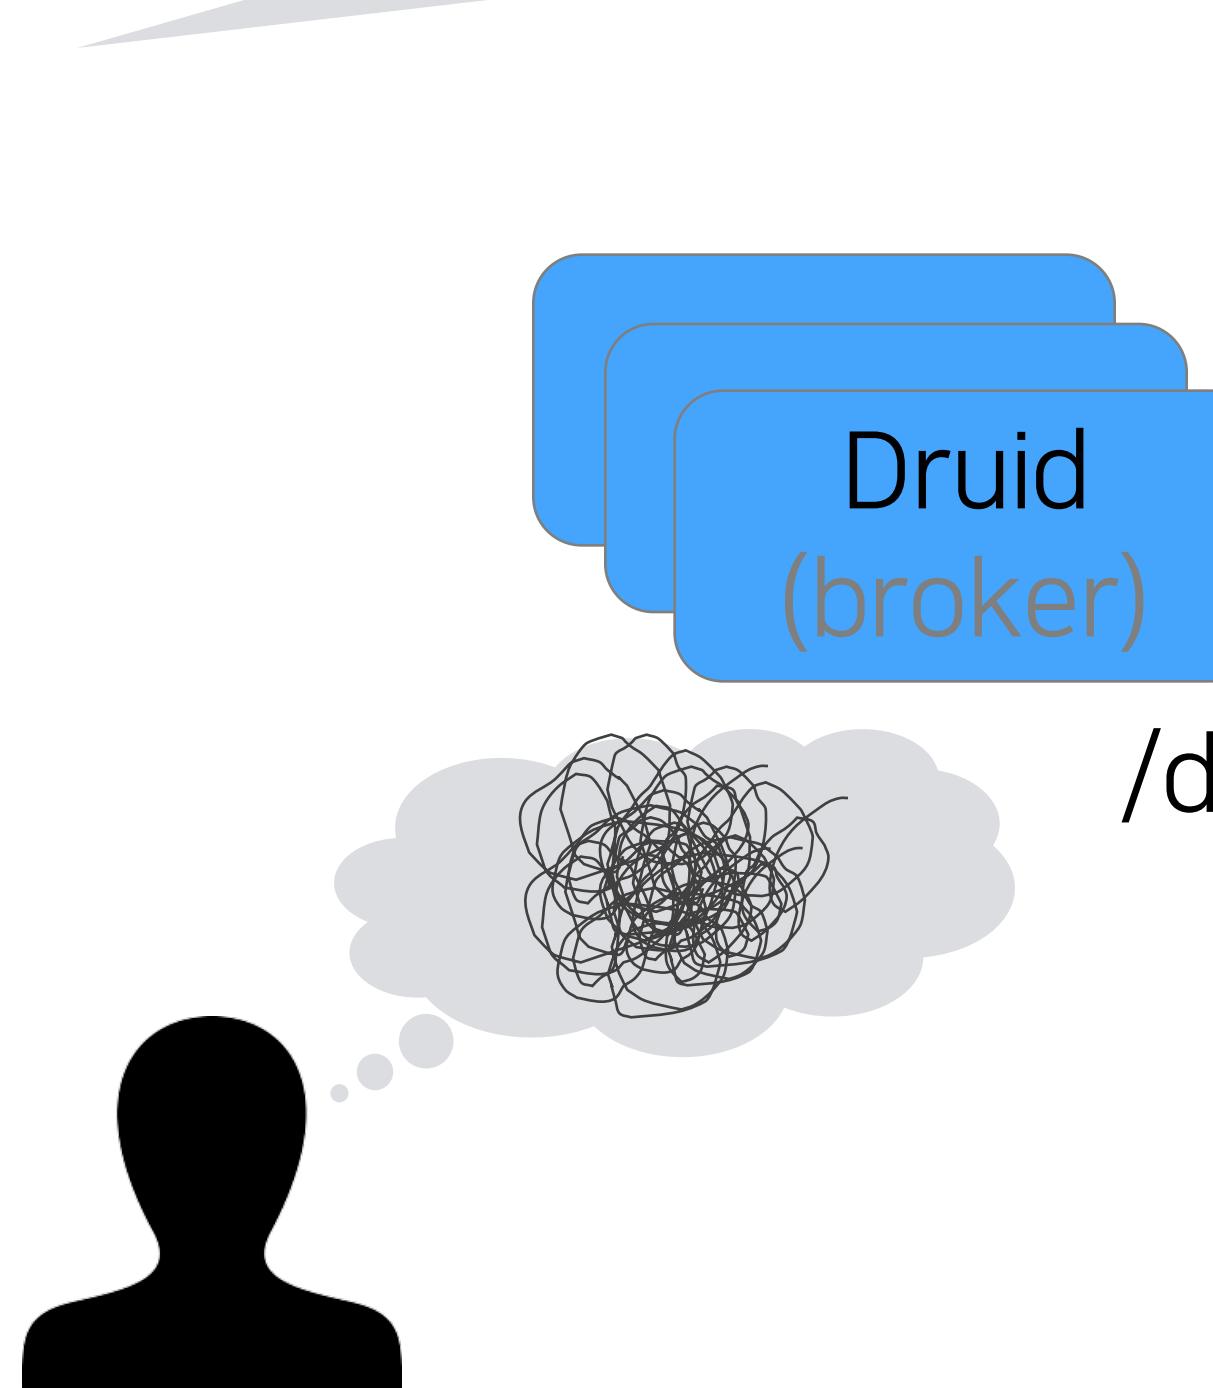


## 4.2 질의성능 향상시키기

### 3) Spark-Druid Connector

다양한 분석을 할 수 있게 해주세요

union, left join, inner join, groupby, having, .....



/druid/v2

/druid/v2/sql

ERROR

```
{  
  "queryType": "topN",  
  "dataSource": "logs",  
  "dimension": "sample_dim",  
  "granularity": "all",  
  "threshold": 5,  
  "aggregation": [...]  
}
```

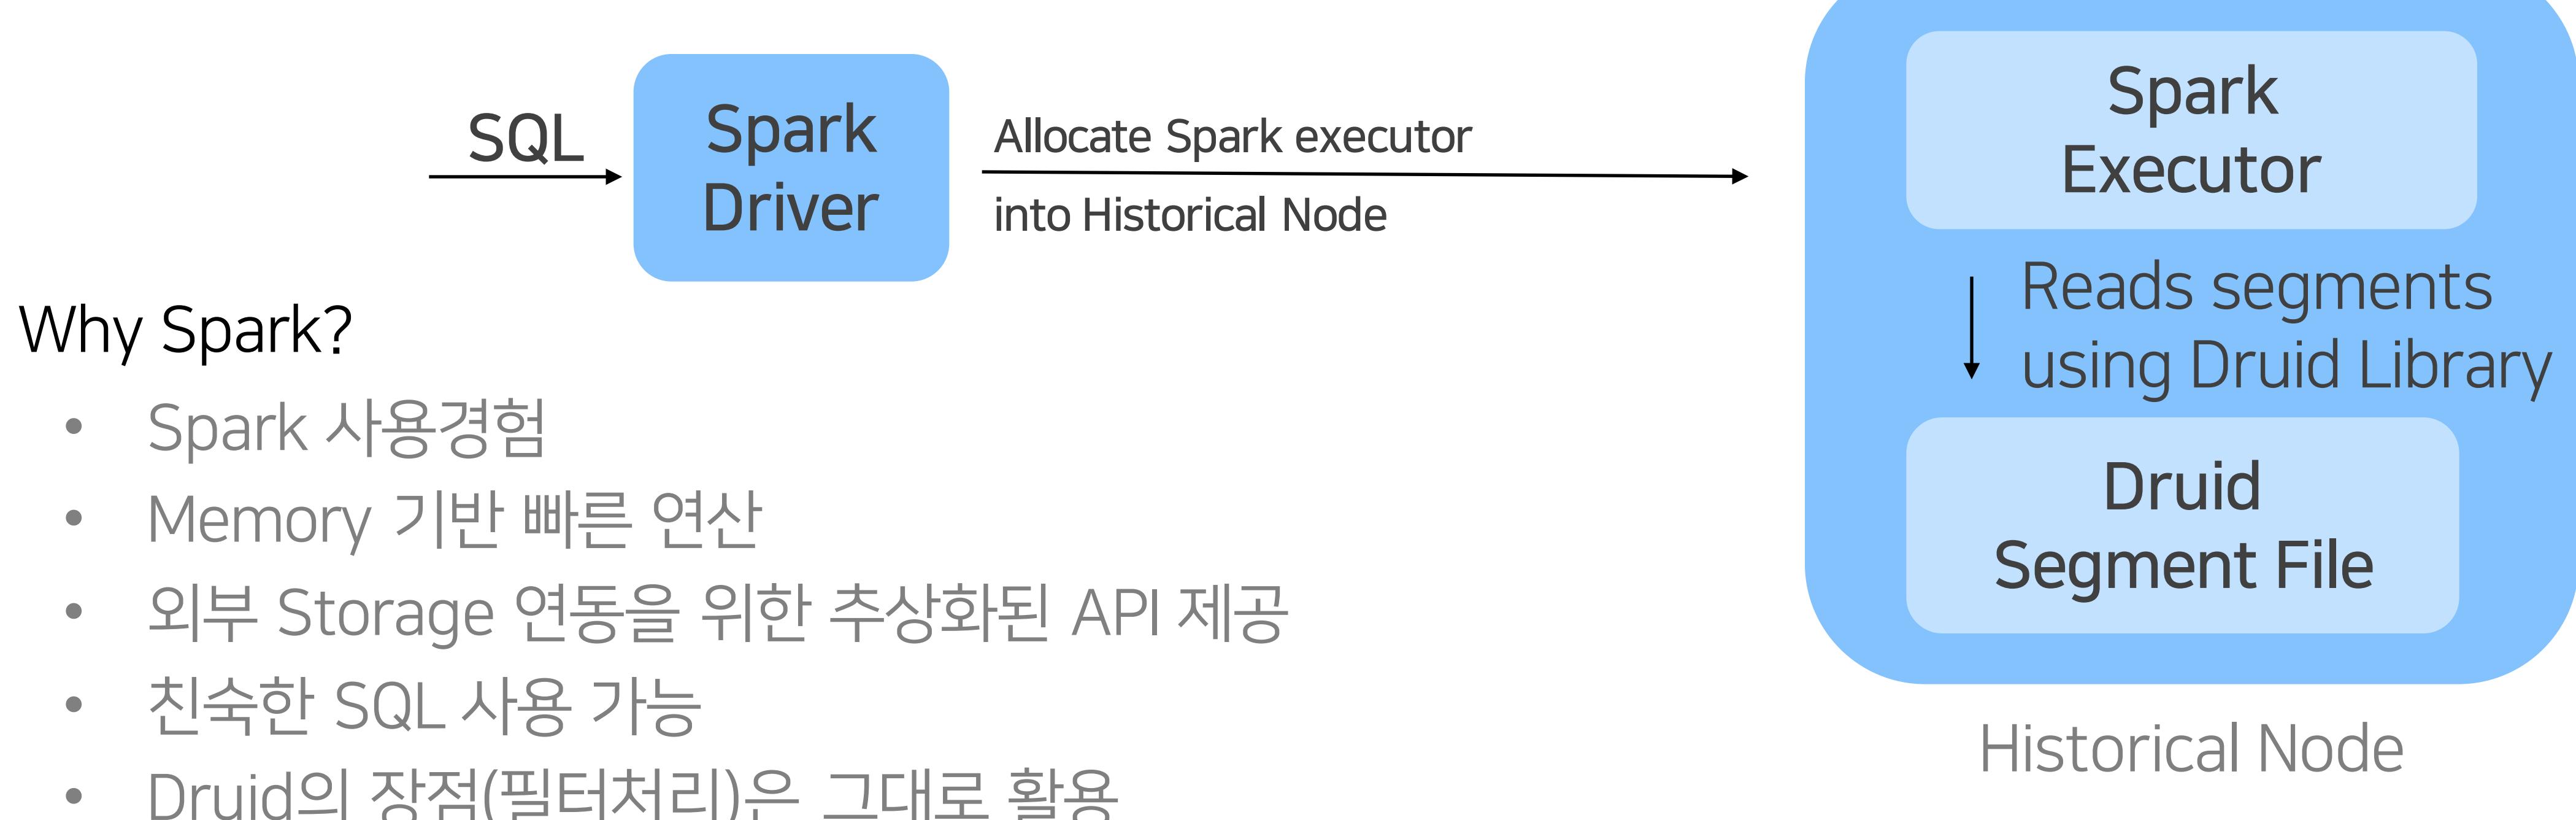
```
{"query": "select dim1 from logs limit 5"}
```

```
{  
  "query":  
    "select dim1 from logs limit 5  
    union all  
    select dim1 from ds2 limit 5"  

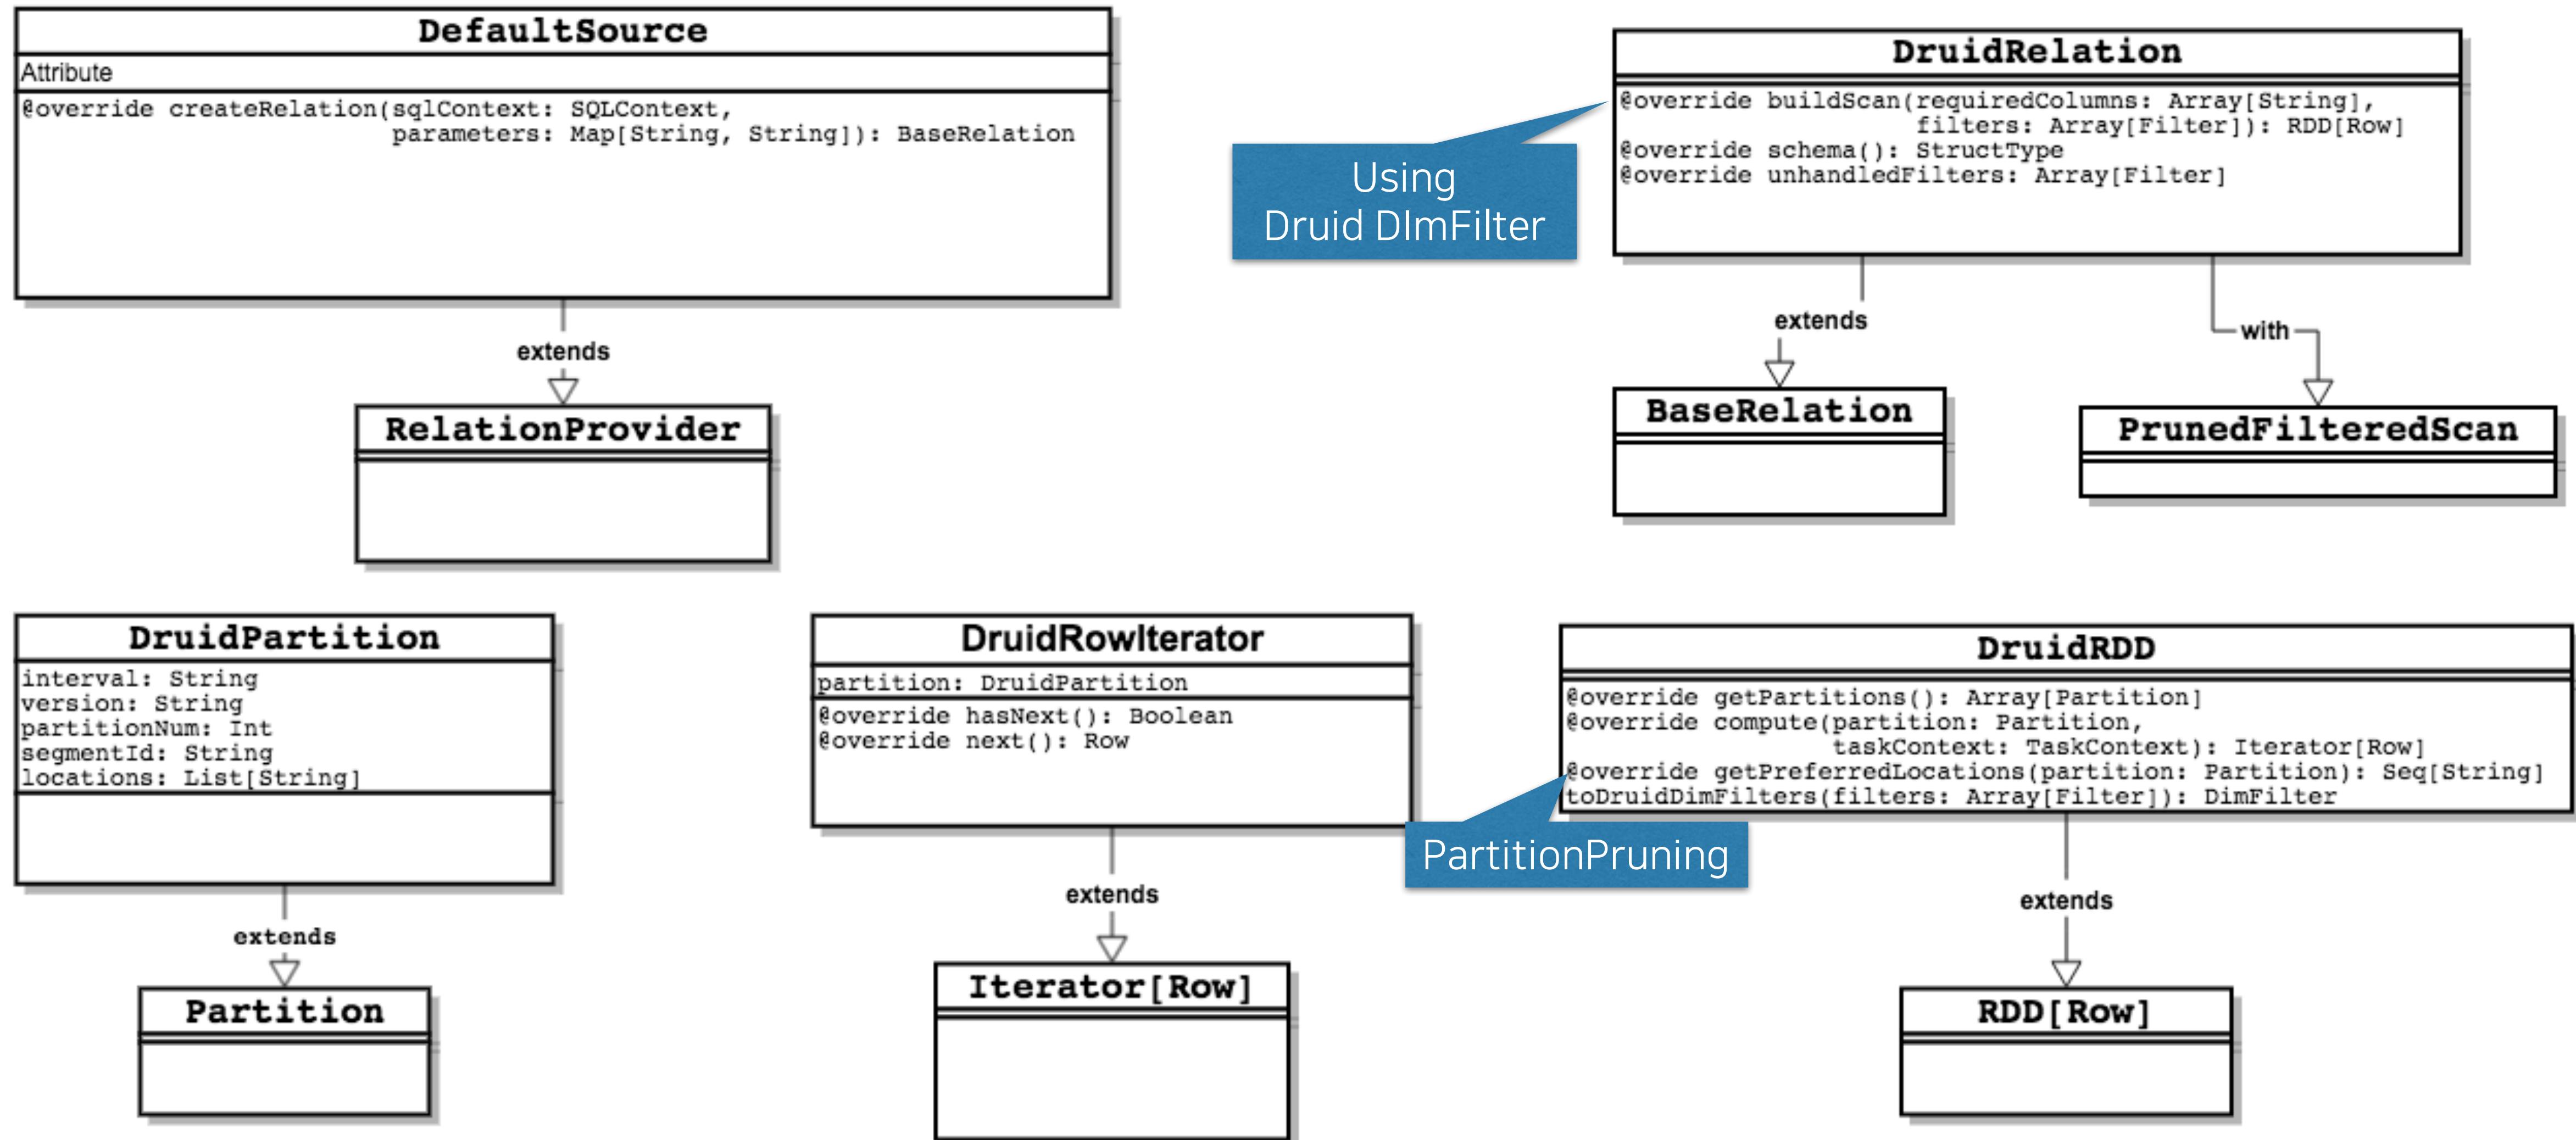
```

## 4.2 질의성능 향상시키기

### 3) Spark-Druid Connector



## 4.2 질의성능 향상시키기



## 4.2 질의성능 향상시키기

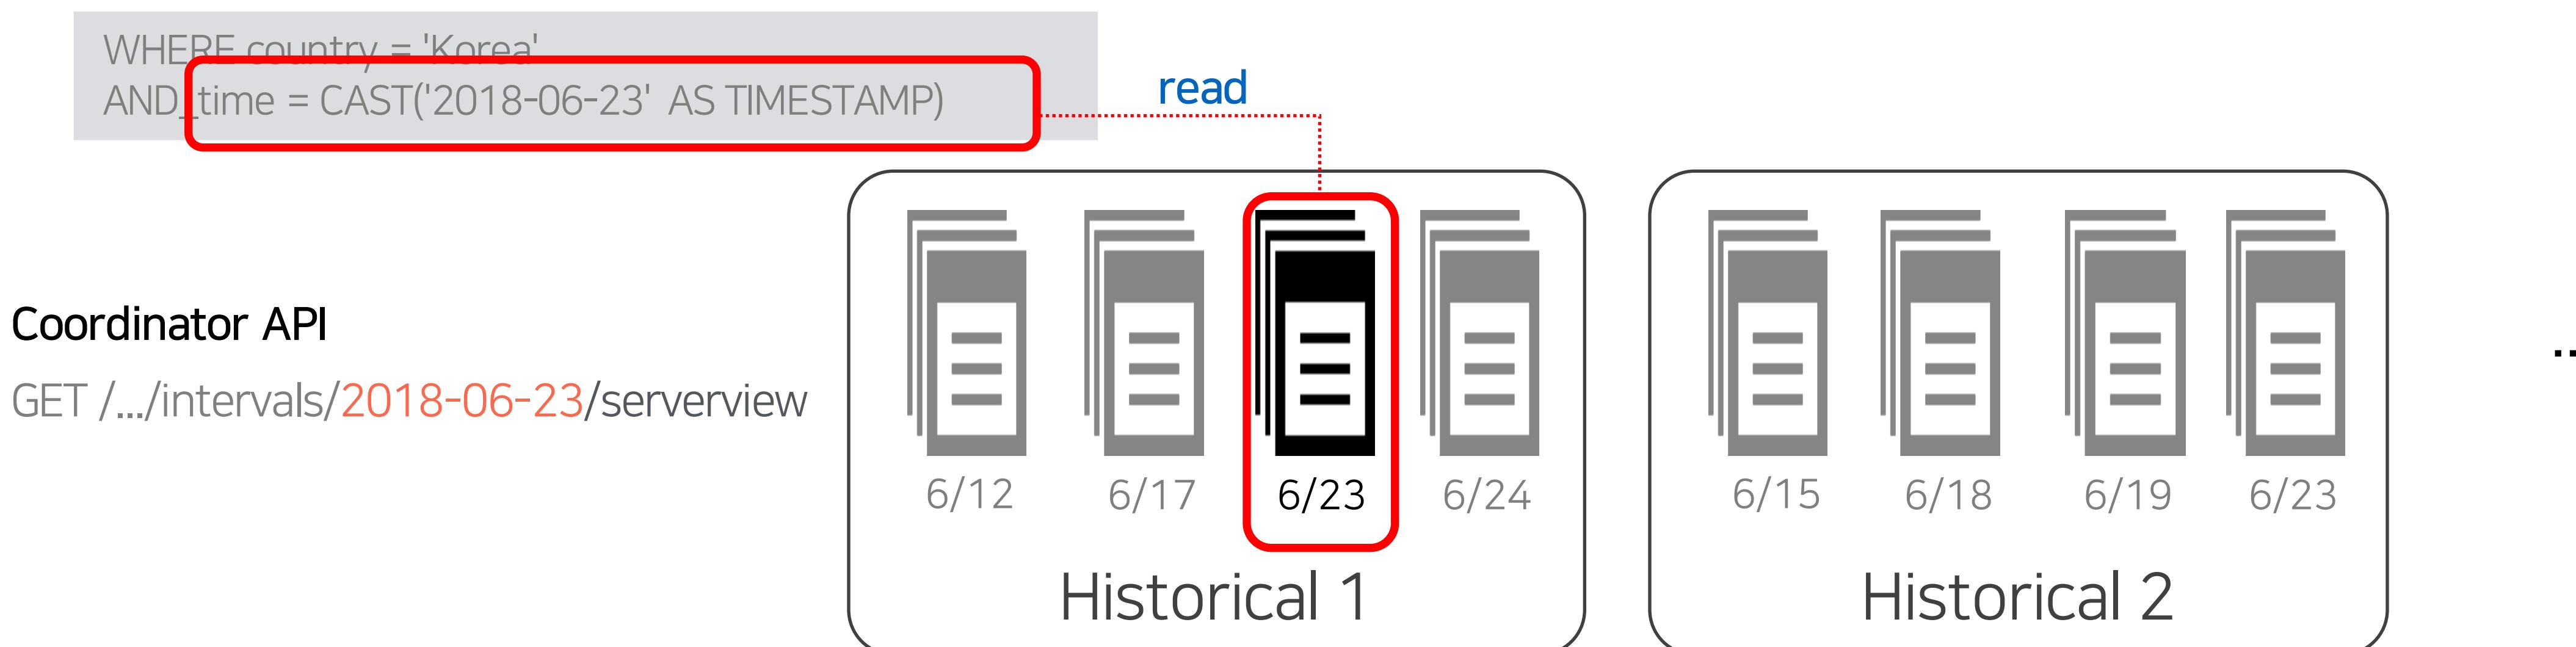
### 3) Spark-Druid Connector

- Locality

Historical node에 download & unzip된 segment file을 read (N/W 절약)

- Segment Partition pruning

SQL의 조건을 분석하여 조회할 데이터가 있는 partition만 read (I/O, Memory 절약)

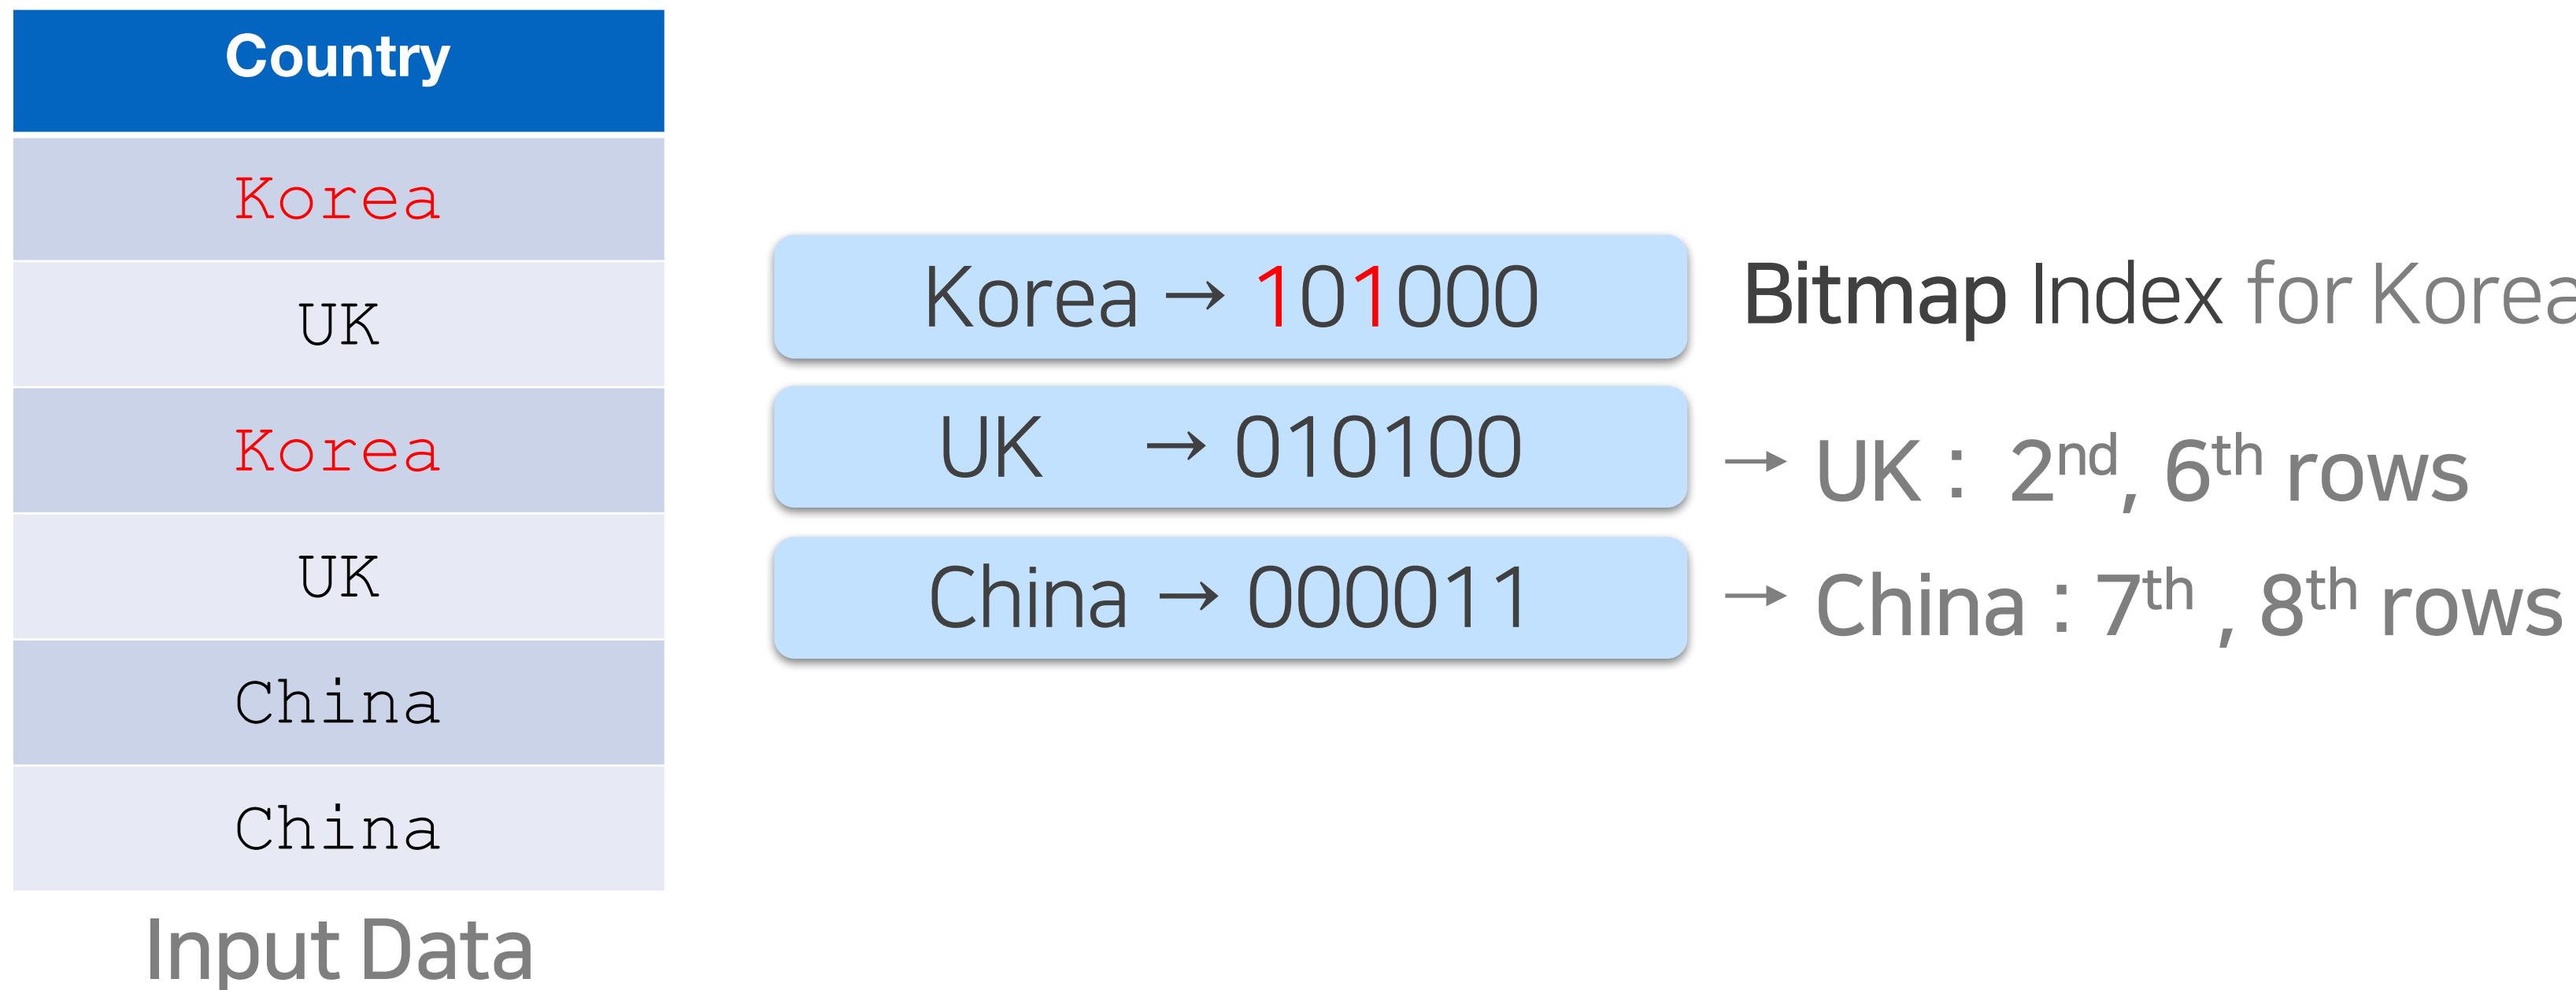


## 4.2 질의성능 향상시키기

### 3) Spark-Druid Connector

- Using Druid Filter

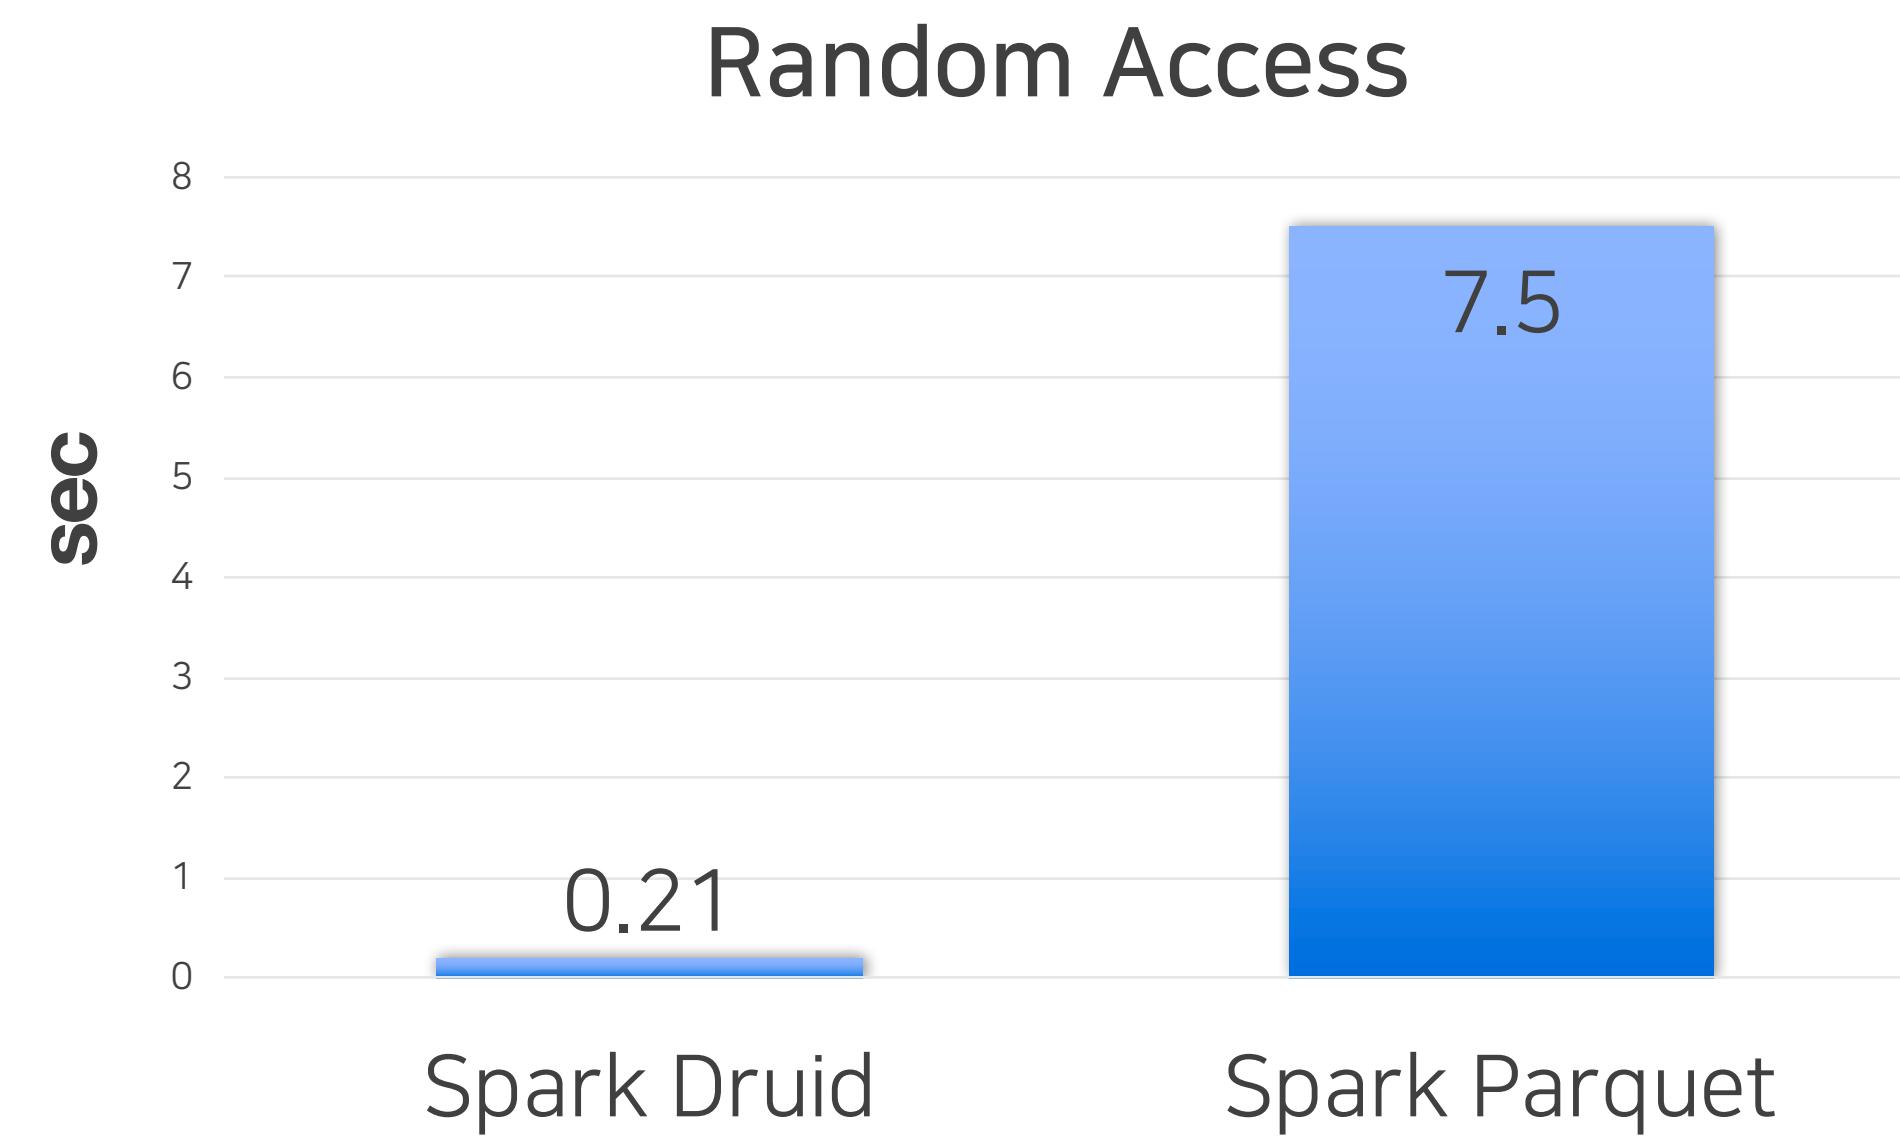
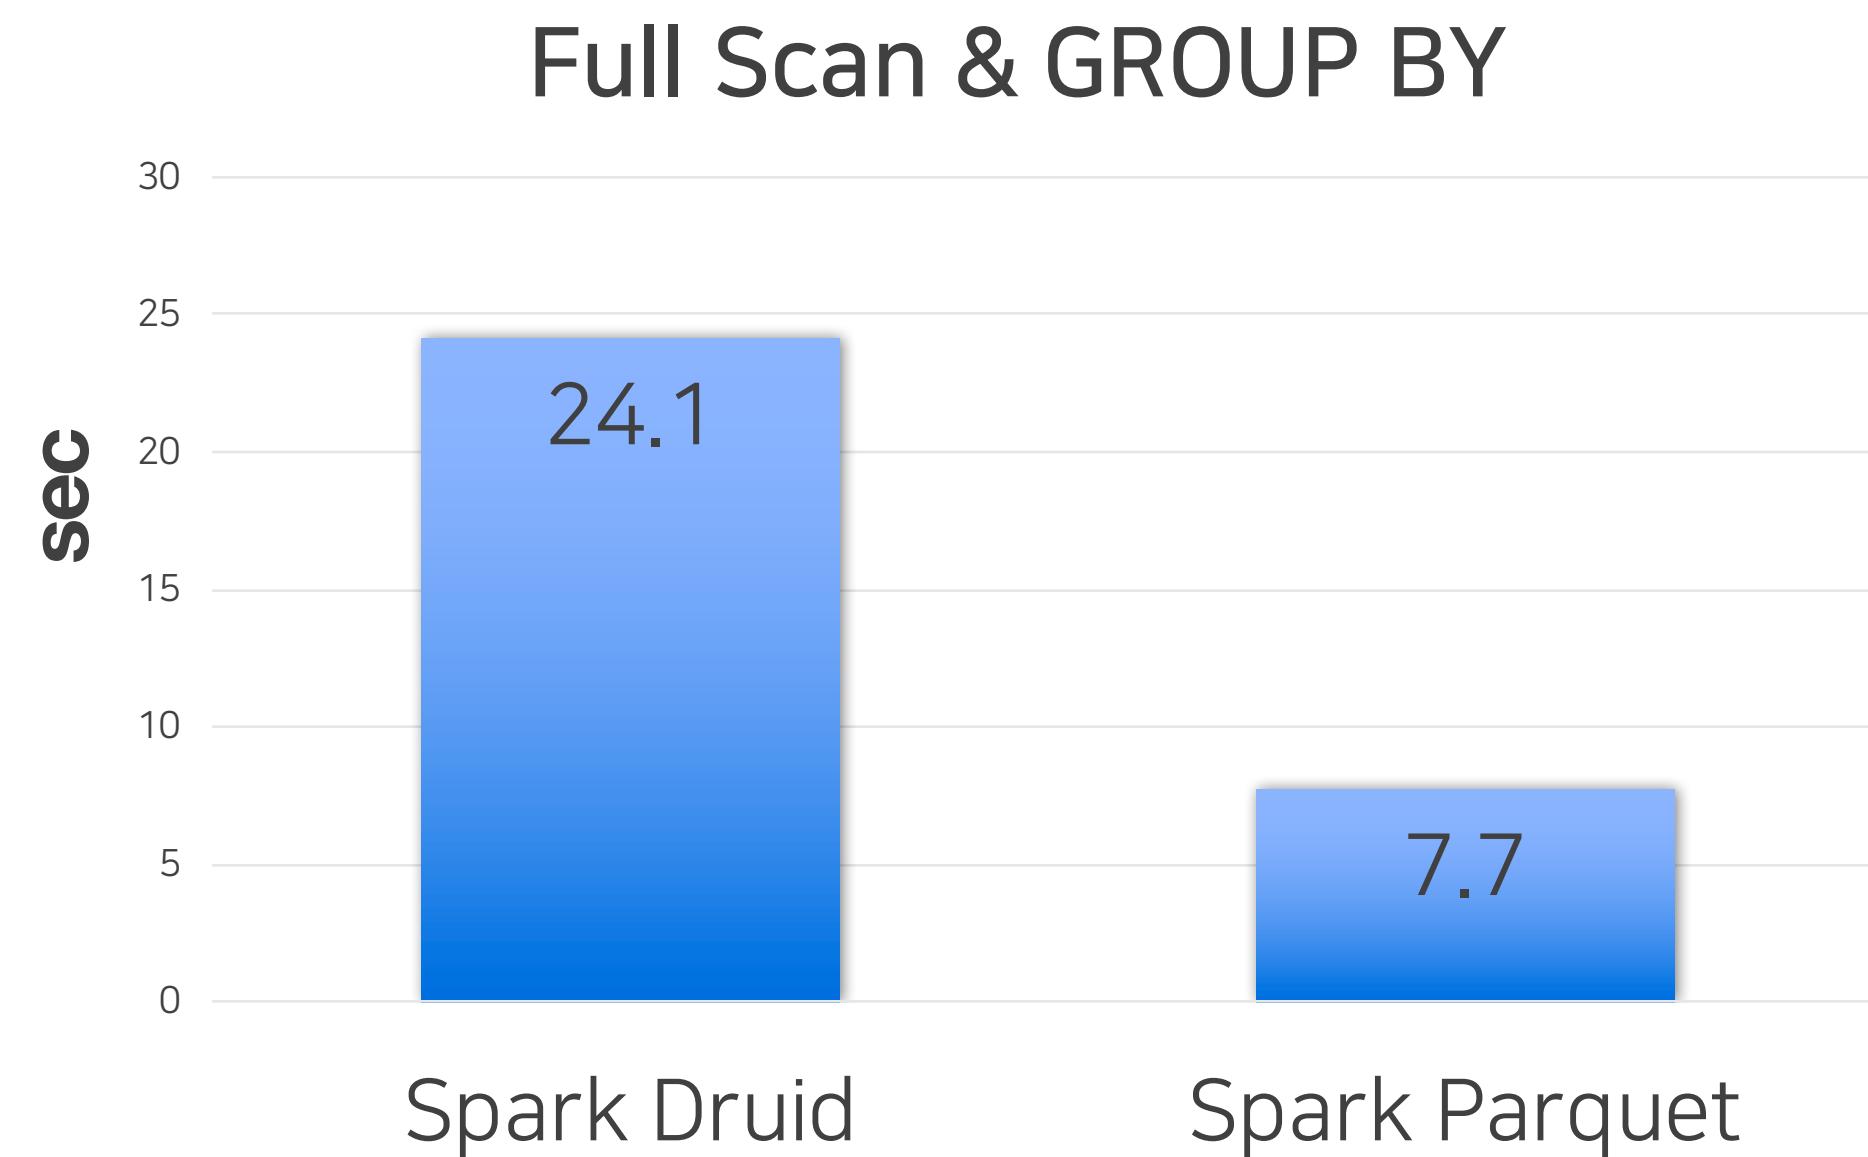
where statement (SQL Filter) -> Spark Filter (load data in memory and filter process)  
-> Druid Filter (load filtered data by druid)



## 4.2 질의성능 향상시키기

### 3) Spark-Druid Connector – 응답시간 비교

(Dataset : 44억건 로그 저장)



# 4.3 저장성능 향상시키기

## 1) Hash Partition 확장

	Ingest Job	특징	아쉬운 점
Single Dimension	<ul style="list-style-type: none"><li>• GroupBy Dim.</li><li>• Dim. Selection</li><li>• Index Generator</li></ul>	<ul style="list-style-type: none"><li>• DimValue 사전순 분포</li><li>• Shard Pruning 가능</li></ul>	<ul style="list-style-type: none"><li>• 중간결과로 인한 Resource 사용 증가</li><li>• Shard 개수 지정 불가</li></ul>
Hash	<ul style="list-style-type: none"><li>• Index Generator</li></ul>	<ul style="list-style-type: none"><li>• 빠른 입수성능</li><li>• 입수 Resource ▼</li></ul>	<ul style="list-style-type: none"><li>• Shard Pruning 불가</li></ul>

Single Dim. Partition처럼  
Hash Partition을  
사용할 수 있지 않을까?

# 4.3 저장성능 향상시키기

## 1) Hash Partition 확장

Ingest

### Index Generator Job

- Shard Spec에 다음정보 전달
  - partitionDimensions
  - hashPrunable

Query

### HashBasedNumberedShardSpec

- getDomain : shard range 반환
- *getHashShardNum* 추가
  - partition value -> shardNum

### HashBasedNumberedShardSpec

- hashPrunable = true
  - partitionDimensions.length = 1
  - && isSingleValueDimension

### DimFilterUtils

- filterShards 설정
  - if (hashPrunable == true)
    - include(*getHashShardNum(v)*)

# 4.3 저장성능 향상시키기

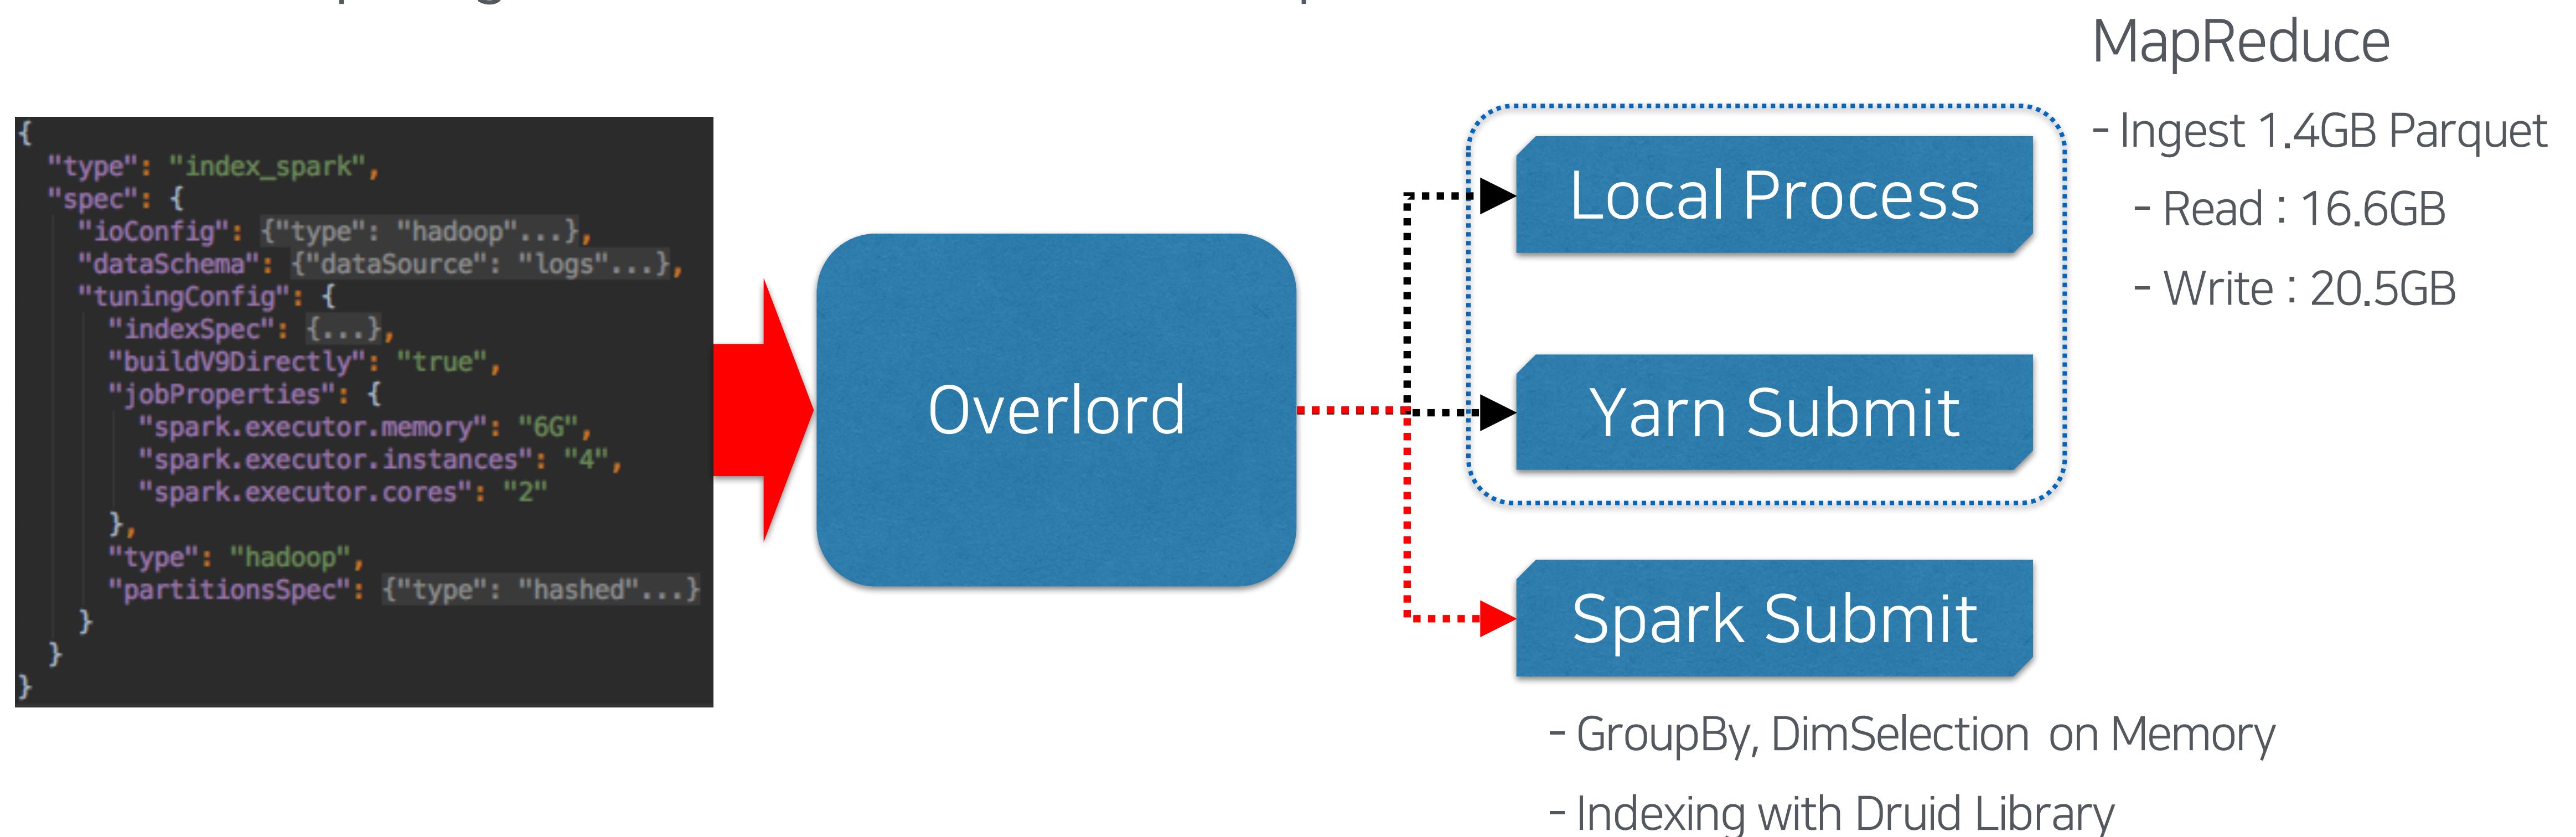
## 1) Hash Partition 확장

	Ingest Job	특징	아쉬운점
Single Dimension	<ul style="list-style-type: none"><li>• GroupBy Dim.</li><li>• Dim Selection</li><li>• Index Generator</li></ul>	<ul style="list-style-type: none"><li>• DimValue 사전순 분포</li><li>• Shard Pruning 가능</li></ul>	<ul style="list-style-type: none"><li>• 중간결과로 인한 Resource 사용 증가</li><li>• Shard 개수 지정 불가</li></ul>
Hash (Partition Dimension 1개)	<ul style="list-style-type: none"><li>• Index Generator</li></ul>	<ul style="list-style-type: none"><li>• 빠른 입수성능</li><li>• 입수 Resource ▼</li><li>• Shard Pruning 가능</li></ul>	<ul style="list-style-type: none"><li>• <del>Shard Pruning 불가</del></li></ul>

## 4.3 저장성능 향상시키기

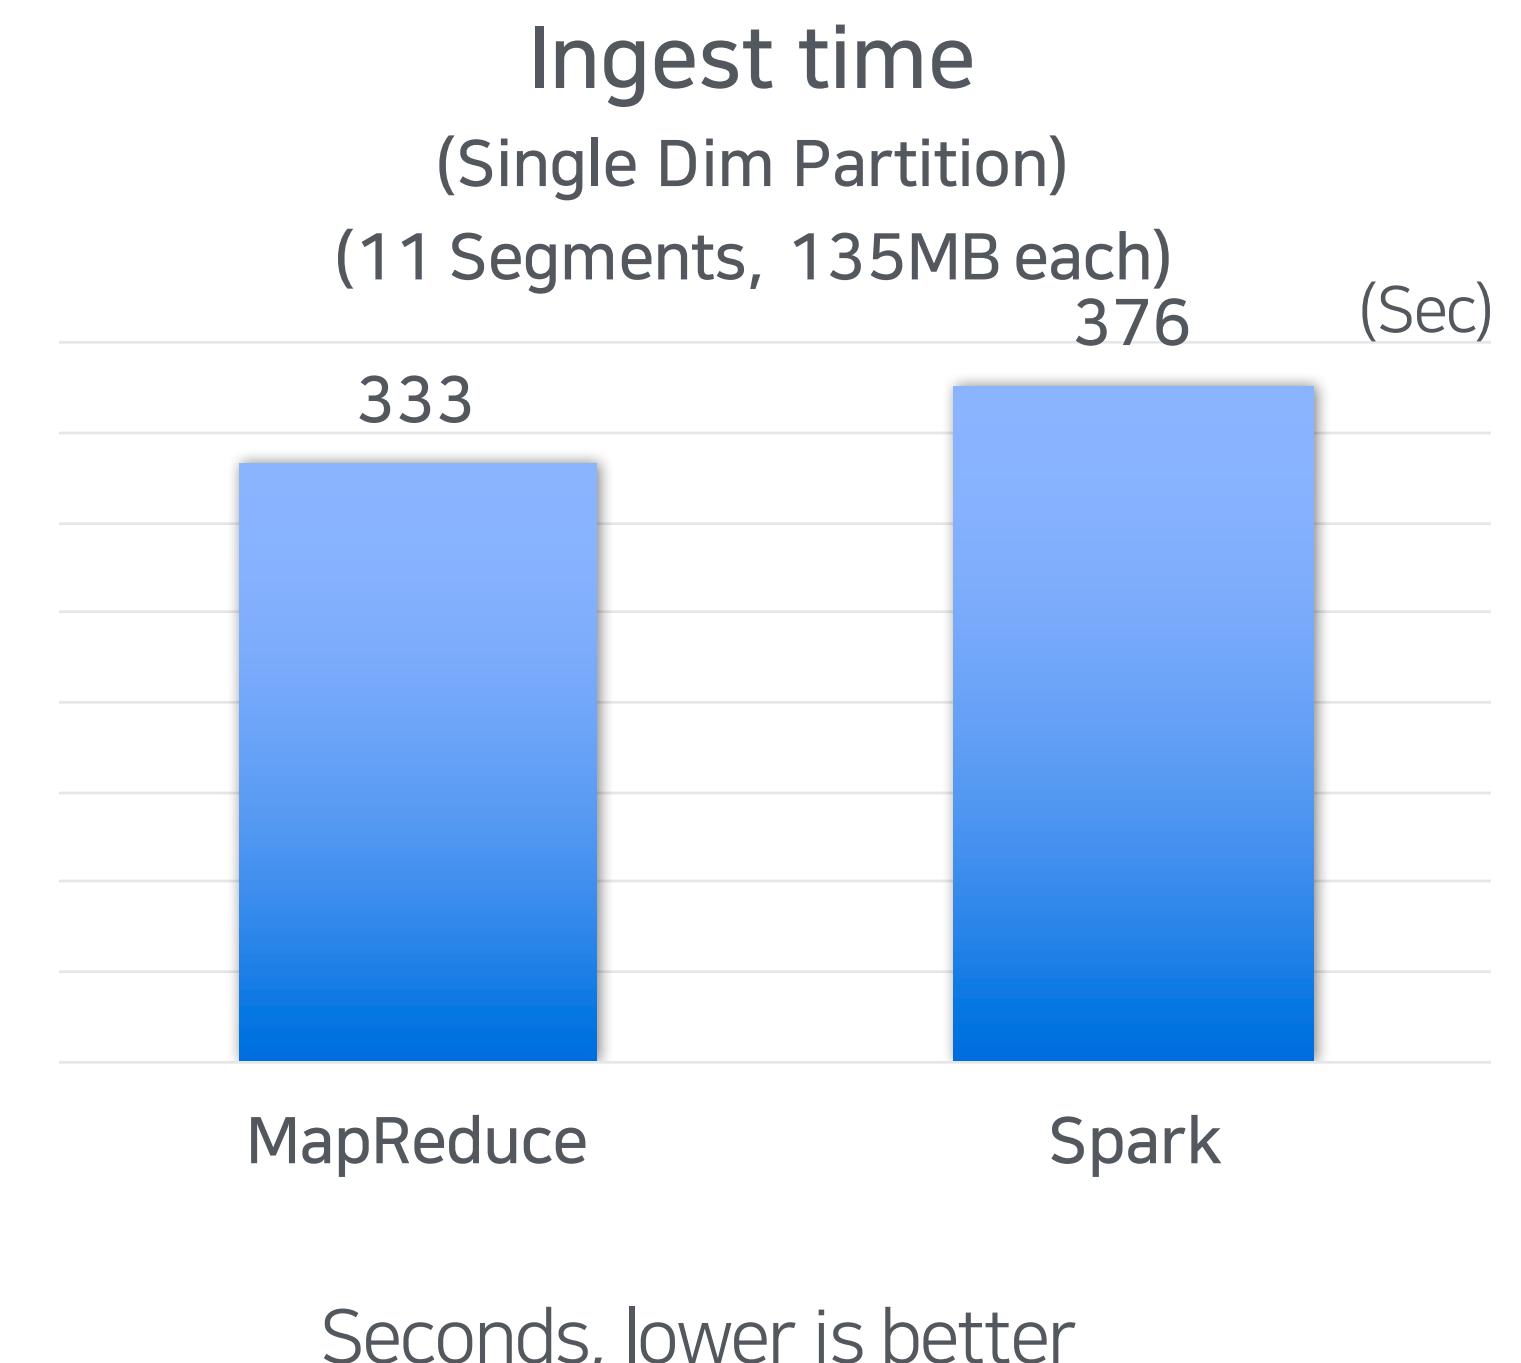
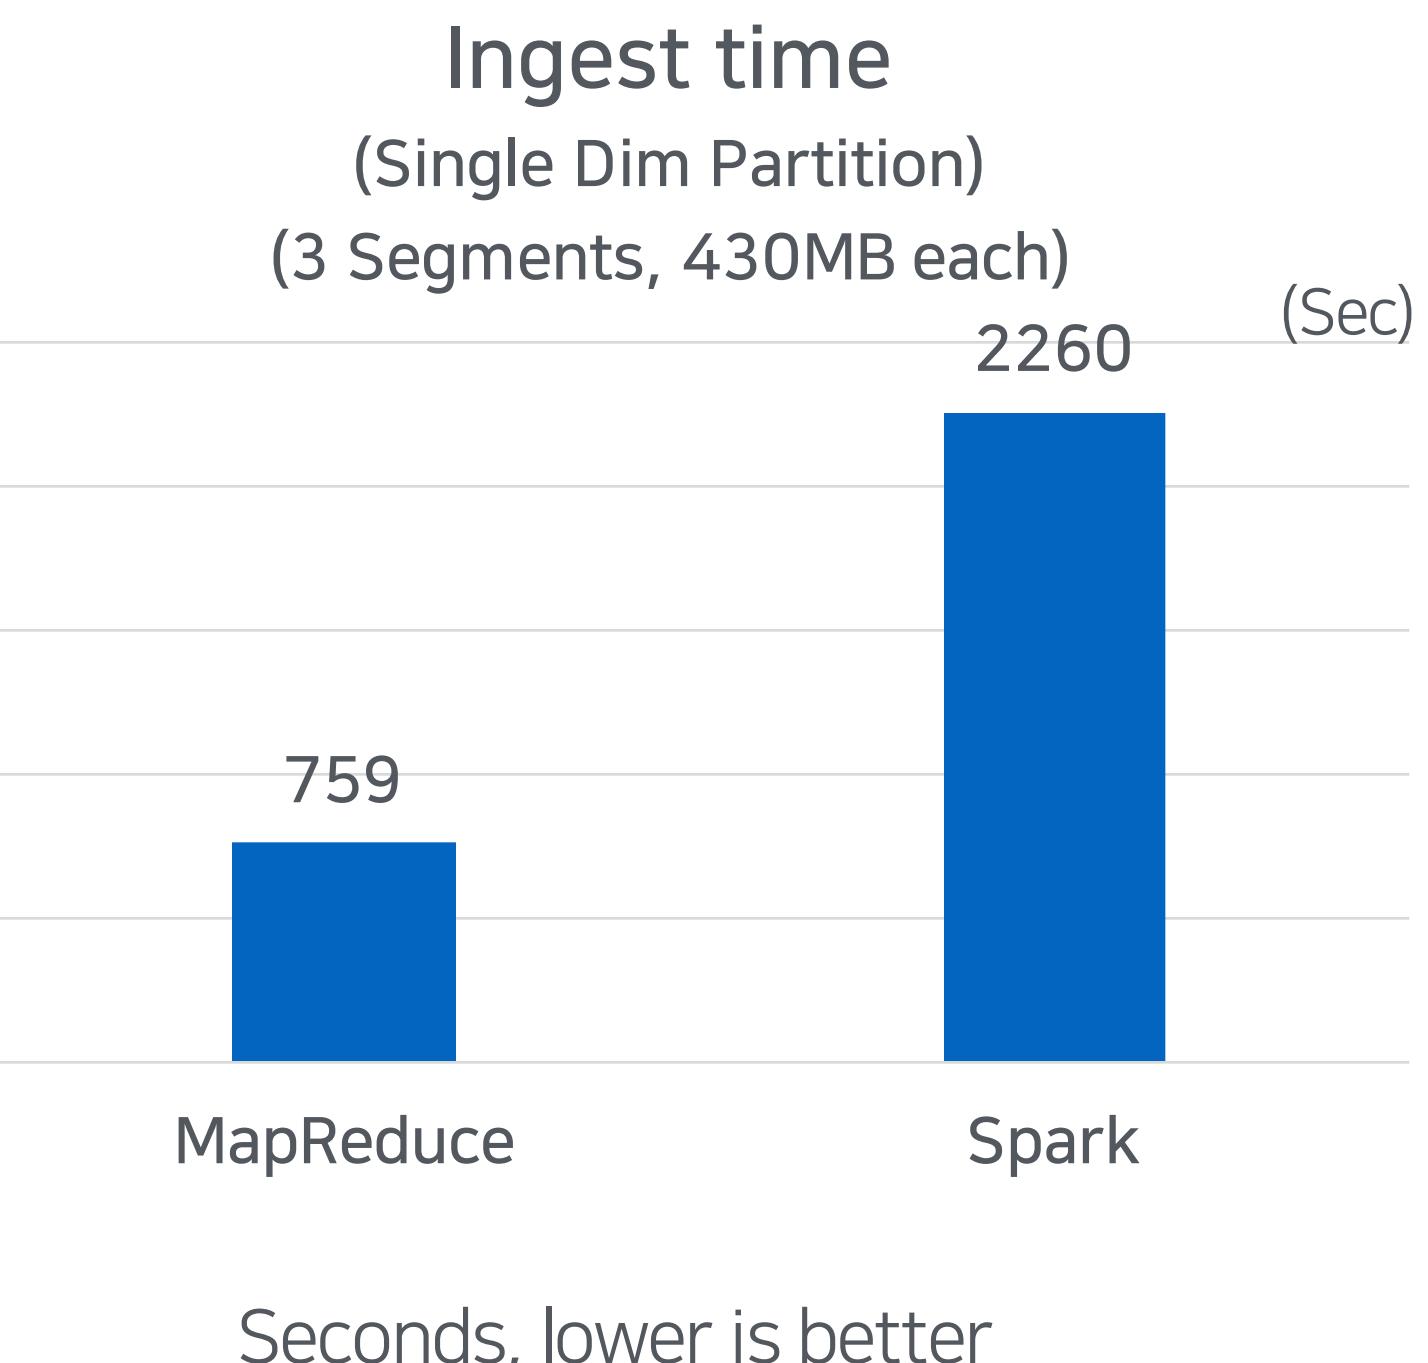
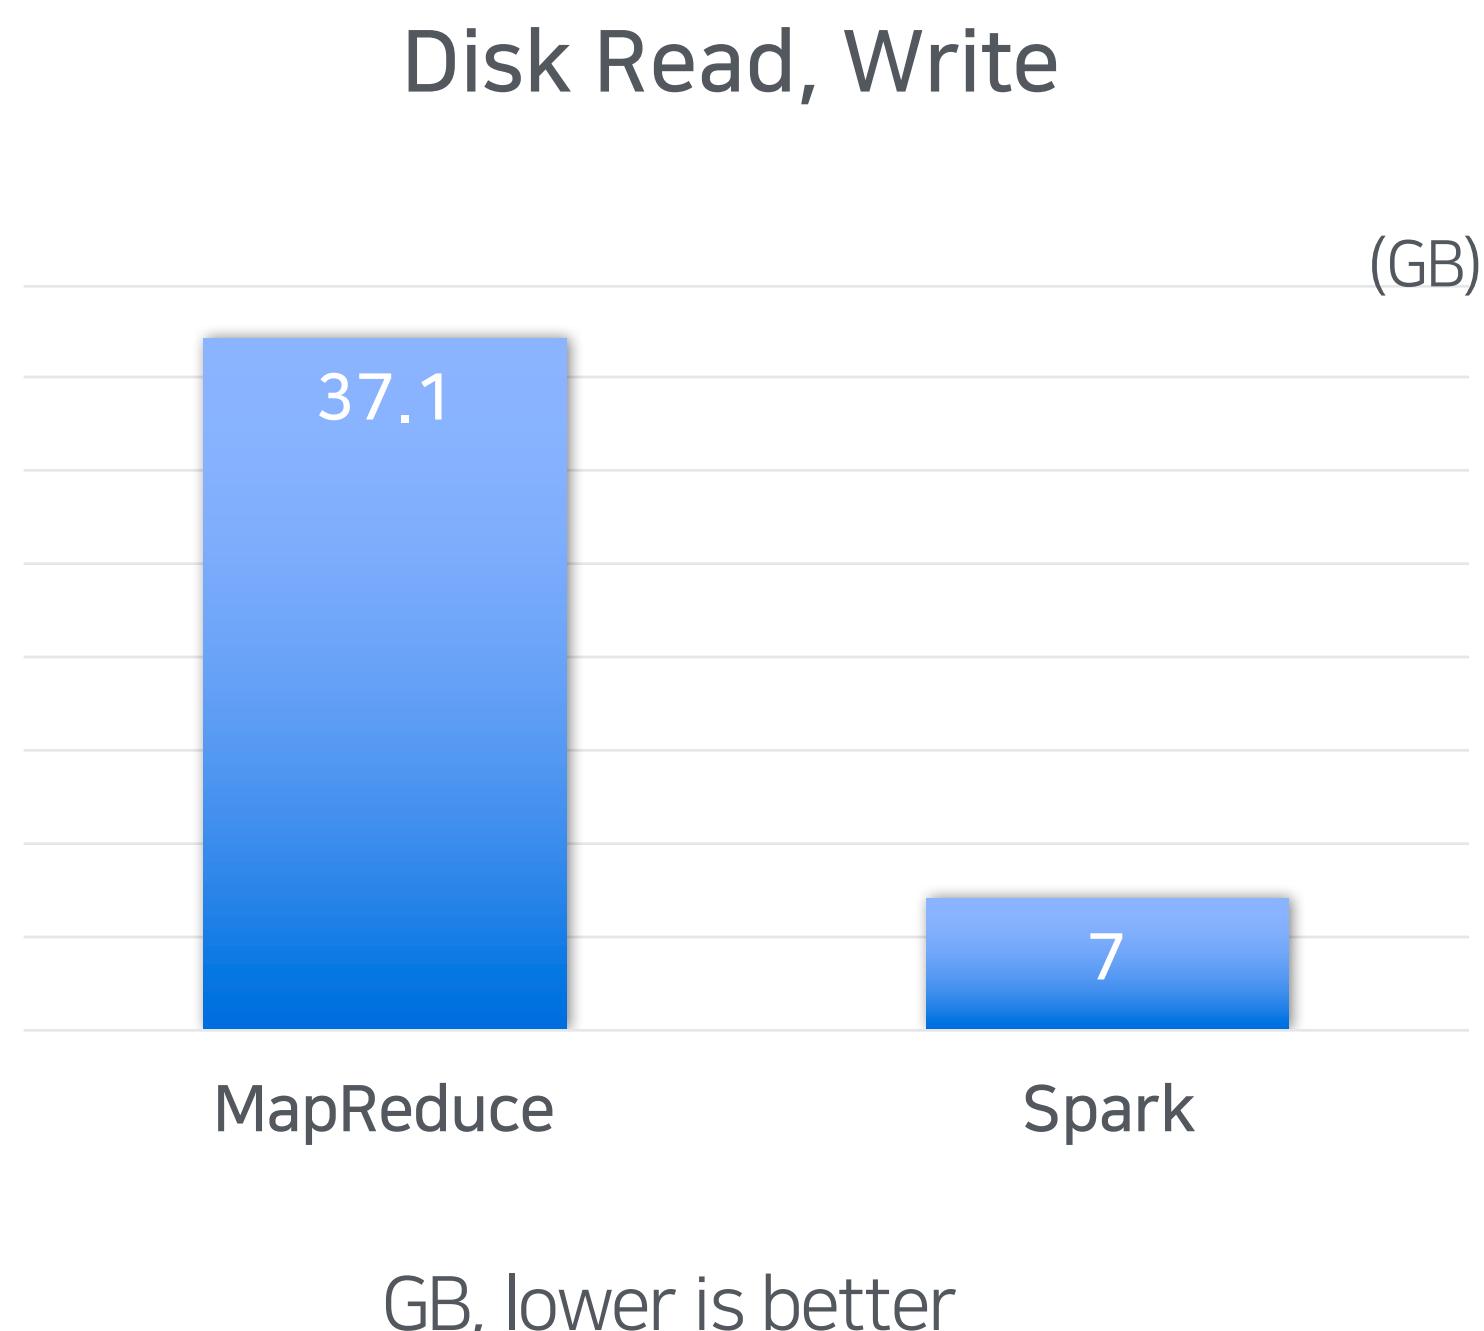
### 2) Druid Spark Batch

- 참고) <https://github.com/metamx/druid-spark-batch>



## 4.3 저장성능 향상시키기

### 2) Druid Spark Batch



## 4.4 편한 운영을 위한 노력

### CDH 연동

- 참고) <https://github.com/knoguchi/cm-druid>

The screenshot shows the Cloudera Manager interface for managing a CDH 5.8.5 cluster. On the left, under the 'Status' tab, there is a list of services: Hosts, Druid-2 (which is highlighted with a red box), HDFS-8, Kafka-6, Spark-8, YARN (MR2 Included), and ZooKeeper-8. On the right, the 'Configuration' page is displayed for the 'Druid-2 on druid' service. The configuration page includes a sidebar with filters for Scope, Category, and Status, and a main area for viewing and editing various configuration properties.

Property	Value
Local Persistent Directory	Druid-2 (Service-Wide) druid.persistent.dir /var/lib/druid
Druid Extensions	Druid-2 (Service-Wide) druid.extensions.loadList ["druid-histogram", "druid-datasources", "druid-lookups-cached-global", "mysql-metadata"]
Druid Startup Logging	Druid-2 (Service-Wide) <input checked="" type="checkbox"/> Properties druid.startup.logging.logProperties
Druid Zookeeper Service Host	Druid-2 (Service-Wide) <input checked="" type="checkbox"/> druid.zk.service.host XXXX
Druid Zookeeper Paths Base	Druid-2 (Service-Wide) druid.zk.paths.base /druid
Druid Metadata Storage Type	Druid-2 (Service-Wide) druid.metadata.storage.type mysql

# Q & A

질문은 Slido에 남겨주세요.

sli.do

#devview

TRACK 1

# Appendix

# Appendix

- About team
- Comparison of Data Size
- Druid Configuration
- Ingest Spec for External Resource (Yarn, HDFS)
- How Kafka Indexing Service Works
- Druid Query Flow

# About team



- [네이버 콘텐츠 통계서비스 소개 및 구현 경험 공유](#)
- [빅데이터 다차원 분석 플랫폼, Kylin](#)
- Web Analytics at Scale with Elasticsearch @ naver.com
  - [Part1](#)
  - [Part2](#)
- [Kudu를 이용한 빅데이터 다차원 분석 시스템 개발](#)
- [Web analytics at scale with Druid at Naver \(Strata 2018\)](#)

# Comparison of Data Size

## Primary data size

Parquet	Apache Druid	Elasticsearch	Apache Kudu
7GB	34GB	50GB	14GB

# Druid Configuration

## Druid Deployment

- 10 Broker Nodes
- 40 Historical Nodes
- 2 Overlord & MiddleManager Nodes
- 2 Coordinator Nodes
- 10 Yarn & HDFS Nodes for Batch Ingestion
- Spark Standalone Cluster runs on Historical Nodes for Locality

# Druid Configuration

## Druid Deployment

- Version : 0.11.0
- H/W Spec for Broker & Historical
  - CPU : 40 cores (w/ hyperthread)
  - RAM : 128GB
  - HDD : SSD w/ RAID 5

# Druid Configuration

## Memory Configuration

Configuration	Broker	Historical
-Xmx	20GB	12GB
-XX:MaxDirectMemorySize	30GB	45GB
druid.processing.numMergeBuffers	10	20
druid.processing.numThreads	20	30
druid.processing.buffer.sizeBytes	512MB	800MB
druid.cache.sizeInBytes	0	5GB
druid.server.http.numThreads	40	40

# Ingest Spec for External Resource

## Use Yarn External Resource for Batch Ingestion

```
"tuningConfig": {  
    "type": "hadoop",  
    "jobProperties": {  
        "yarn.resourcemanager.hostname" : "host1.com",  
        "yarn.resourcemanager.address" : "host1.com:8032",  
        "yarn.resourcemanager.scheduler.address": "host1.com:8030",  
        "yarn.resourcemanager.webapp.address": "host1.com:8088",  
        "yarn.resourcemanager.resource-tracker.address": "host1.com:8031",  
        "yarn.resourcemanager.admin.address": "host1.com:8033"  
    }  
}
```

# Ingest Spec for External Resource

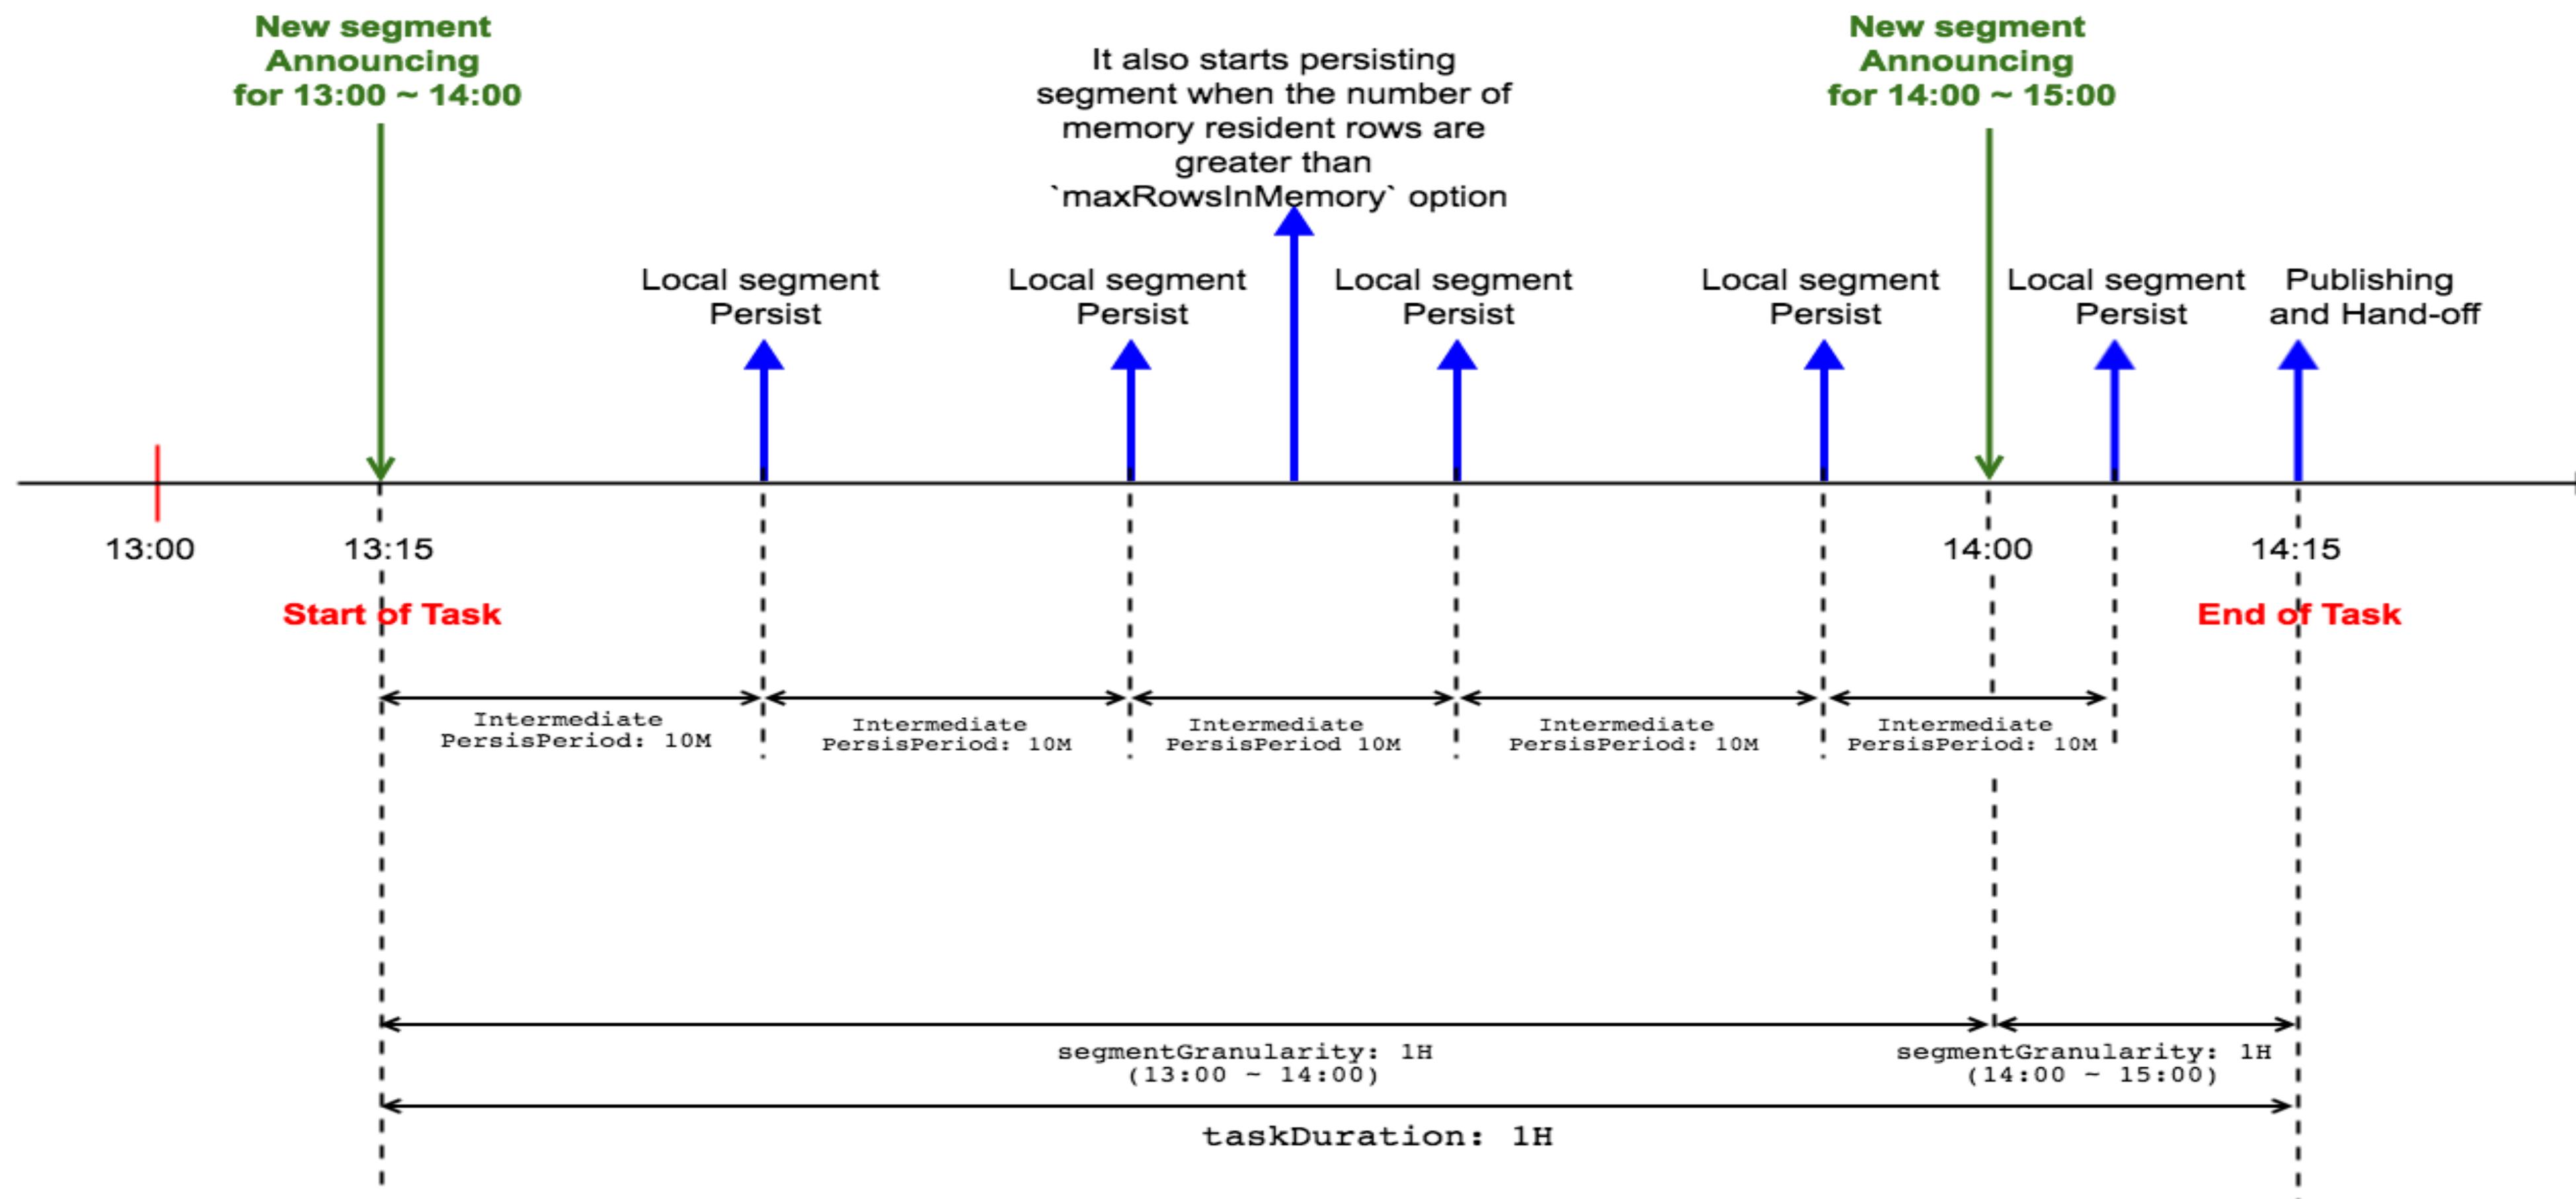
Use External HDFS for intermediate M/R output

```
"tuningConfig": {  
    "type": "hadoop",  
    "jobProperties": {  
        "fs.defaultFS": "hdfs://DEFAULT_FS:8020",  
        "dfs.namenode.http-address": "NAMENODE:50070",  
        "dfs.namenode.https-address": "NAMENODE:50470",  
        "dfs.namenode.servicerpc-address": "NAMENODE:8022"  
    }  
}
```

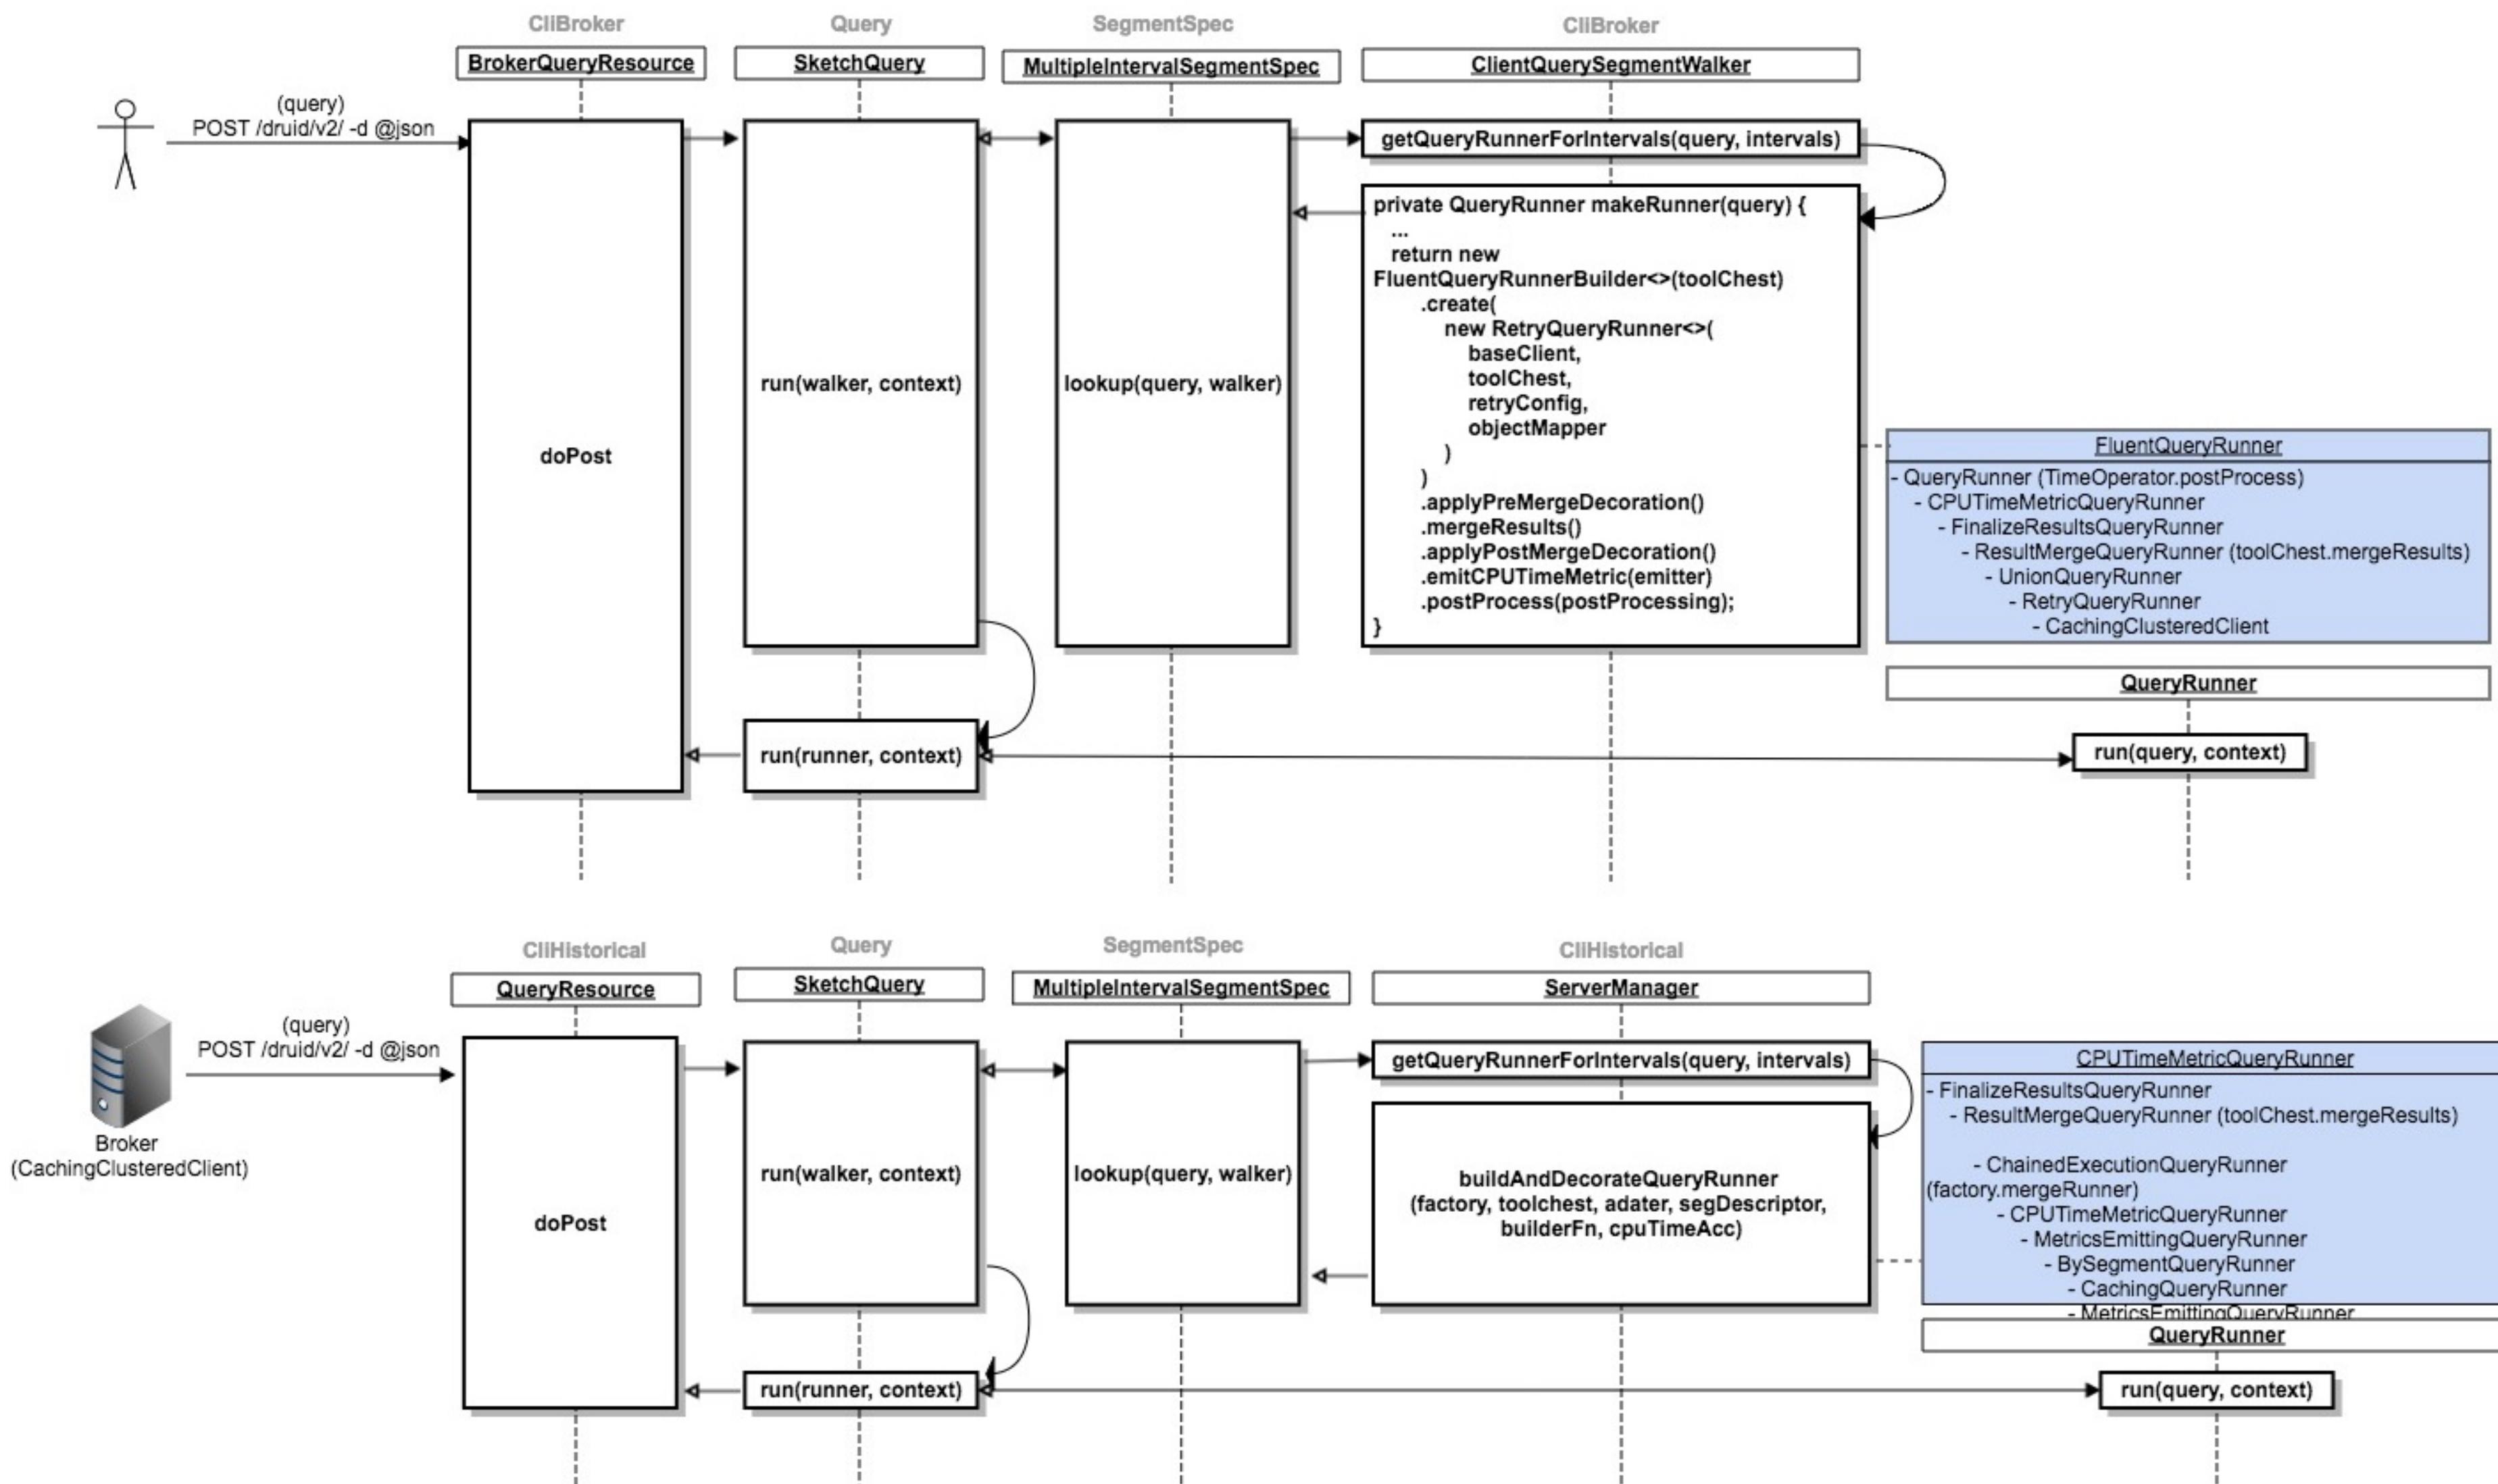
# How Kafka Indexing Service Works

[Kafka Indexing Service Conf.]

```
segmentGranularity: 1H
taskDuration: 1H
intermediatePersistPeriod: 10M
maxRowsInMemory: 600000
maxRowsPerSegment: 7000000
```



# Druid Query Flow



**Thank you**

